

**ECE 5730**  
**Memory Systems**  
**Spring 2009**

**Cache Power Management**



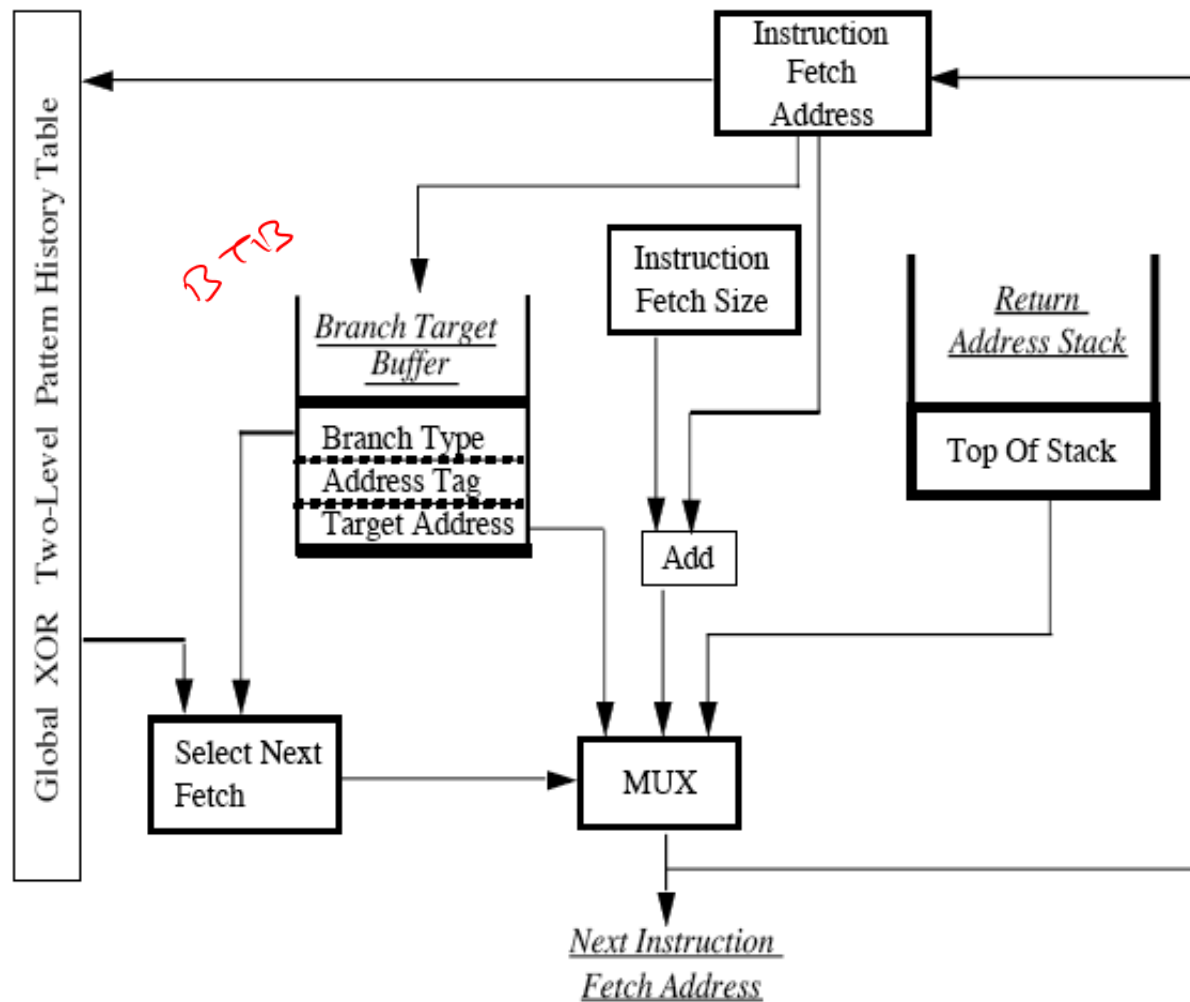
Cornell University

# Announcements

- Quiz 5 on Tuesday
- Quiz 4
  - Average = 8.9/10
- No office hours today

# I-cache Way Prediction

- Conventional instruction fetch mechanism

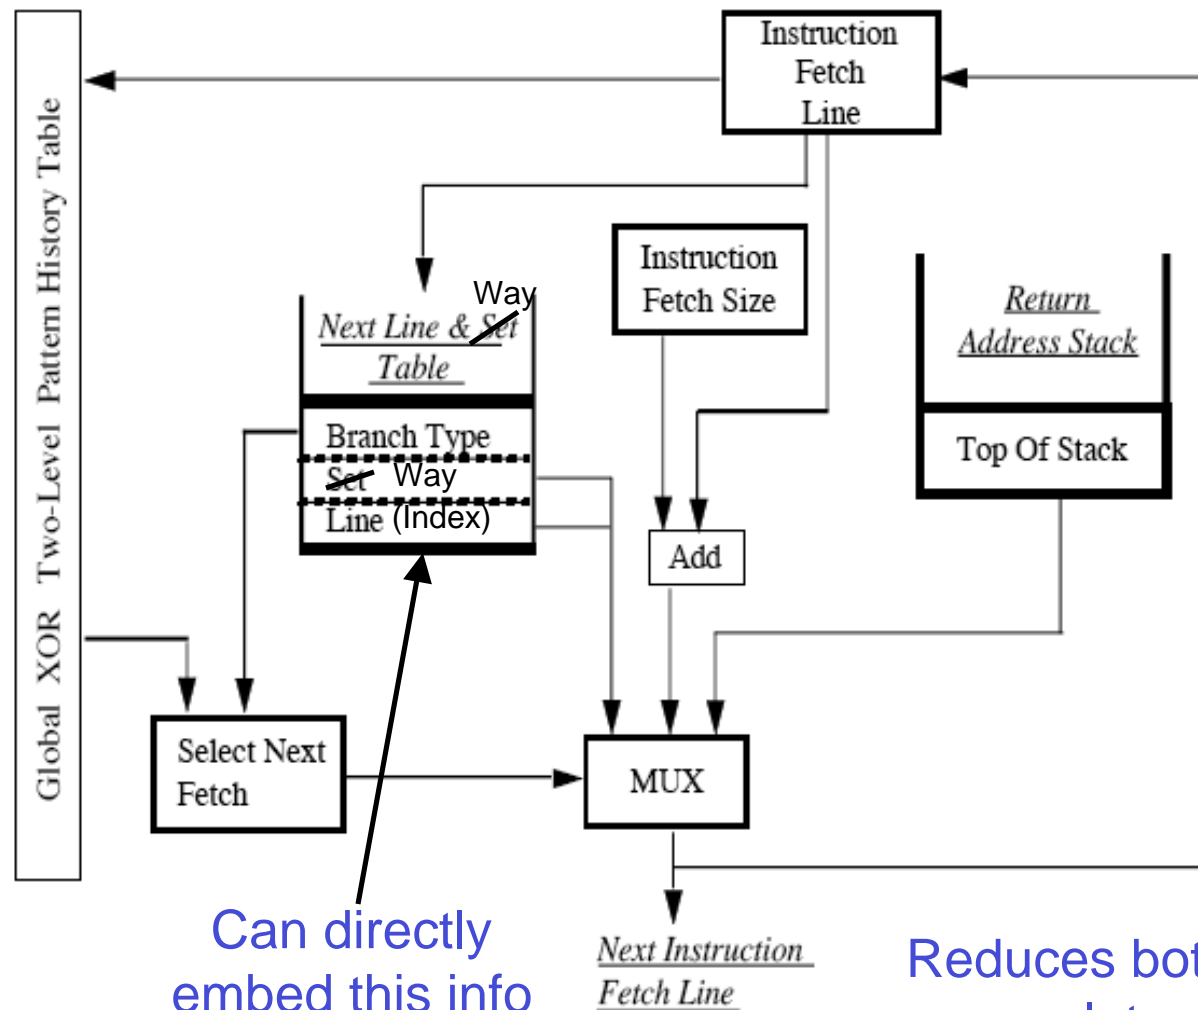


# I-cache Way Prediction

← BTB does not predict cache stuff. This scheme fixes

- Predicting the way, and the next line *index*

that by  
replacing the  
BTB



Can directly  
embed this info  
into the I-cache

Reduces both  
access latency  
and energy!

# Cache Power Management

- **Circuit techniques**

- Transistor sizing, multi- $V_t$ , low-swing bitlines, etc

- **Microarchitecture techniques**

- Static techniques: banking, phased tag/data access, way prediction, CAM-tags

- today* { – **Dynamic techniques: gated-Vdd, cache decay, drowsy caches**

- **Compiler techniques**

- Data partitioning to enable sleep mode

# Dynamic Techniques for Leakage

- Static microarchitectural techniques address dynamic power by reducing activity

- But lots of leaking idle cache rows!

leakage current causes power loss (up to 50% in a cache)

- Three example microarchitectural approaches

- **Gated-Vdd**

- Gate the supply-to-ground path → adds more transistors in the way to minimize leakage

- **Cache decay**

as Gated-Vdd

- Same gating mechanism, but different control policy

- **Drowsy caches** → prevents dynamic power loss (no misses since data is retained)

- Reduce the Vdd in order to retain cell state

→ drop the voltage, retains the data

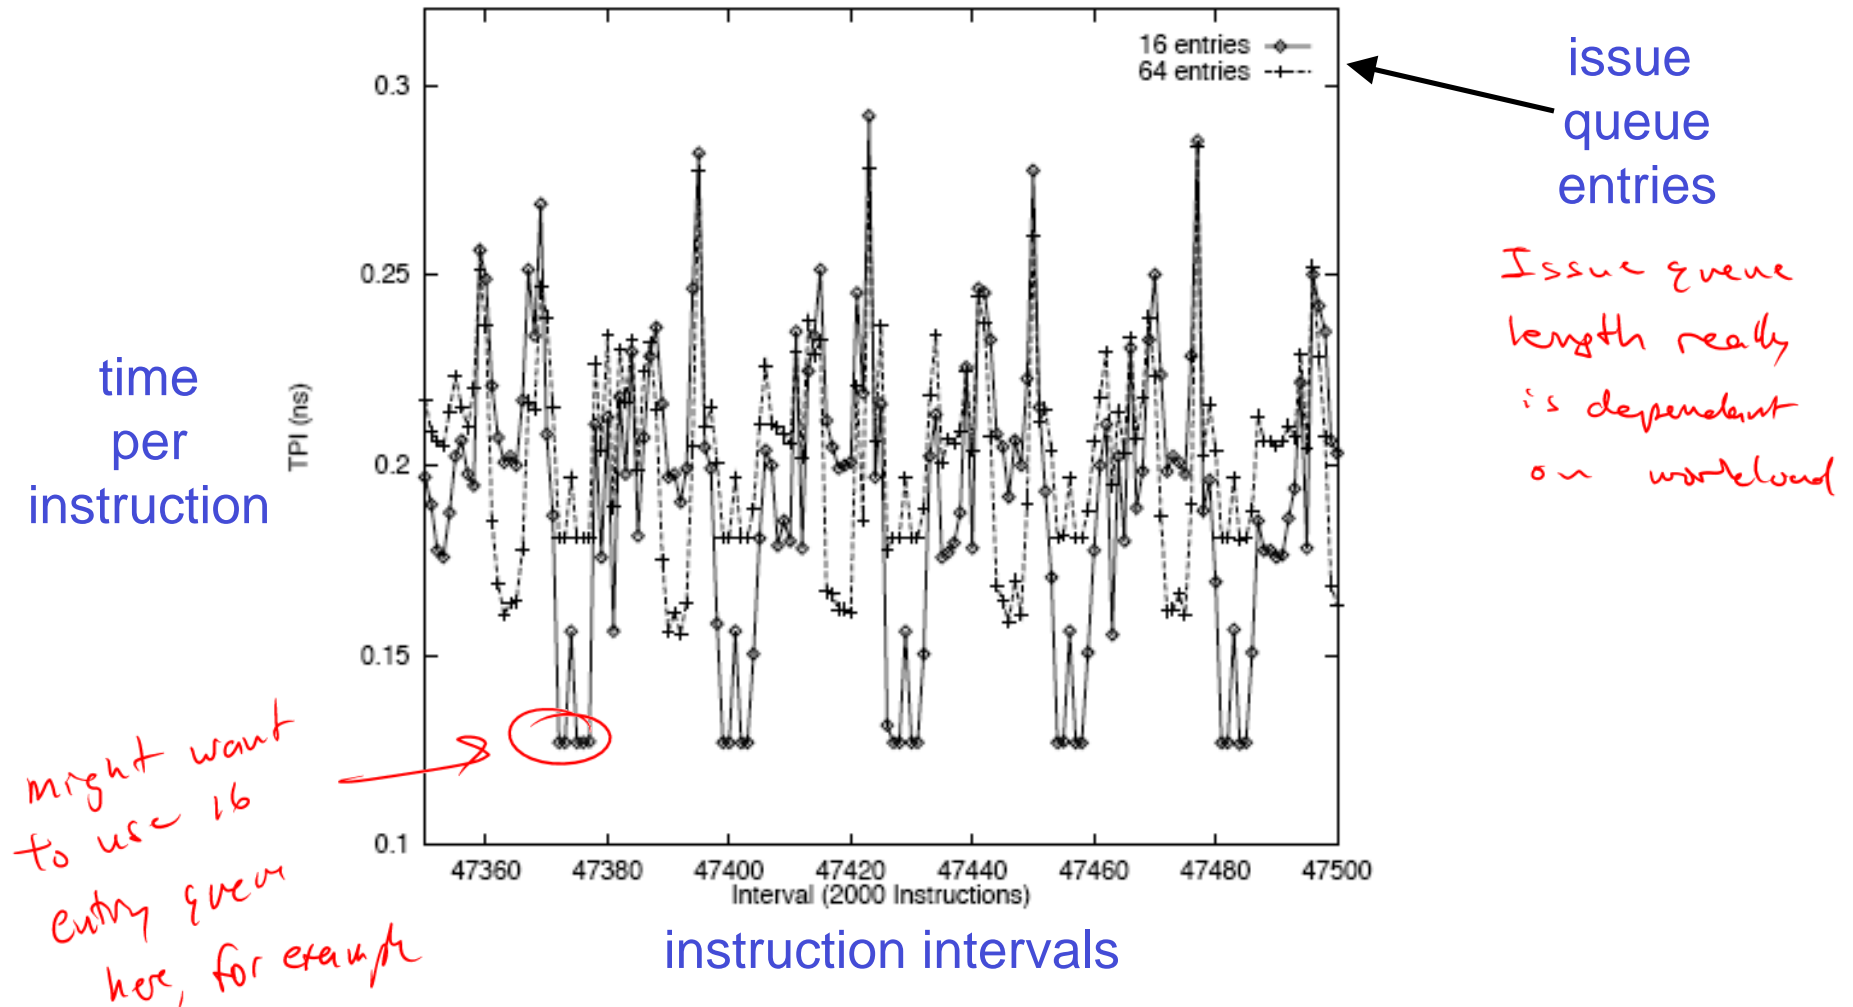
still some wake up latency, but faster than a miss

doesn't retain data  
↑

prevents static power loss, but dynamic power loss on miss.

# Resizable Hardware

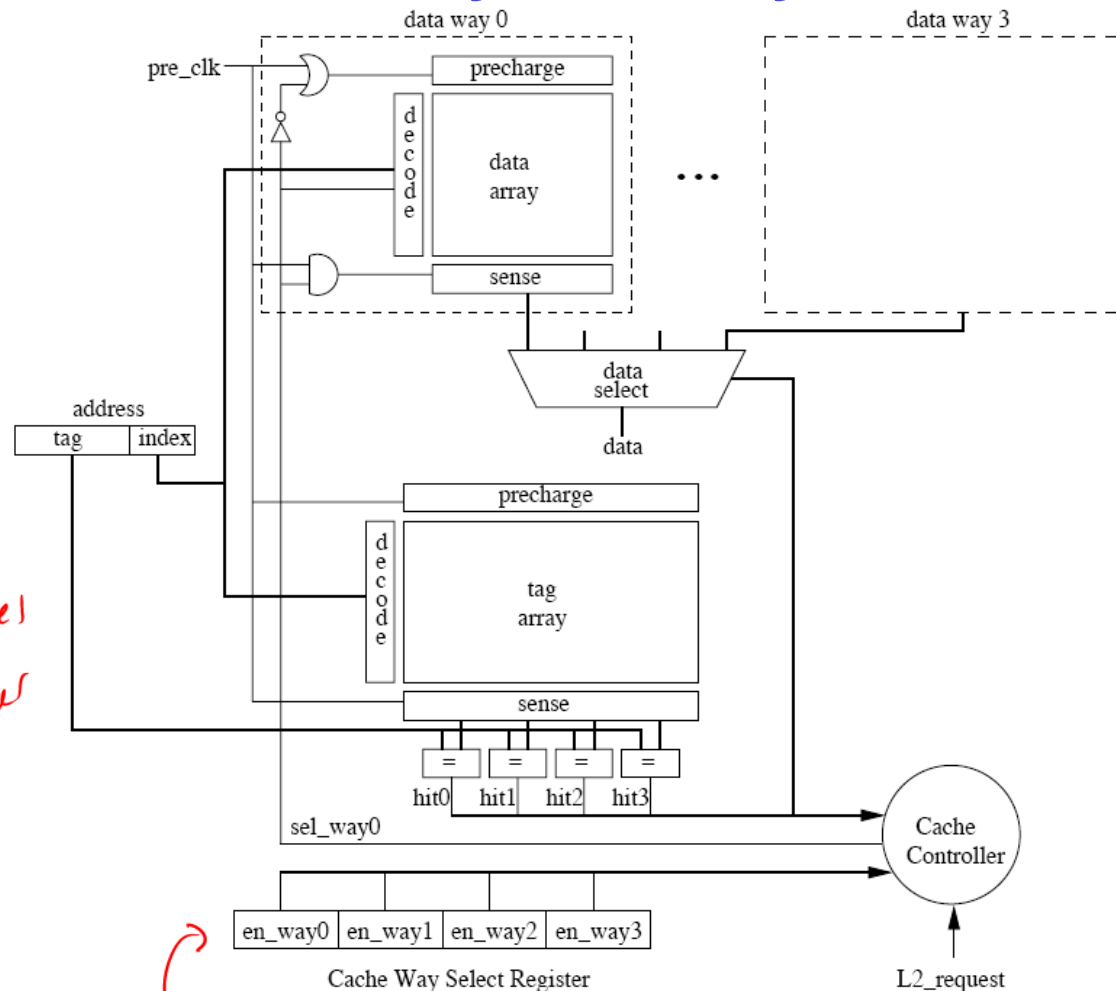
- Demand for hardware resources varies among applications, and within a single application



[Albonesi98]

# Resizable Caches

- *Resizable caches* turn off portions of the cache that are not heavily used by the running program



this address  
dynamic power

enable bits for cache ways  
could thus switch from direct-mapped cache  
all the way to 4-way set associative.




# Gated-Vdd

✓ finer grained control  
by set instead of way  
as in previous slide

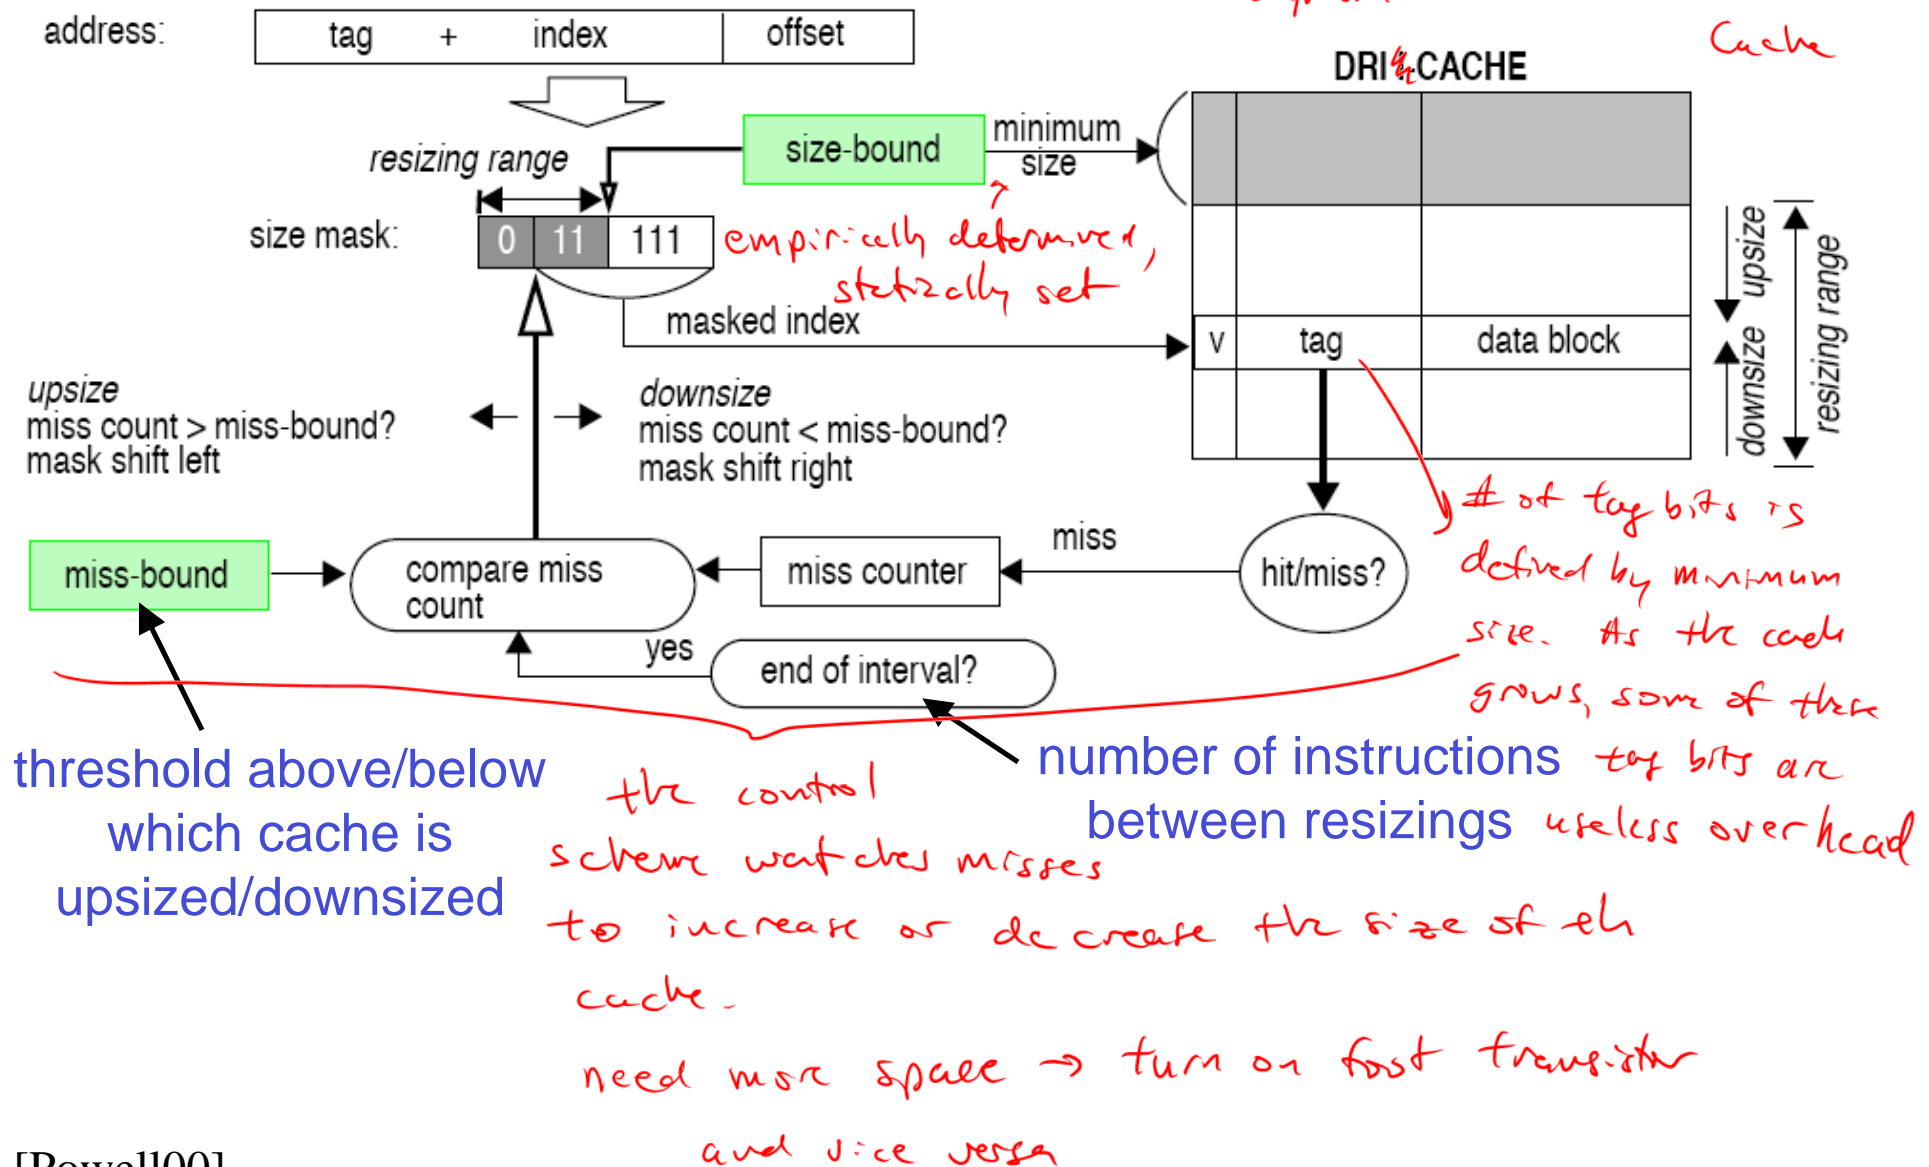
- Dynamically resize the cache (number of sets)
- Sets are disabled by gating the path between Vdd and ground (“stacking effect”)

Even though the  
first transistor  
leaks, by adding  
another series  
transistor you  
add series  
“resistance”  
thus limiting the  
leakage current



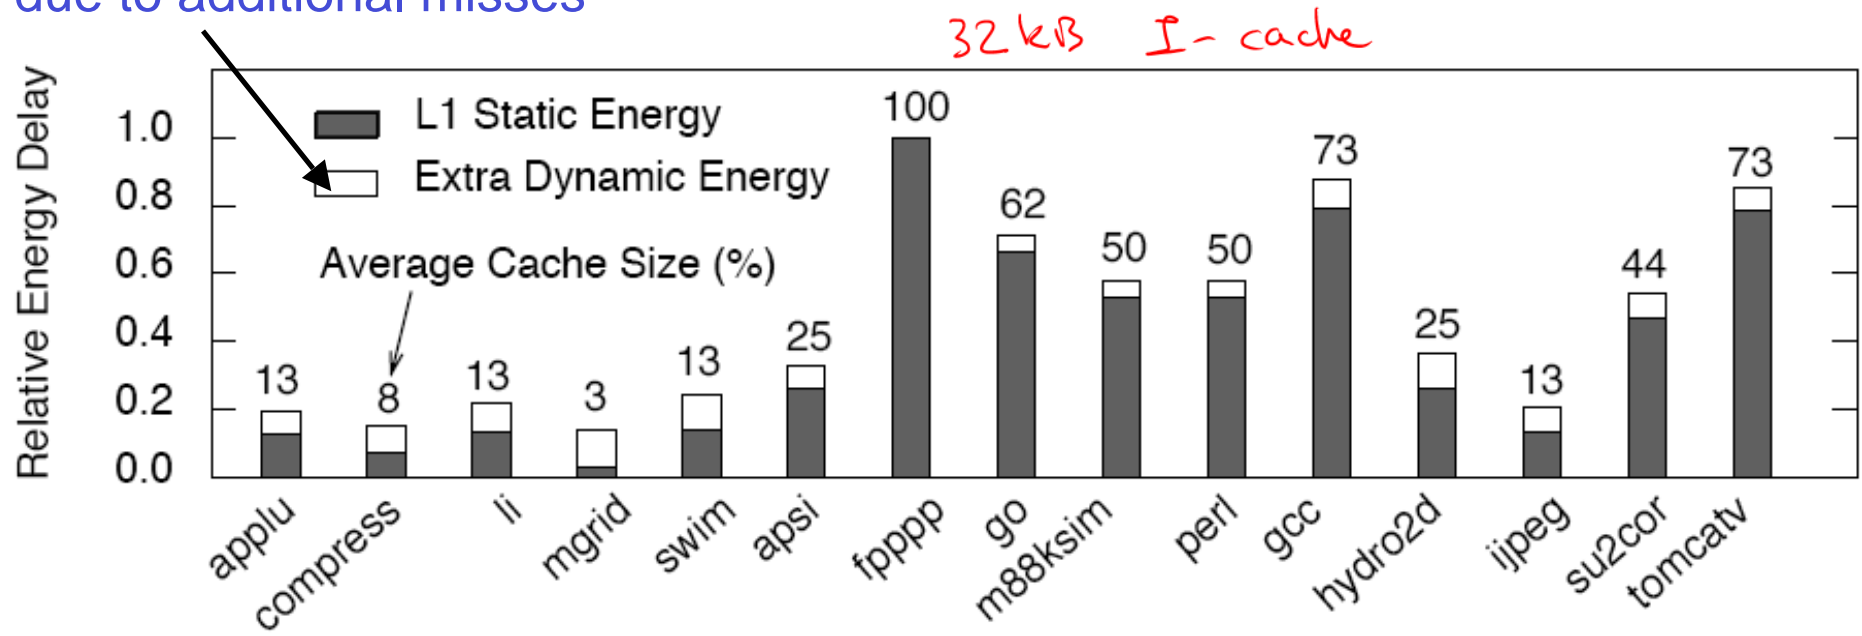
# Gated-Vdd Microarchitecture

*Dynamic Resizable Instruction Cache*



# Gated-Vdd I-cache Effectiveness

essentially performance (energy)  
due to additional misses

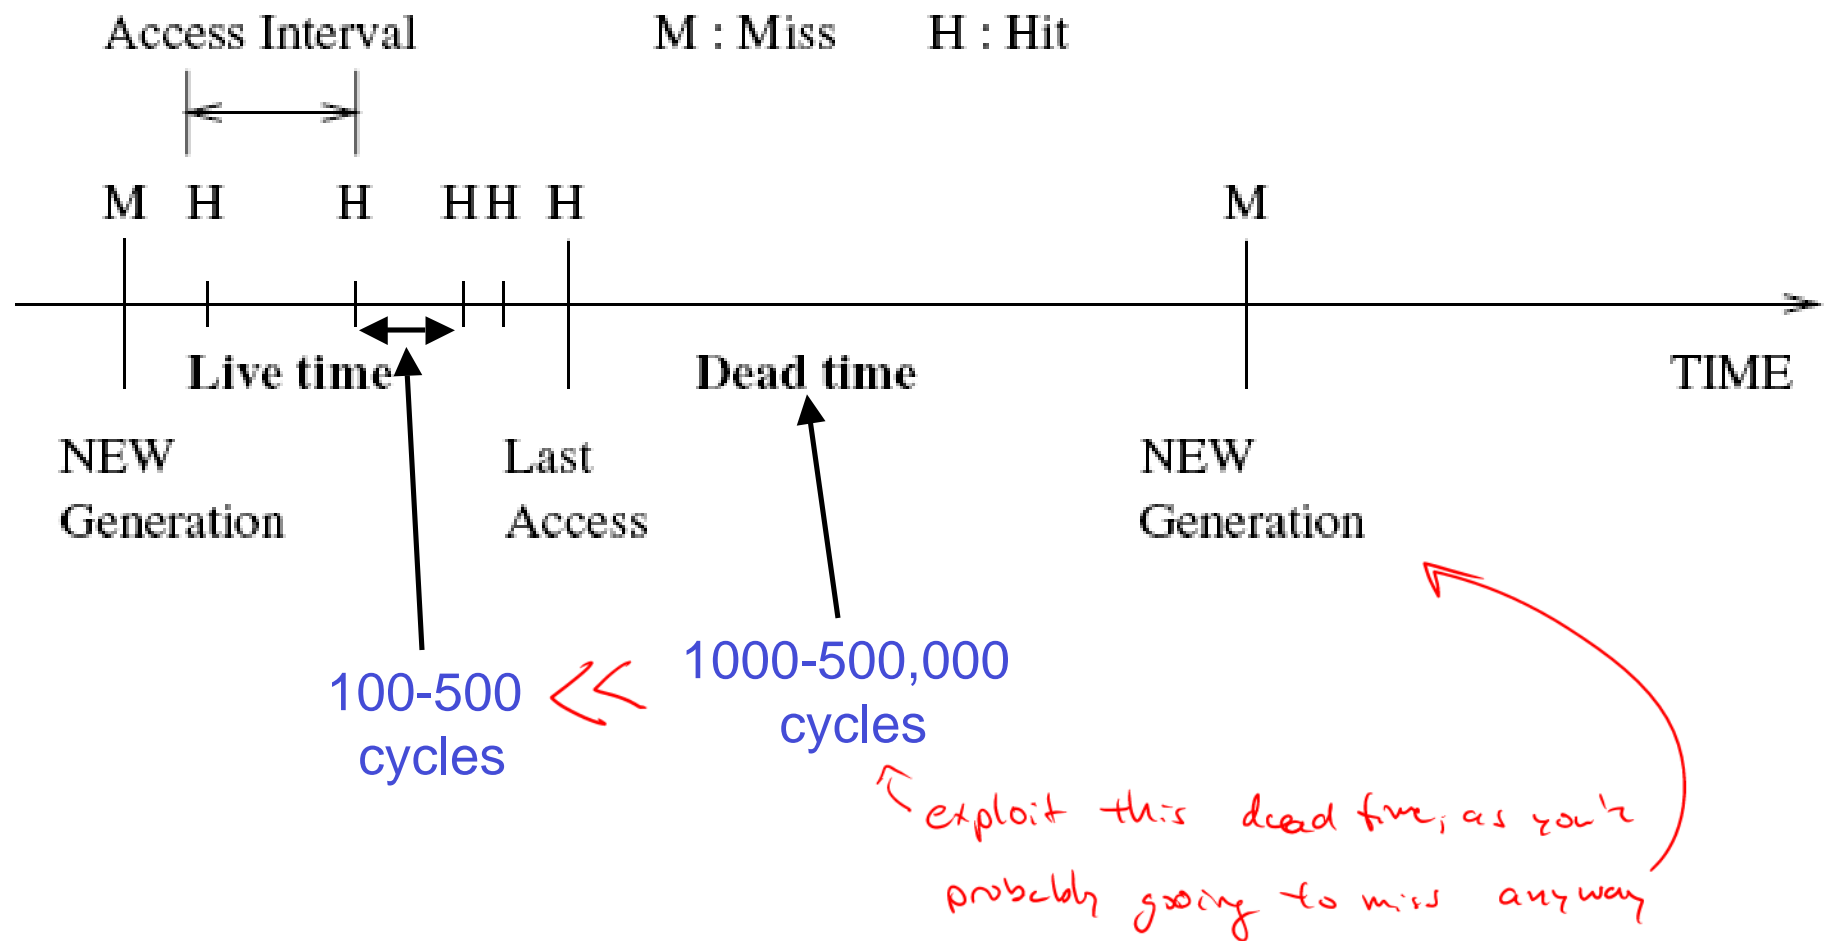


over time, observe the behavior of things in a cache.

access typically occurs  
in bursts

# Cache Decay


- Exploits *generational behavior* of cache contents



# Cache Decay


- Fraction of time cache lines are “dead”

*a lot of the time, cache lines are dead*



32KB L1 D-cache

gray code:  
only one bit  
changes on  
update

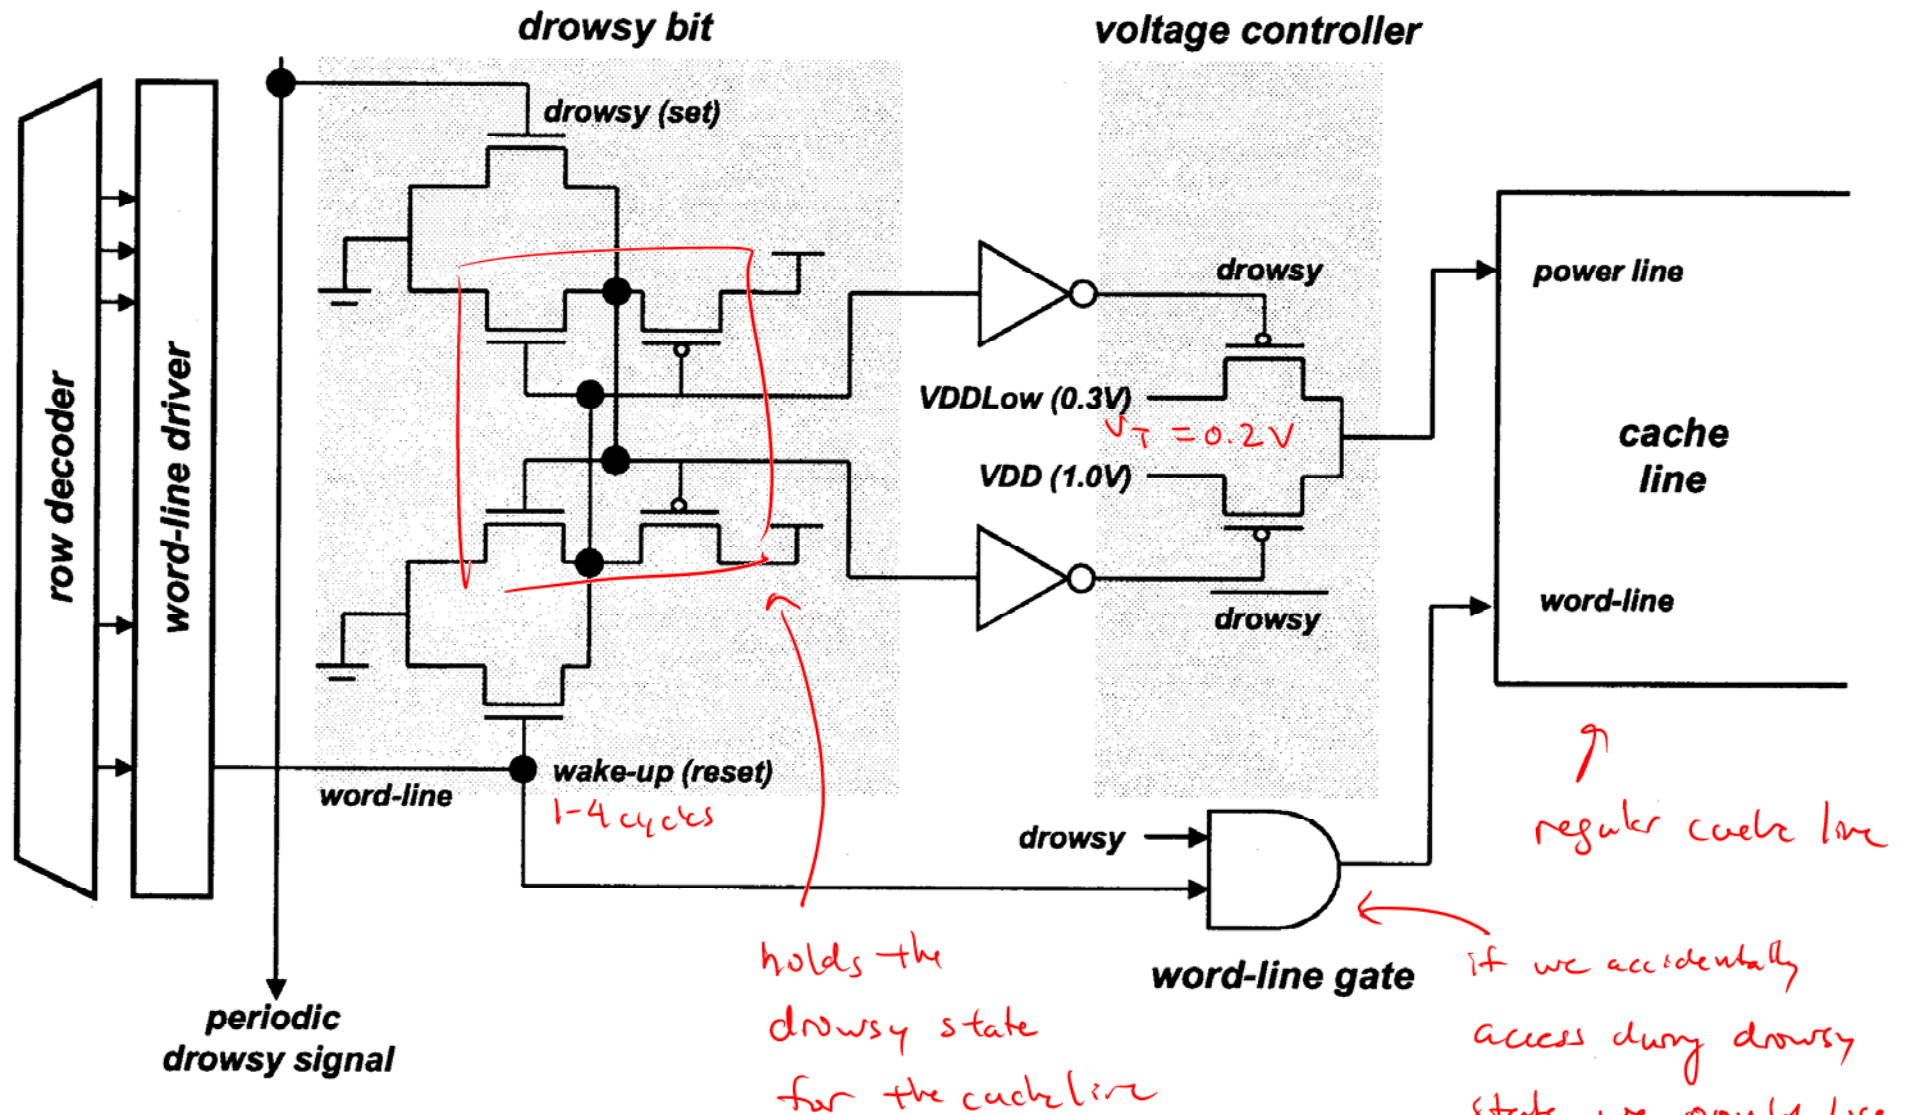


turn it back on, on access

# Drowsy Caches

- **Gated-Vdd cells lose their state**
  - Instructions/data must be refetched
  - Dirty data must be first written back
- **By *dynamically scaling Vdd*, cell is put into a *drowsy* state where it retains its value**
  - Leakage drops superlinearly with reduced Vdd *← exponential*
  - Cell can be fully restored in a few cycles
  - Much lower misprediction cost than gated-Vdd, but noise susceptibility and less reduction in leakage *← misprediction on drowsy, not hit/miss* **vs gated Vdd**


# Drowsy Cache Organization






# Drowsy Cache Effectiveness

(assumed 1-cycle  
wakeup time)




drowsy cache  
turns drowsy  
also

# Drowsy Cache Performance Cost

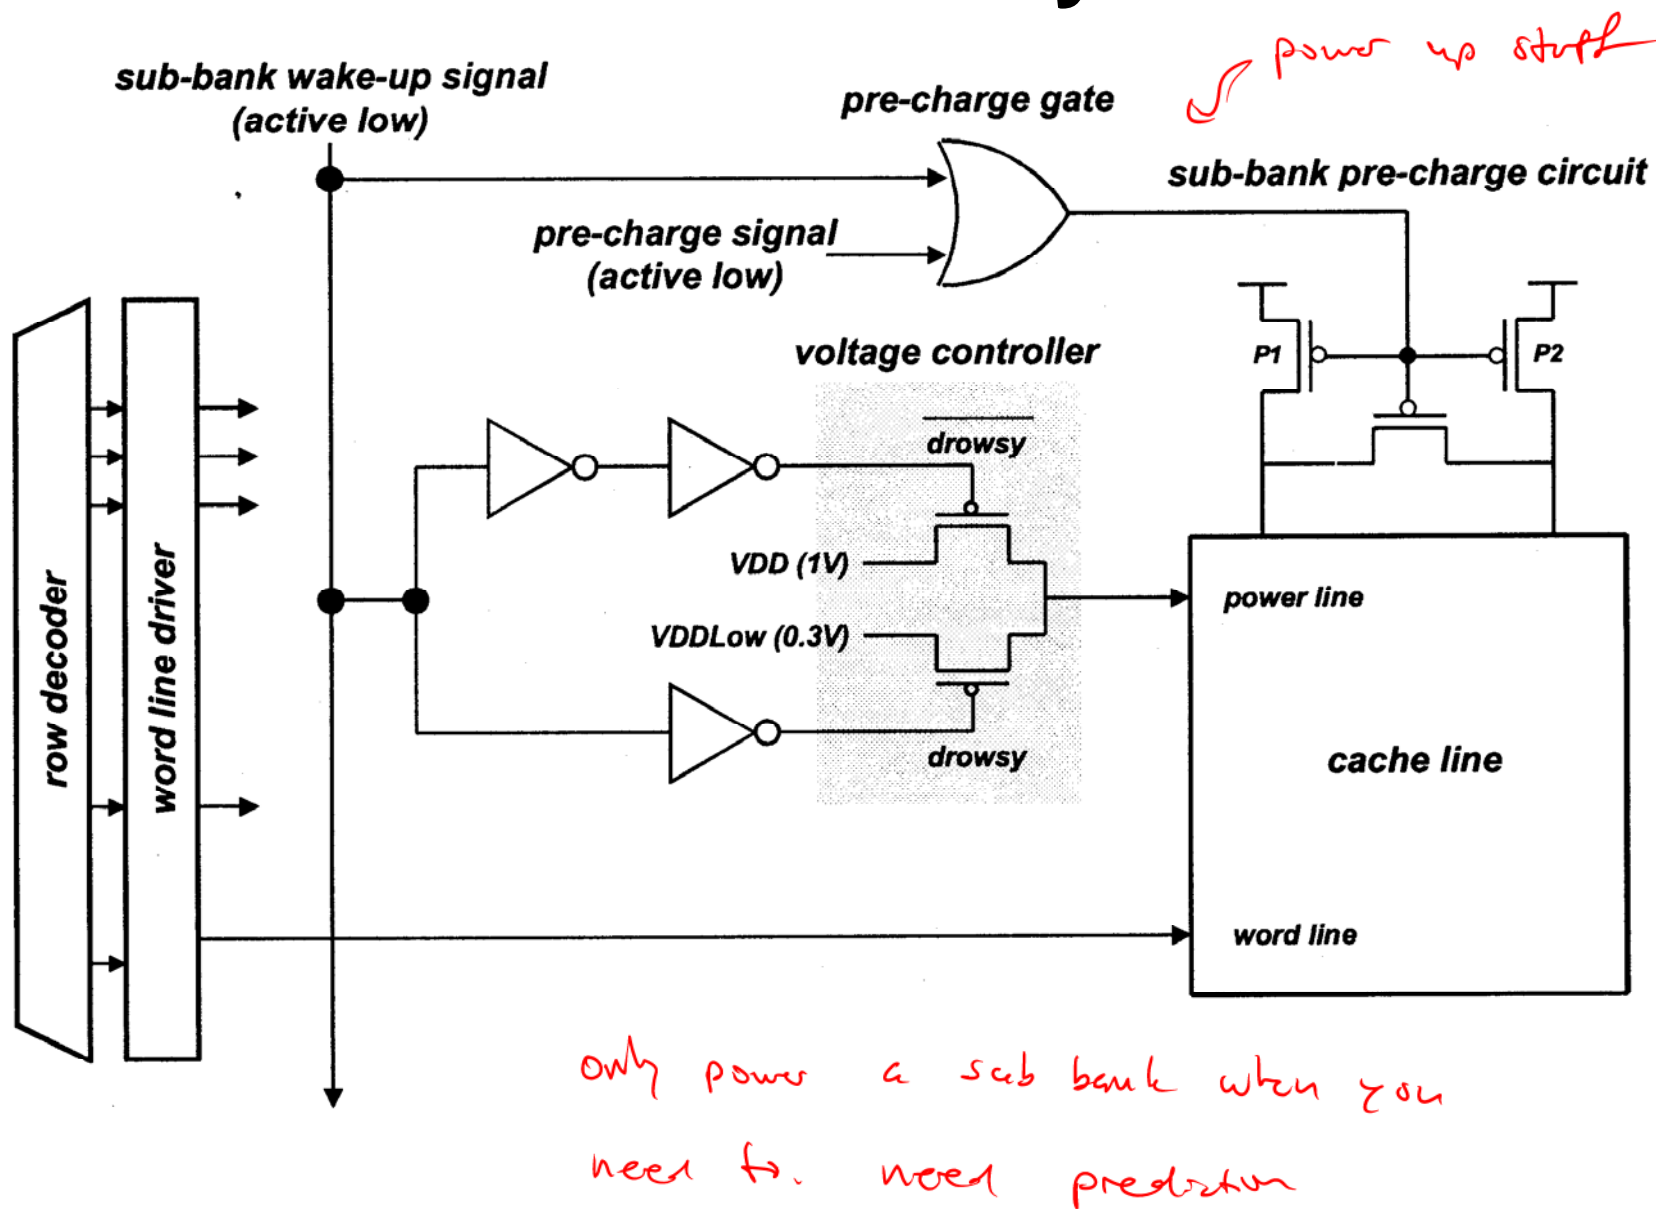


*drowsy cache  
is bad for  
I-cache.  
more temporal  
locality for  
data, must  
do something  
new for  
I cache*

# Sub-banked Drowsy I-cache



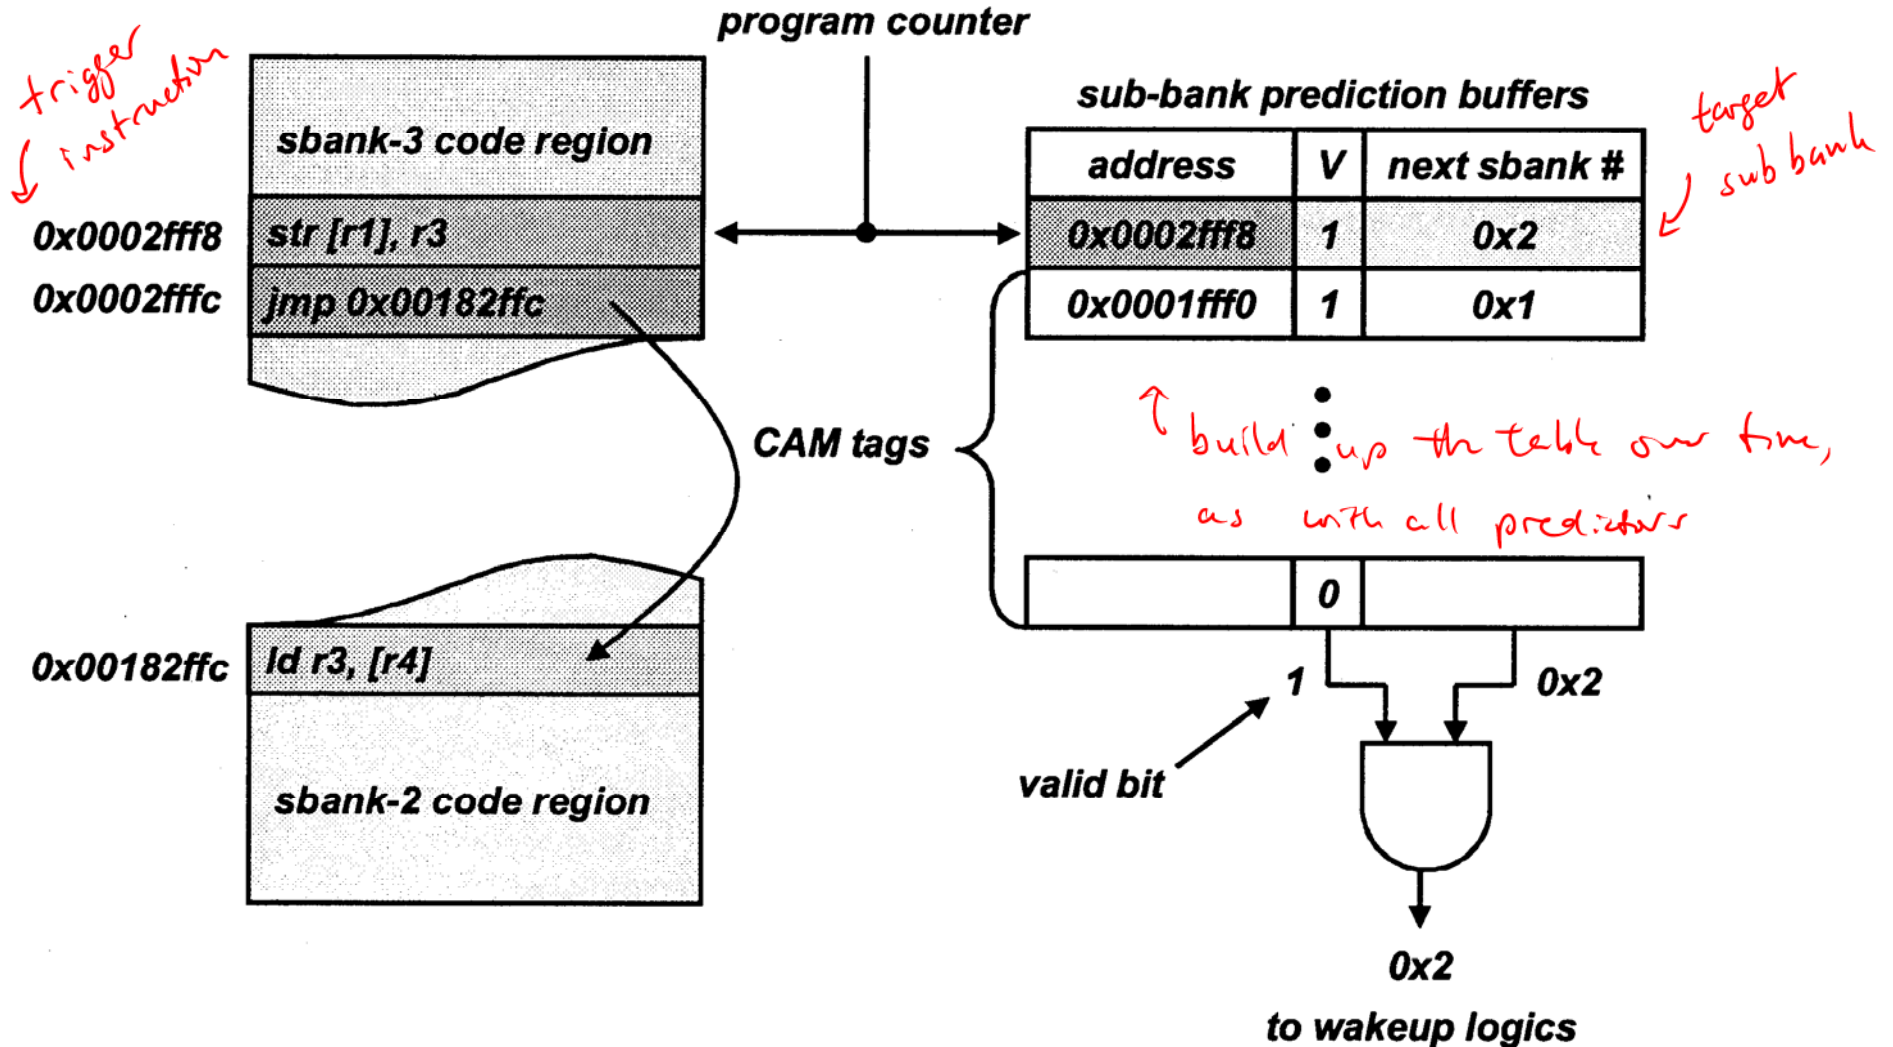
# Sub-banked Drowsy I-cache





# Next Sub-bank Prediction


- Predict fetch transition to a new sub-bank



# Next Sub-bank Prediction

- Can embed predictions with cache tags

*instead of building a whole predictor table, just add the bits to the cache tag*



~75% prediction accuracy

# Drowsy D-cache vs. D-cache Decay

*slightly worse performance* (pointing to Drowsy D-cache)

*lots more leakage* (pointing to D-cache Decay)

FP	Run-time impact (%)		Norm. leakage (%)		INT	Run-time impact (%)		Norm. leakage (%)	
	drowsy	decay	drowsy	decay		drowsy	decay	drowsy	decay
ammp	0.54	0.12	34 (22)	81	bzip2	0.77	0.38	34 (22)	35
applu	0.50	0.00	37 (26)	41	crafty	0.46	1.08	39 (28)	114
apsi	0.32	0.02	35 (23)	31	eon	0.36	0.13	38 (26)	38
art	0.61	0.01	38 (27)	44	gap	0.46	0.90	32 (20)	31
equake	0.03	0.04	33 (21)	10	gcc	0.06	0.00	53 (45)	96
facerec	0.15	0.03	32 (20)	23	gzip	2.07	1.23	34 (22)	85
fma3d	0.14	0.87	32 (20)	141	mcf	0.66	0.04	34 (22)	35
galgel	0.31	0.00	51 (42)	68	parser	1.17	1.75	35 (22)	63
lucas	0.66	0.00	34 (22)	25	perl	0.87	6.97	36 (24)	224
mesa	0.21	0.09	34 (22)	22	twolf	0.87	0.29	35 (22)	36
mgrid	0.49	0.00	40 (28)	42	vortex	0.25	0.78	36 (25)	126
sixtrack	0.34	0.18	39 (27)	56	vpr	0.89	0.84	36 (24)	76
swim	0.61	0.00	37 (25)	34	avg	0.57	0.64	37 (25)	59
wupwise	0.20	0.01	34 (22)	50					

↑  
with high Vt transistors

# Drowsy I-cache vs. I-cache Decay

FP	Run-time impact (%)		Norm. leakage (%)		INT	Run-time impact (%)		Norm. leakage (%)	
	drowsy	decay	drowsy	decay		drowsy	decay	drowsy	decay
ammp	0.00 (99)	1.83	25	28	bzip2	0.83 (52)	0.69	24	8
applu	0.00 (99)	0.22	24	10	crafty	3.75 (71)	6.26	26	118
apsi	2.01 (67)	0.14	25	25	eon	2.54 (70)	0.31	25	39
art	0.00 (99)	0.21	24	1	gap	1.93 (70)	0.55	25	40
equake	0.71 (87)	0.07	25	19	gcc	0.03 (78)	2.37	24	15
facerec	0.39 (76)	0.26	24	5	gzip	1.00 (87)	0.24	25	6
fma3d	4.52 (74)	0.08	26	54	mcf	0.19 (81)	0.23	24	3
galgel	0.00 (27)	0.04	24	1	parser	0.51 (72)	0.44	24	8
lucas	0.01 (97)	0.00	24	2	perl	2.34 (62)	1.84	25	68
mesa	0.67 (86)	0.37	25	23	twolf	0.30 (74)	2.98	24	56
mgrid	0.01 (99)	0.10	24	19	vortex	4.34 (63)	8.43	26	147
sixtrack	1.23 (86)	0.45	25	15	vpr	0.00 (99)	0.16	24	5
swim	0.00 (99)	0.28	24	9	avg	0.79 (74)	0.99	25	23
wupwise	0.79 (79)	0.32	24	20					

[Kim04]

*% accuracy of the predictor*



# Cache Power Management

- **Circuit techniques**

- Transistor sizing, multi- $V_t$ , low-swing bitlines, etc

- **Microarchitecture techniques**

- Static techniques: banking, phased tag/data access, way prediction, CAM-tags
- Dynamic techniques: gated-Vdd, cache decay, drowsy caches

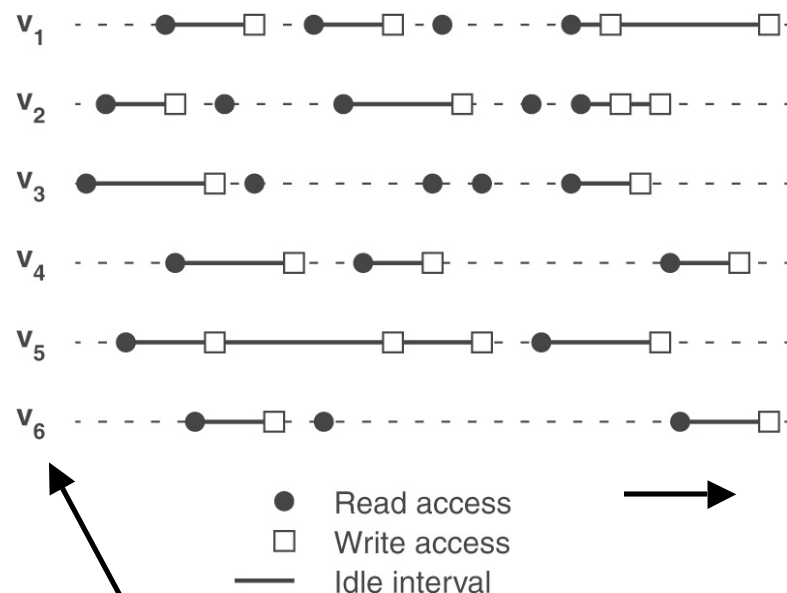
- **Compiler techniques**

- Data partitioning to enable sleep mode

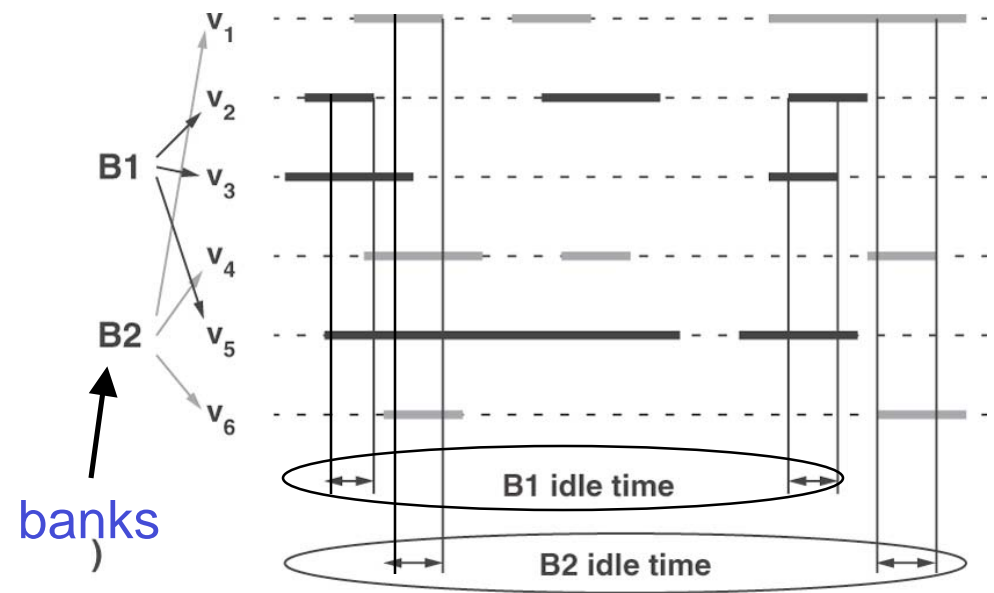
→ add sleep enable/disable instructions to an ISA  
drowsy

# Compiler-based Data Partitioning

- Multiple D-cache banks, each with sleep mode
- Lifetime analysis used to assign commonly idle data to the same bank



variables



# Next Time

## Cache Case Studies