
Chapter 3: Computer Arithmetic

- **Not implement ALU**
 - **Domain knowledge**
- **Do think about**
 - **Representing numbers (and ALU)**
 - **ISA from data perspective**
- **(optional) subword parallelism**

Big Picture

- ❑ Part 1: what is computer, CSE, computer architecture?
 - Fundamental concepts and principles
- ❑ Part 2: ISA ([externals](#))
 - Ch. 1: performance
 - Exe. time, benchmark, model, RISC, power, multicore
 - Ch. 2: language of computer; ISA
 - What is a good ISA? Today's RISC-style ISA (MIPS)
 - Ch. 3: computer arithmetic ([key internal building block](#))
 - Data representation and ALU, ISA – data perspective
- ❑ Part 3: implementation of ISA ([internals](#))
 - Ch. 4: processor
 - Ch. 5: memory system

Numbers

- ❑ Data types (INT and FP)
 - Numbers int,fp
 - Text, audio, image, video, and so on int
- ❑ Representing numbers and ALU
 - Positive and negative integers
 - Numbers are finite (overflow)
 - Scientific numbers (too big or too small)
- ❑ ISA from data perspective

Possible Representations

Sign Magnitude: One's Complement Two's Complement

$$000 = +0$$

$$001 = +1$$

$$010 = +2$$

$$011 = +3$$

$010(2) + 110(-2)$
 $= 100$
 \rightarrow correct

$$100 = -0$$

$$101 = -1$$

$$110 = -2$$

$$111 = -3$$

$010(2) + 101(-2)$
 $= 111 (-0)$

$$000 = +0$$

$$001 = +1$$

$$010 = +2$$

$$011 = +3$$

$$100 = -3$$

$$101 = -2$$

$$110 = -1$$

$$111 = -0$$

$$000 = +0$$

$$001 = +1$$

$$010 = +2$$

$$011 = +3$$

$$100 = -4$$

$$101 = -3$$

$$110 = -2$$

$$111 = -1$$

- ❑ Issues: balance, number of zeros, ease of operations
- ❑ Which one is best? Why?

Possible Representations

Biased Notation:

$$000 = -3$$

$$001 = -2$$

$$010 = -1$$

$$011 = 0$$

$$100 = +1$$

$$101 = +2$$

$$110 = +3$$

$$111 = +4$$

(Bias 3)

Biased Notation:

$$000 = -4$$

$$001 = -3$$

$$010 = -2$$

$$011 = -1$$

$$100 = 0$$

$$101 = +1$$

$$110 = +2$$

$$111 = +3$$

(Bias 4)

Unsigned:

$$000 = +0$$

$$001 = +1$$

$$010 = +2$$

$$011 = +3$$

$$100 = +4$$

$$101 = +5$$

$$110 = +6$$

$$111 = +7$$

MIPS (반복)

□ 32 bit signed numbers:

| | | | | | | | | | | | |
|------|------|------|------|------|------|------|------|-----------------|---|--------------------------------|---------------|
| 0000 | 0000 | 0000 | 0000 | 0000 | 0000 | 0000 | 0000 | $_{\text{two}}$ | = | 0_{ten} | |
| 0000 | 0000 | 0000 | 0000 | 0000 | 0000 | 0000 | 0001 | $_{\text{two}}$ | = | $+ 1_{\text{ten}}$ | |
| 0000 | 0000 | 0000 | 0000 | 0000 | 0000 | 0000 | 0010 | $_{\text{two}}$ | = | $+ 2_{\text{ten}}$ | |
| ... | | | | | | | | | | | |
| 0111 | 1111 | 1111 | 1111 | 1111 | 1111 | 1111 | 1110 | $_{\text{two}}$ | = | $+ 2,147,483,646_{\text{ten}}$ | <i>maxint</i> |
| 0111 | 1111 | 1111 | 1111 | 1111 | 1111 | 1111 | 1111 | $_{\text{two}}$ | = | $+ 2,147,483,647_{\text{ten}}$ | |
| 1000 | 0000 | 0000 | 0000 | 0000 | 0000 | 0000 | 0000 | $_{\text{two}}$ | = | $- 2,147,483,648_{\text{ten}}$ | <i>minint</i> |
| 1000 | 0000 | 0000 | 0000 | 0000 | 0000 | 0000 | 0001 | $_{\text{two}}$ | = | $- 2,147,483,647_{\text{ten}}$ | |
| 1000 | 0000 | 0000 | 0000 | 0000 | 0000 | 0000 | 0010 | $_{\text{two}}$ | = | $- 2,147,483,646_{\text{ten}}$ | |
| ... | | | | | | | | | | | |
| 1111 | 1111 | 1111 | 1111 | 1111 | 1111 | 1111 | 1101 | $_{\text{two}}$ | = | $- 3_{\text{ten}}$ | |
| 1111 | 1111 | 1111 | 1111 | 1111 | 1111 | 1111 | 1110 | $_{\text{two}}$ | = | $- 2_{\text{ten}}$ | |
| 1111 | 1111 | 1111 | 1111 | 1111 | 1111 | 1111 | 1111 | $_{\text{two}}$ | = | $- 1_{\text{ten}}$ | |

Two's Complement Operations (반복)

- ❑ Negating two's complement number: invert all bits and add 1
 - Remember: “negate” and “invert” are quite different!
- ❑ Signed and unsigned numbers
 - Memory address: start at 0 and end at the largest (e.g., FFFF)
 - Negative address make no sense
 - Some programming languages make this distinction
 - ⇒ C: **int** versus **unsigned int**
 - MIPS comparison instructions
 - ⇒ signed: **slt, slti**
 - ⇒ unsigned: **sltu, sltiu**

Signed and unsigned loads (반복)

- ❑ Signed load: copy sign repeatedly to fill rest of register
 - e.g., two's complement numbers
 - ❑ Unsigned load: fill with 0s to left of data
 - When bit pattern is unsigned
 - ❑ MIPS instructions
 - **lb** versus **lbu**
 - **lh** versus **lhu**
 - **lw**
- † No concern to store

Addition & Subtraction

- ❑ Just like in grade school (carry/borrow 1s)

$$\begin{array}{r} 0111 \\ + 0110 \\ \hline \end{array} \qquad \begin{array}{r} 0111 \\ - 0110 \\ \hline \end{array} \qquad \begin{array}{r} 0110 \\ - 0101 \\ \hline \end{array}$$

- ❑ Two's complement operations easy

- subtraction using addition of negative numbers

$$\begin{array}{r} 0111 \\ + 1010 \\ \hline \end{array}$$

- ❑ Arithmetic overflow (result too large for finite word):

- e.g., adding two n-bit numbers not yield an n-bit number

$$\begin{array}{r} 0111 \\ + 0001 \\ \hline 1000 \\ \hline \end{array}$$

Two's Complement Arithmetic Overflow

- ❑ No overflow when adding a positive and a negative number
- ❑ Overflow occurs when the value affects the sign:
 - overflow when adding two positives yields a negative
 - or, adding two negatives gives a positive
- ❑ Consider the operations $A + B$, and $A - B$
 - Can overflow occur if B is 0 ? no
 - Can overflow occur if A is 0 ? yes $B=2^n$
- ❑ Overflow detection using a single exclusive-or gate

two's complement overflow
 0 1
 1 0
 carry bit
 exclusive or
 detect
 -> detect
 (exception) interrupt

| | | |
|-------|------|------|
| 70 | 0110 | 0110 |
| + 80 | 0101 | 0000 |
| ----- | | |
| 150 | 1001 | 0110 |

| | | |
|-------|------|------|
| -70 | 1011 | 1010 |
| -80 | 1011 | 0000 |
| ----- | | |
| -150 | 0110 | 1010 |

Effects of Overflow

- ❑ An exception (interrupt) occurs
 - Control jumps to predefined address for exception
 - Interrupted address is saved for possible resumption
- ❑ Given overflow, what does OS do?
- ❑ Arithmetic overflow: program bug
 - What do programmers should do? variable
 - Flight control vs. homework assignment
- ❑ What about unsigned integers?
 - Commonly used for memory addresses (by compiler)
 - 2's complement arithmetic overflows are ignored
- ❑ MIPS solution: provide two kinds of arithmetic instructions

add, addi, sub

addu, addiu, subu

// do not detect overflow

h/w

MIPS Assembly Language: add, subtract

MIPS assembly language

| Category | Instruction | Example | Meaning | Comments |
|------------|--------------------------------|---------------------|----------------------|--|
| Arithmetic | add | add \$s1,\$s2,\$s3 | $\$s1 = \$s2 + \$s3$ | Three operands; overflow detected |
| | subtract | sub \$s1,\$s2,\$s3 | $\$s1 = \$s2 - \$s3$ | Three operands; overflow detected |
| | add immediate | addi \$s1,\$s2,100 | $\$s1 = \$s2 + 100$ | + constant; overflow detected |
| | add unsigned | addu \$s1,\$s2,\$s3 | $\$s1 = \$s2 + \$s3$ | Three operands; overflow undetected |
| | subtract unsigned | subu \$s1,\$s2,\$s3 | $\$s1 = \$s2 - \$s3$ | Three operands; overflow undetected |
| | add immediate unsigned | addiu \$s1,\$s2,100 | $\$s1 = \$s2 + 100$ | + constant; overflow undetected |
| | move from coprocessor register | mfc0 \$s1,\$epc | $\$s1 = \epc | Used to copy Exception PC plus other special registers |


Build ALU

(Topic 1 에서 **abstraction** 이라는 개념을
설명하며 어느 정도 다룬 내용임)

32-bit ALU Design




□ Operations

- Arithmetic: add, subtract, multiply, divide
- Logical: bitwise AND, OR, NOT




32-bit ALU Design (Abstraction)

- Build 1-bit ALU
 - Use 32 of them in parallel





1-bit ALU Design (디지털논리설계; Abstraction)

□ 1-bit ADD, AND, OR




Multiplexor (Data Selector; Abstraction)

- 2-to-1 MUX, 4-to-1 MUX (c.f., Demultiplexer)





1-bit Full Adder Design (디지털논리설계)



| A | B | C _{in} | S | C _{out} |
|---|---|-----------------|---|------------------|
| 0 | 0 | 0 | 0 | 0 |
| 0 | 0 | 1 | 1 | 0 |
| 0 | 1 | 0 | 1 | 0 |
| 0 | 1 | 1 | 0 | 1 |
| 1 | 0 | 0 | 1 | 0 |
| 1 | 0 | 1 | 0 | 1 |
| 1 | 1 | 0 | 0 | 1 |
| 1 | 1 | 1 | 1 | 1 |

$$C_{out} = a \cdot b + a \cdot C_{in} + b \cdot C_{in}$$
$$sum = a \text{ xor } b \text{ xor } C_{in}$$

1-bit Full Adder Design (Abstraction)




$$C_{out} = a \cdot b + a \cdot c_{in} + b \cdot c_{in}$$
$$sum = a \text{ xor } b \text{ xor } c_{in}$$

XOR (Exclusive-OR) Gate (Abstraction)

$$\square C = A \text{ XOR } B = A \oplus B$$

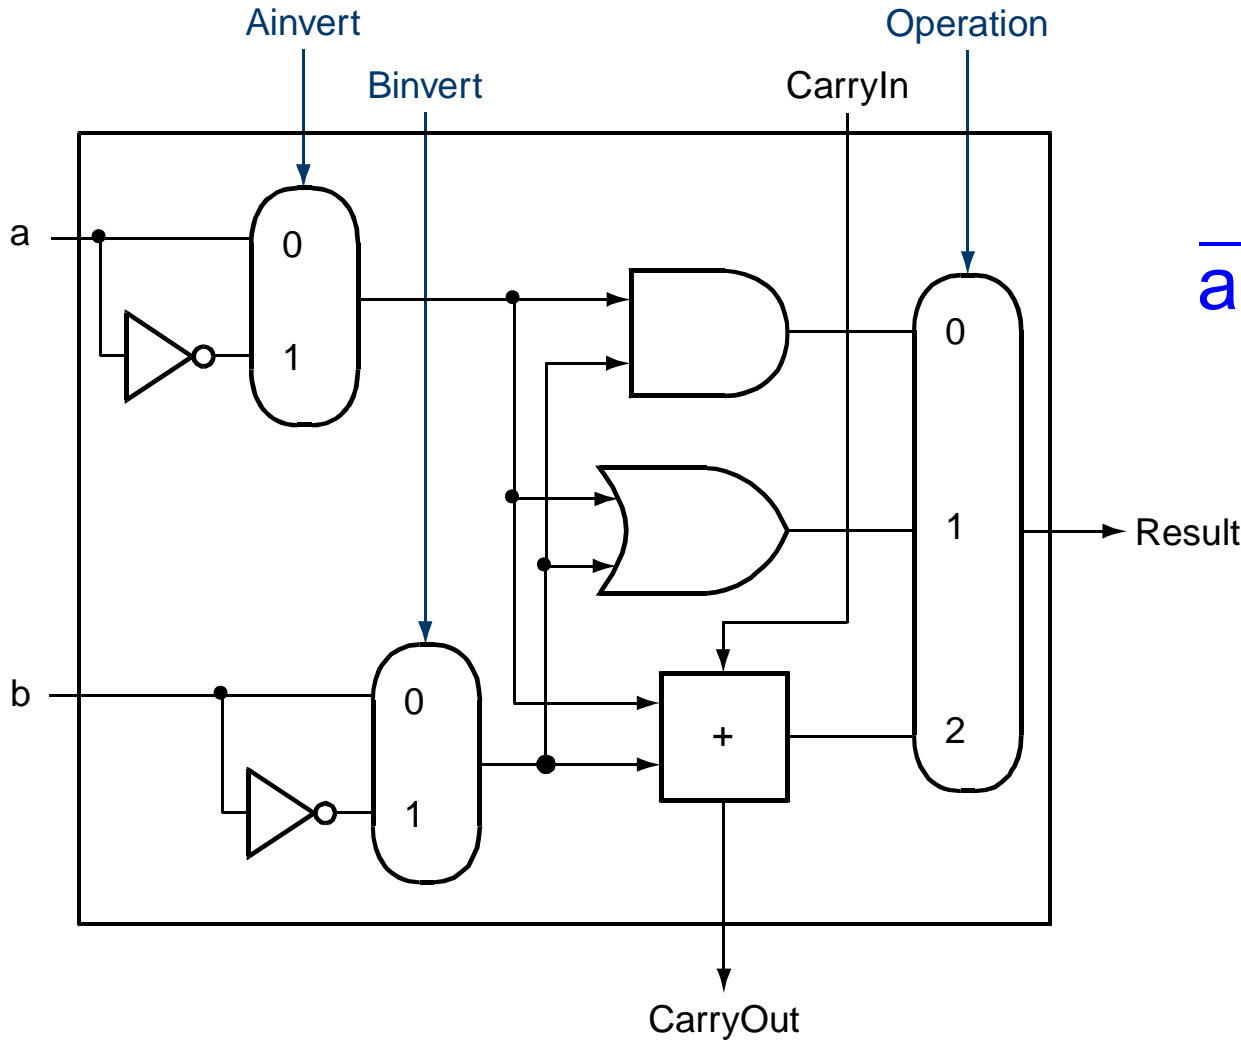
| p | q | $p \oplus q$ |
|----------|----------|--------------------------------|
| 1 | 1 | 0 |
| 1 | 0 | 1 |
| 0 | 1 | 1 |
| 0 | 0 | 0 |



$$C = A \cdot \underline{B} + \underline{A} \cdot B$$

Adding a NOR function

❑ Can also choose to invert **a**. How do we get “**a NOR b**” ?



$$\overline{a \text{ or } b} = \overline{a \text{ and } b}$$

Add Binary Numbers

- Same with decimal add

Carry out Carry in

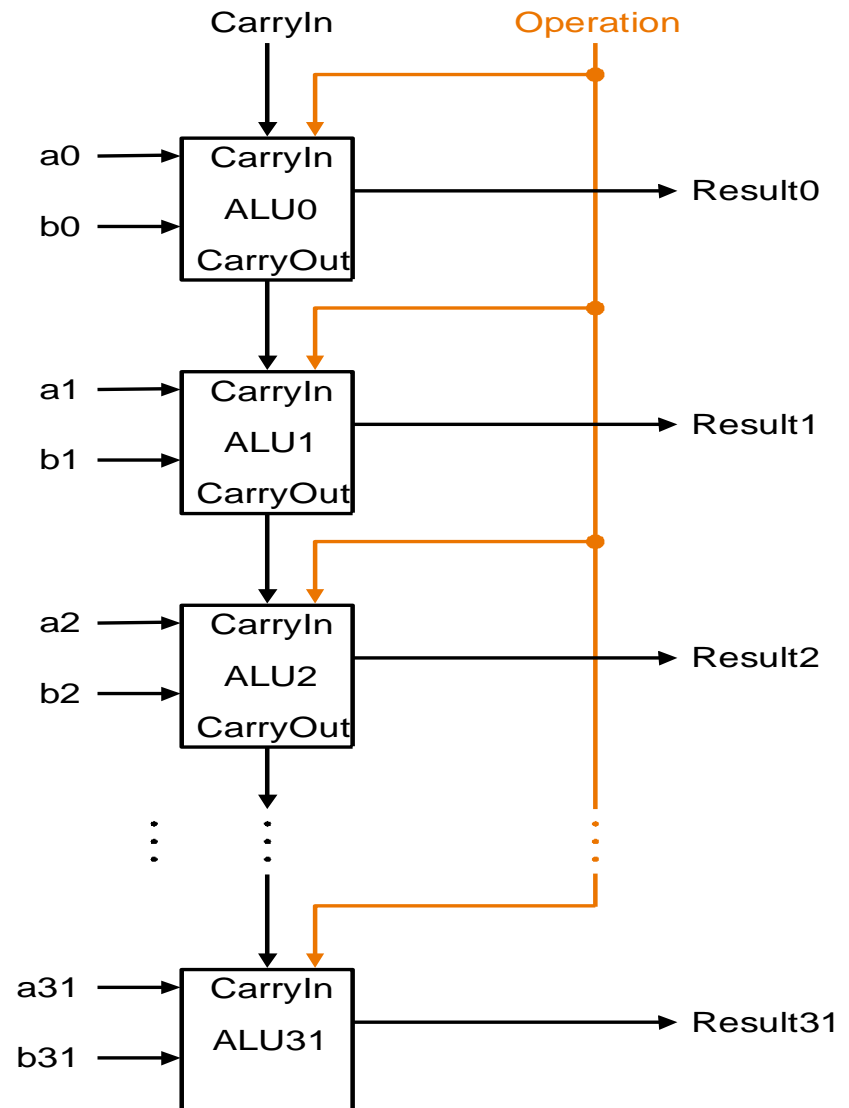


1 0 0 0 Carry

$9_{10} = 1 0 0 1$

$12_{10} = 1 1 0 0$


1 0 1 0 1 = 21_{10}

32-bit ALU Design (Abstraction)



What about subtraction ($a - b$) ?

- ❑ Two's complement approach: negate b and add
 - Negate: bitwise NOT, then add 1




and, or, +, -

Tailoring the ALU to the MIPS

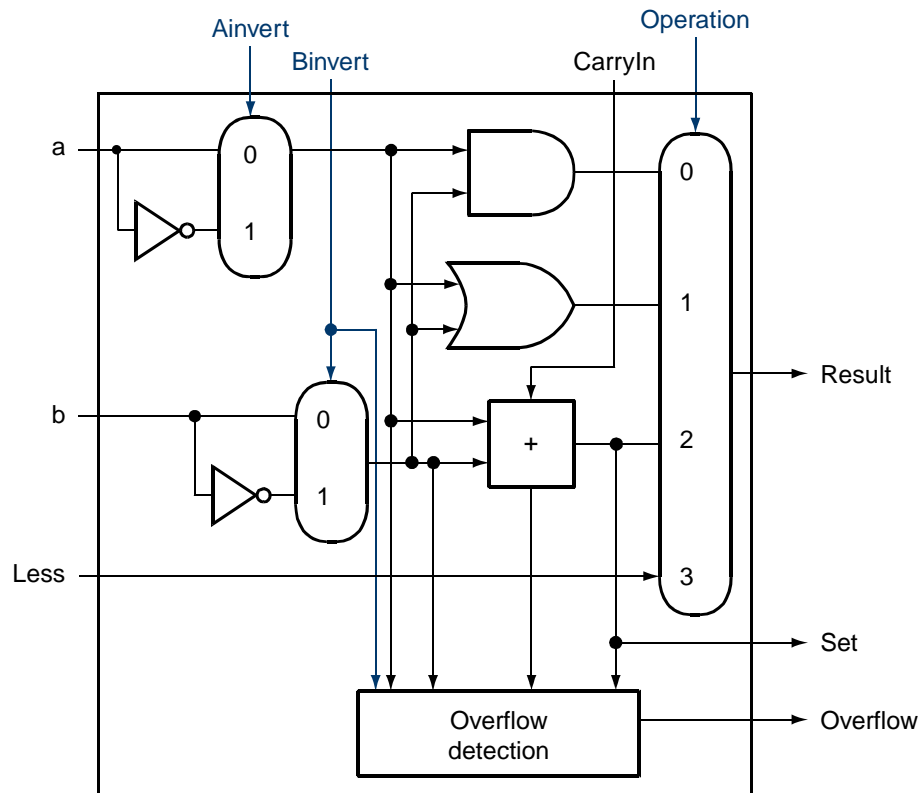
- ❑ Need to support the set-on-less-than instruction (**slt**)
 - Remember: **slt** is an arithmetic instruction
 - Produces a 1 if $rs < rt$ and 0 otherwise
 - Use subtraction: $(a-b) < 0$ implies $a < b$

slt \$t0, \$t1, \$t2

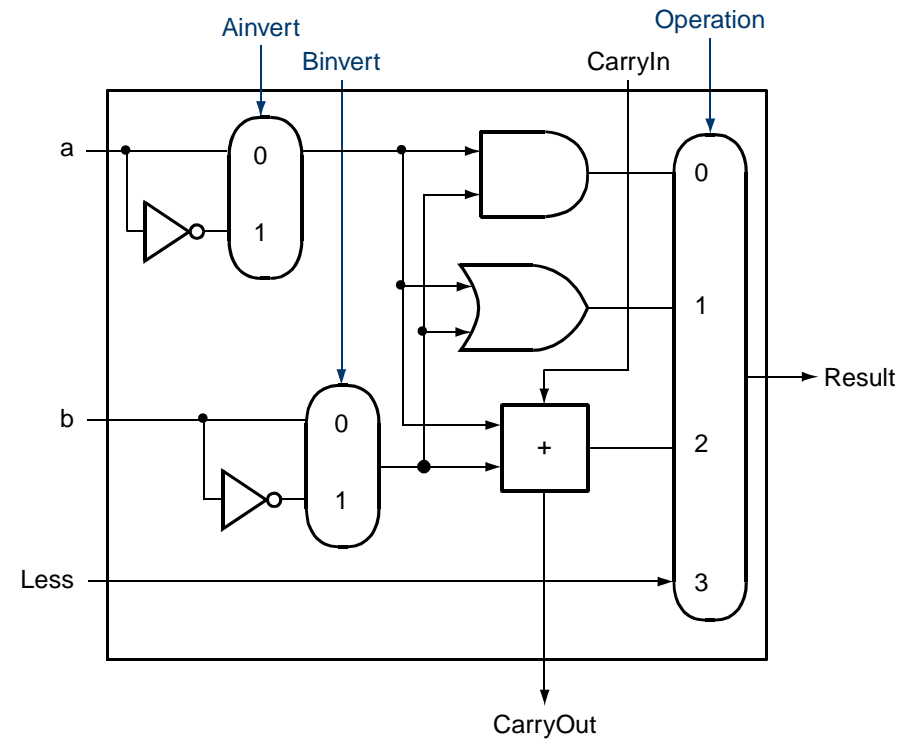


Supporting slt

❑ Can we figure out the idea?



Adding most significant bits



Adding all other bits

Tailoring the ALU to the MIPS

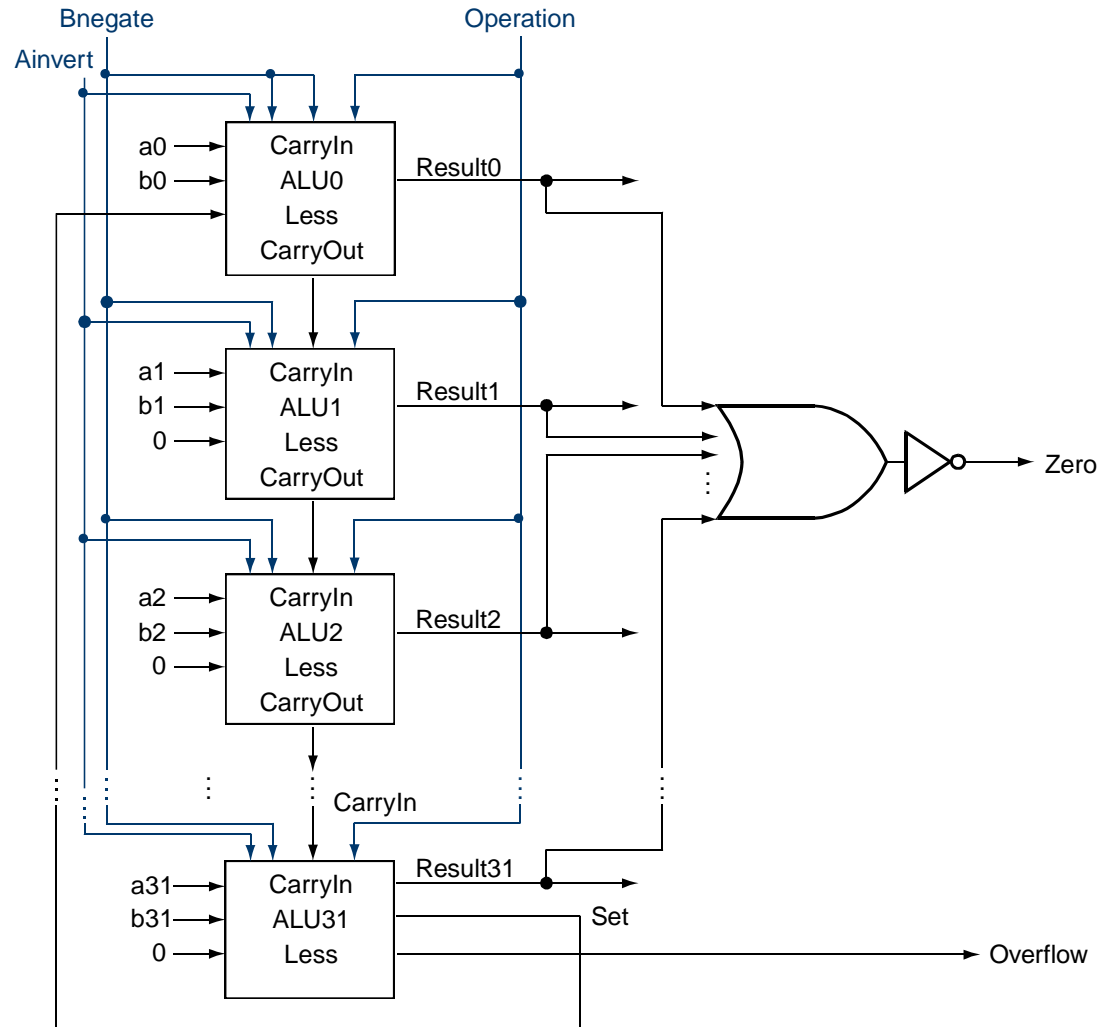
- ❑ Need to support test for equality (beq \$t5, \$t6, \$t7)
 - use subtraction: $(a-b) = 0$ implies $a = b$
가

Test for equality

□ Notice control lines:

0000 = and
0001 = or
0010 = add
0110 = subtract
0111 = slt
1100 = NOR

† *Note: zero is a 1 when
the result is zero!*



Conclusion

- ❑ We can build an ALU to support the MIPS instruction set
 - Key idea: use multiplexor to select the output we want
 - Can efficiently perform subtraction using 2's complement
 - Can replicate a 1-bit ALU to produce a 32-bit ALU
- ❑ Our primary focus: comprehension, however,
 - Clever changes to organization can improve performance (similar to using better algorithms in software)
 - Will see this in multiplication, let's look at addition now

Problem: ripple carry adder is slow

- ❑ Is a 32-bit ALU as fast as a 1-bit ALU?
- ❑ Is there more than one way to do addition?
 - Two extremes: ripple carry and sum-of-products

† Can you see the ripple? How could you get rid of it?

$$c_1 = b_0c_0 + a_0c_0 + a_0b_0$$

$$c_2 = b_1c_1 + a_1c_1 + a_1b_1 \quad c_2 =$$

$$c_3 = b_2c_2 + a_2c_2 + a_2b_2 \quad c_3 =$$

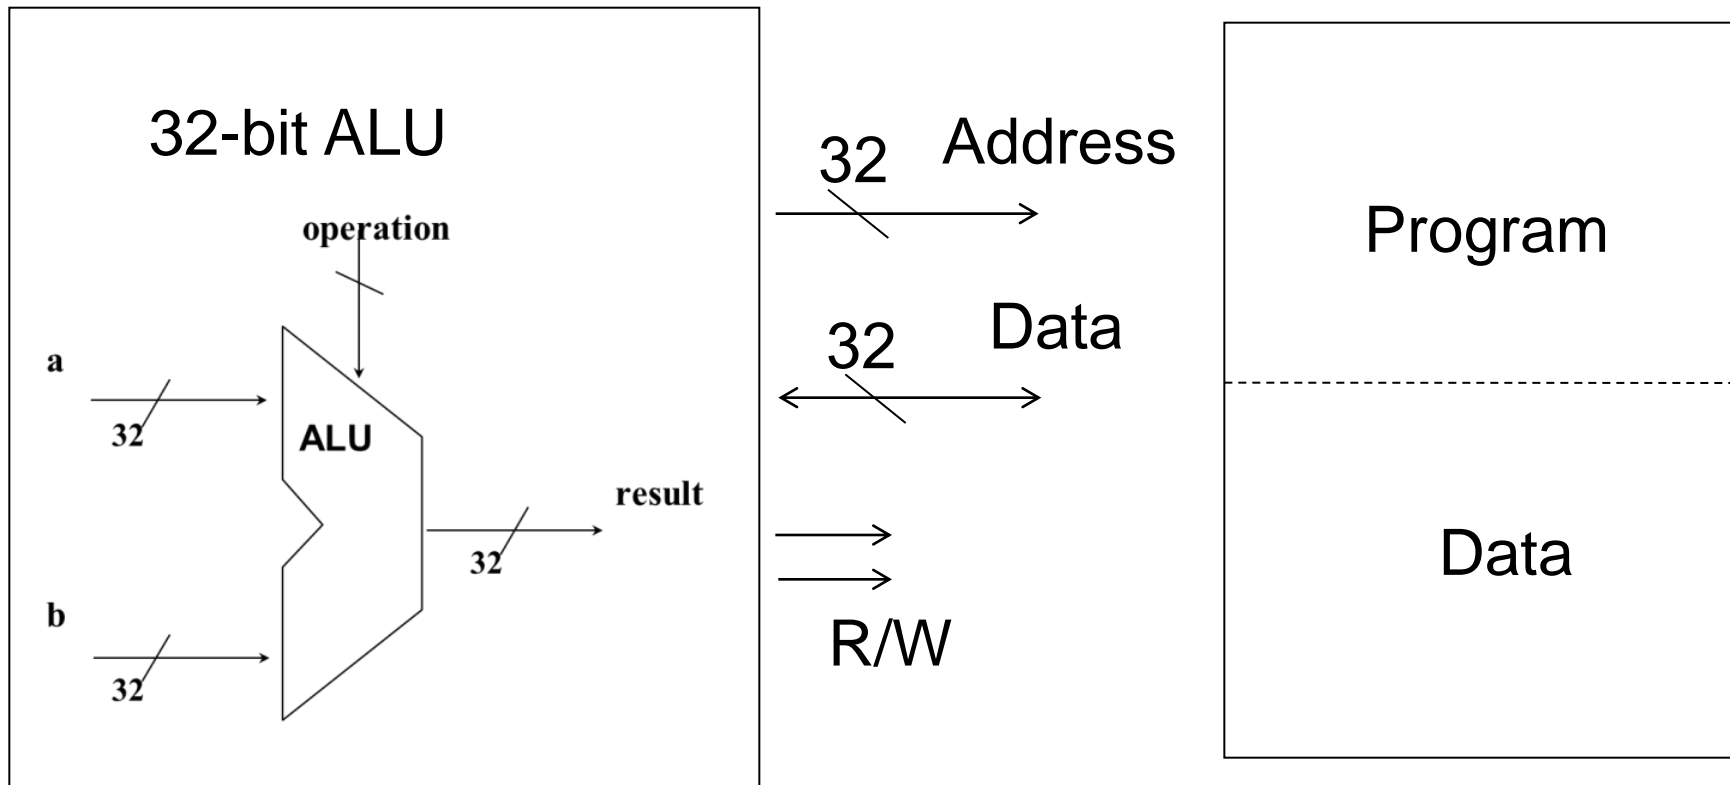
$$c_4 = b_3c_3 + a_3c_3 + a_3b_3 \quad c_4 =$$

Not feasible! Why?

32-bit Computer

CPU

Memory: 32-bit wide



I/O: Monitor/keyboard, LAN-Internet, ...

Computer Arithmetic and ALU

- ❑ Can you imagine
 - Many algorithms for addition,
 - Multiplication, division, FP operations as well
- ❑ Focus of building computer in 1945
 - Faster ALU
 - How to represent numbers
- ❑ Computer arithmetic
 - Matured in 1950s and 1960s
- ❑ Today, can buy ALU as IP (Intellectual Property)

Arithmetic for Multimedia

- ❑ Graphics and media processing operates on vectors of 8-bit and 16-bit data
 - Use 64-bit adder, with partitioned carry chain
 - Operate on 8×8-bit, 4×16-bit, or 2×32-bit vectors
 - SIMD (single-instruction, multiple-data), SSE
 - Subword parallelism
- ❑ Saturating operations
 - On overflow, result is largest representable value
 - c.f. 2s-complement modulo arithmetic
 - e.g., clipping in audio, saturation in video

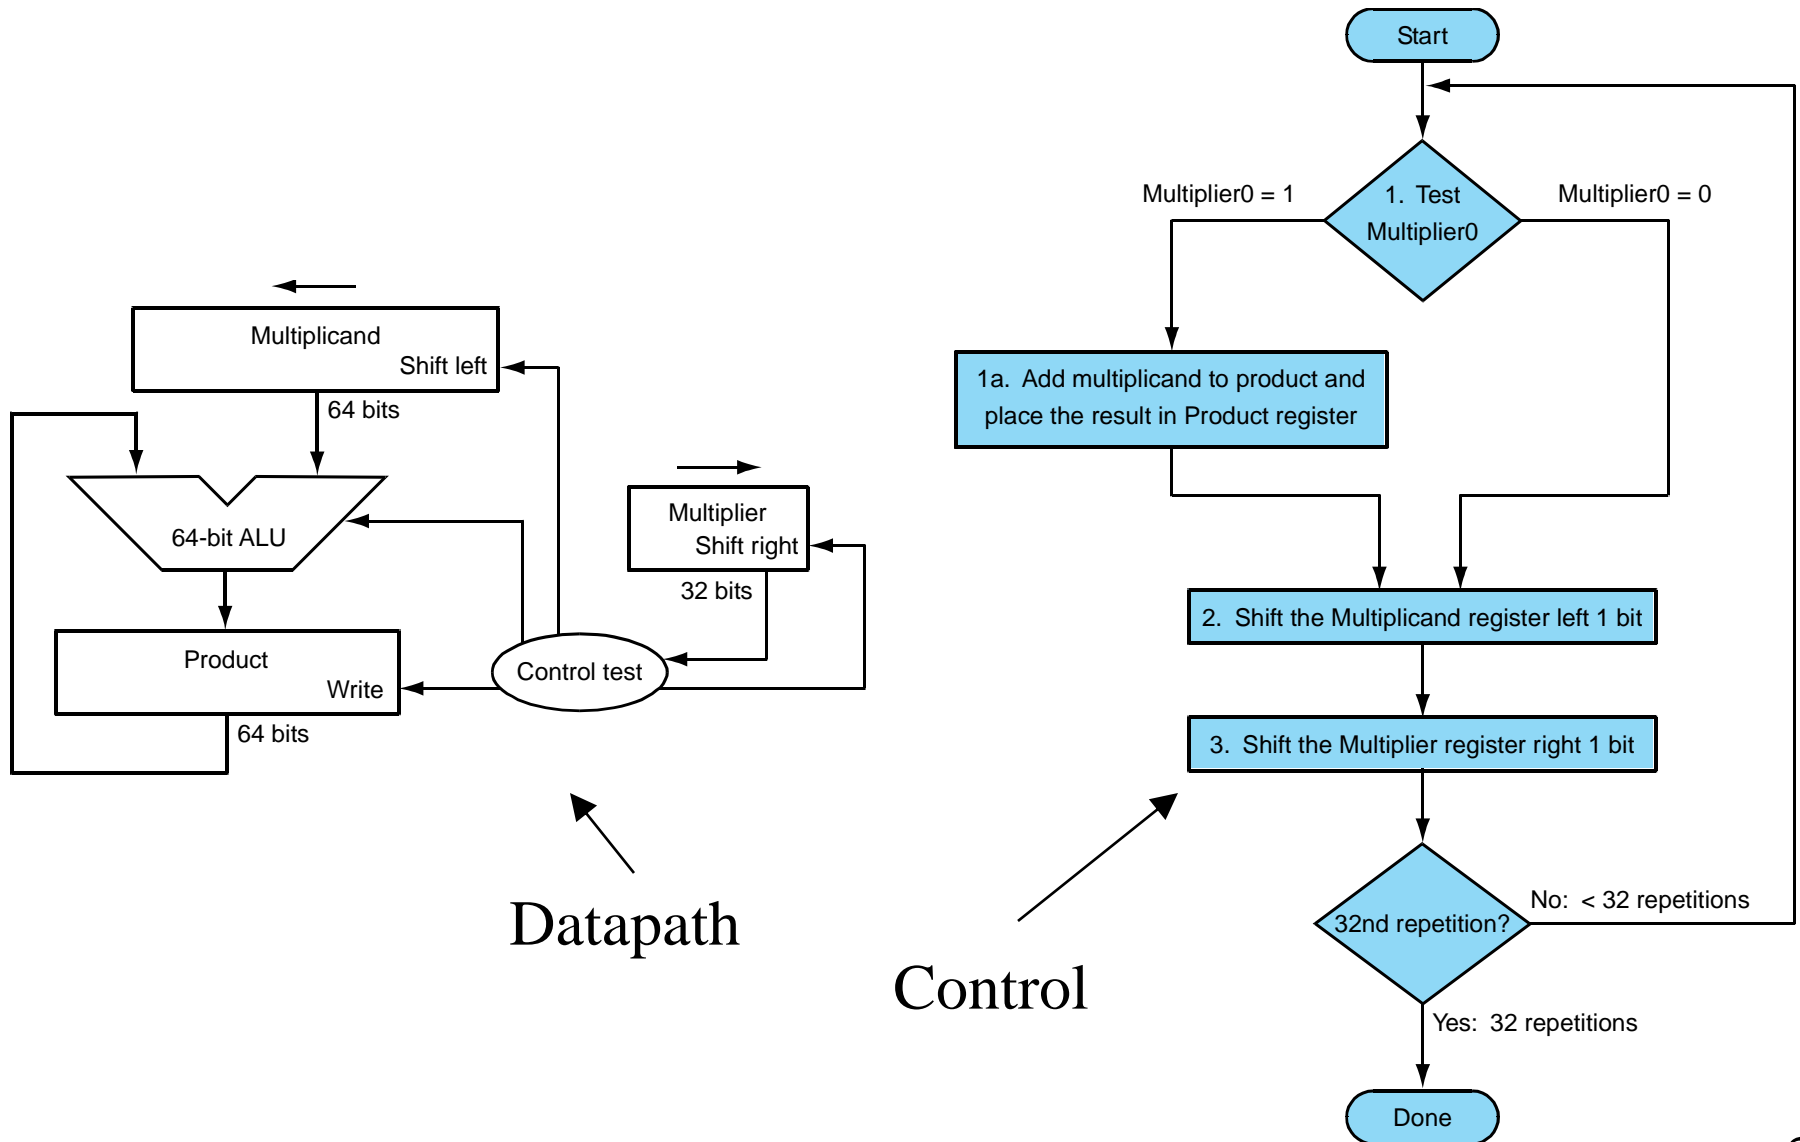
Multiplication

- ❑ More complicated than addition
 - Accomplished via shifting and addition
- ❑ More time and more area
- ❑ Let's look at 3 versions based on a gradeschool algorithm

$$\begin{array}{r} 0010 \quad (\text{multiplicand}) \\ \underline{\quad} \times \underline{1011} \quad (\text{multiplier}) \\ \hline \end{array}$$

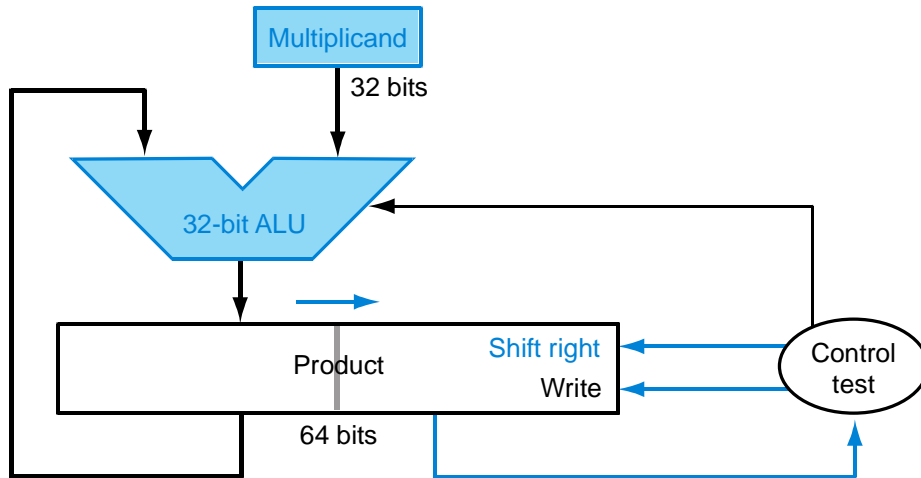
- ❑ Negative numbers: convert and multiply
 - There are better techniques, we won't look at them

Multiplication: Implementation

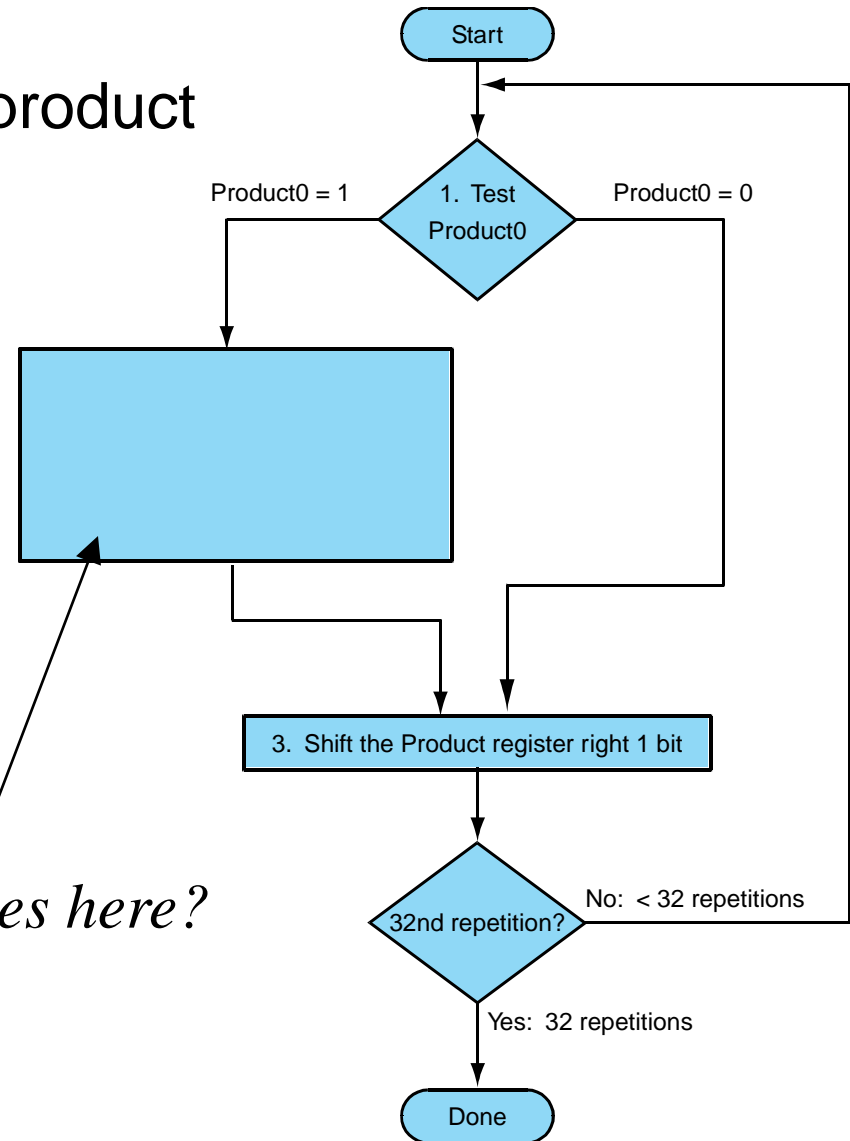


Final Version

- ❑ Multiplier starts in right half of product



What goes here?



MIPS Assembly Language: multiply, divide

| Category | Instruction | Example | Meaning | Comments |
|------------|--------------------------------|---------------------|--|-------------------------------------|
| Arithmetic | add | add \$s1,\$s2,\$s3 | $\$s1 = \$s2 + \$s3$ | Three operands; overflow detected |
| | subtract | sub \$s1,\$s2,\$s3 | $\$s1 = \$s2 - \$s3$ | Three operands; overflow detected |
| | add immediate | addi \$s1,\$s2,100 | $\$s1 = \$s2 + 100$ | + constant; overflow detected |
| | add unsigned | addu \$s1,\$s2,\$s3 | $\$s1 = \$s2 + \$s3$ | Three operands; overflow undetected |
| | subtract unsigned | subu \$s1,\$s2,\$s3 | $\$s1 = \$s2 - \$s3$ | Three operands; overflow undetected |
| | add immediate unsigned | addiu \$s1,\$s2,100 | $\$s1 = \$s2 + 100$ | + constant; overflow undetected |
| | move from coprocessor register | mfc0 \$s1,\$epc | $\$s1 = \epc | Copy Exception PC + special regs |
| | multiply | mult \$s2,\$s3 | Hi, Lo = $\$s2 \times \$s3$ | 64-bit signed product in Hi, Lo |
| | multiply unsigned | multu \$s2,\$s3 | Hi, Lo = $\$s2 \times \$s3$ | 64-bit unsigned product in Hi, Lo |
| | divide | div \$s2,\$s3 | Lo = $\$s2 / \$s3$, Hi = $\$s2 \bmod \$s3$ | Lo = quotient, Hi = remainder |
| | divide unsigned | divu \$s2,\$s3 | Lo = $\$s2 / \$s3$, Hi = $\$s2 \bmod \$s3$ | Unsigned quotient and remainder |
| | move from Hi | mfhi \$s1 | $\$s1 = \text{Hi}$ | Used to get copy of Hi |
| | move from Lo | mflo \$s1 | $\$s1 = \text{Lo}$ | Used to get copy of Lo |

Floating-Point Numbers

(Why the name? Fixed-point numbers?)

Floating Point (a brief look)

- ❑ We need a way to represent
 - numbers with fractions, e.g., 3.1416
 - very small numbers, e.g., 0.71×10^{-78}
 - very large numbers, e.g., 3.15576×10^{69}
- ❑ Representation:
 - sign, exponent, significand: $(-1)^{\text{sign}} \cdot \text{significand} \cdot 2^{\text{exponent}}$
 - more bits for significand gives more accuracy
 - more bits for exponent increases range
- ❑ IEEE 754 floating point standard:
 - single precision: 8 bit exponent, 23 bit significand
 - double precision: 11 bit exponent, 52 bit significand

IEEE 754 floating-point standard

- ❑ Leading “1” bit of significand is implicit
- ❑ Exponent is “biased” to make sorting easier
 - all 0s is smallest exponent all 1s is largest
 - bias of 127 for single precision and 1023 for double precision
 - summary: $(-1)^{\text{sign}} \cdot (1 + \text{significand}) \cdot 2^{\text{exponent} - \text{bias}}$
- ❑ Example:
 - decimal: $-.75 = - (\frac{1}{2} + \frac{1}{4})$
 - binary: $-.11 = -1.1 \times 2^{-1}$ normalized implicit 1 (bit)
 - floating point: exponent = 126 = 01111110
 - IEEE single precision:

10111111 01000000 00000000 00000000

IEEE 754 floating-point standard

| | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | |
|--------|----------|--------|--------|--------|--------|--------|--------|--------|----------|--------|--------|--------|--------|--------|--------|--------|--------|--------|--------|--------|--------|---|---|---|---|---|---|---|---|---|---|
| 3 1 | 3 0 | 2 9 | 2 8 | 2 7 | 2 6 | 2 5 | 2 4 | 2 3 | 2 2 | 2 1 | 2 0 | 1 9 | 1 8 | 1 7 | 1 6 | 1 5 | 1 4 | 1 3 | 1 2 | 1 1 | 1 0 | 9 | 8 | 7 | 6 | 5 | 4 | 3 | 2 | 1 | 0 |
| S | exponent | | | | | | | | fraction | | | | | | | | | | | | | | | | | | | | | | |

1 bit

8 bits

23 bits

| | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | |
|--------|----------|--------|--------|--------|--------|--------|--------|--------|--------|--------|--------|----------|--------|--------|--------|--------|--------|--------|--------|--------|--------|---|---|---|---|---|---|---|---|---|---|
| 3 1 | 3 0 | 2 9 | 2 8 | 2 7 | 2 6 | 2 5 | 2 4 | 2 3 | 2 2 | 2 1 | 2 0 | 1 9 | 1 8 | 1 7 | 1 6 | 1 5 | 1 4 | 1 3 | 1 2 | 1 1 | 1 0 | 9 | 8 | 7 | 6 | 5 | 4 | 3 | 2 | 1 | 0 |
| S | exponent | | | | | | | | | | | fraction | | | | | | | | | | | | | | | | | | | |

1 bit

11 bits

20 bits

| | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | |
|----------------------|--|--|--|--|--|--|--|--|--|--|--|--|--|--|--|--|--|--|--|--|--|--|--|--|--|--|--|--|--|--|--|
| fraction (continued) | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | |
|----------------------|--|--|--|--|--|--|--|--|--|--|--|--|--|--|--|--|--|--|--|--|--|--|--|--|--|--|--|--|--|--|--|

32 bits

Single-Precision Range

❑ Exponents 00000000 and 11111111 reserved

❑ Smallest value

- Exponent: 00000001

- $\Rightarrow \text{actual exponent} = 1 - 127 = -126$

- Fraction: 000...00 \Rightarrow significand = 1.0

- $\pm 1.0 \times 2^{-126} \approx \pm 1.2 \times 10^{-38}$

❑ Largest value

- exponent: 11111110

- $\Rightarrow \text{actual exponent} = 254 - 127 = +127$

- Fraction: 111...11 \Rightarrow significand ≈ 2.0

- $\pm 2.0 \times 2^{+127} \approx \pm 3.4 \times 10^{+38}$

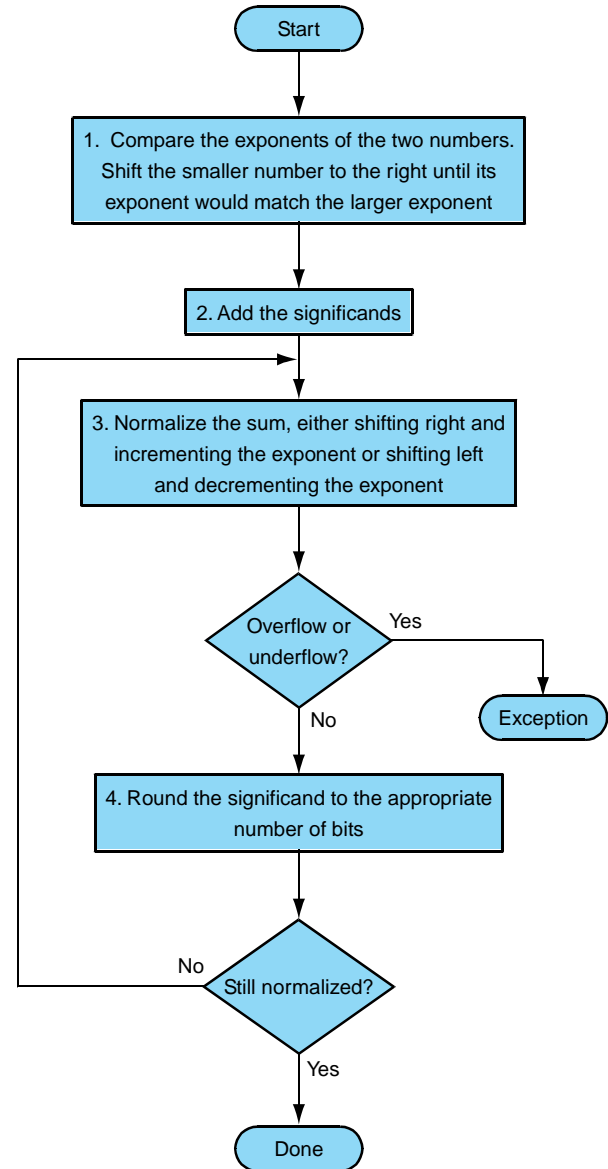
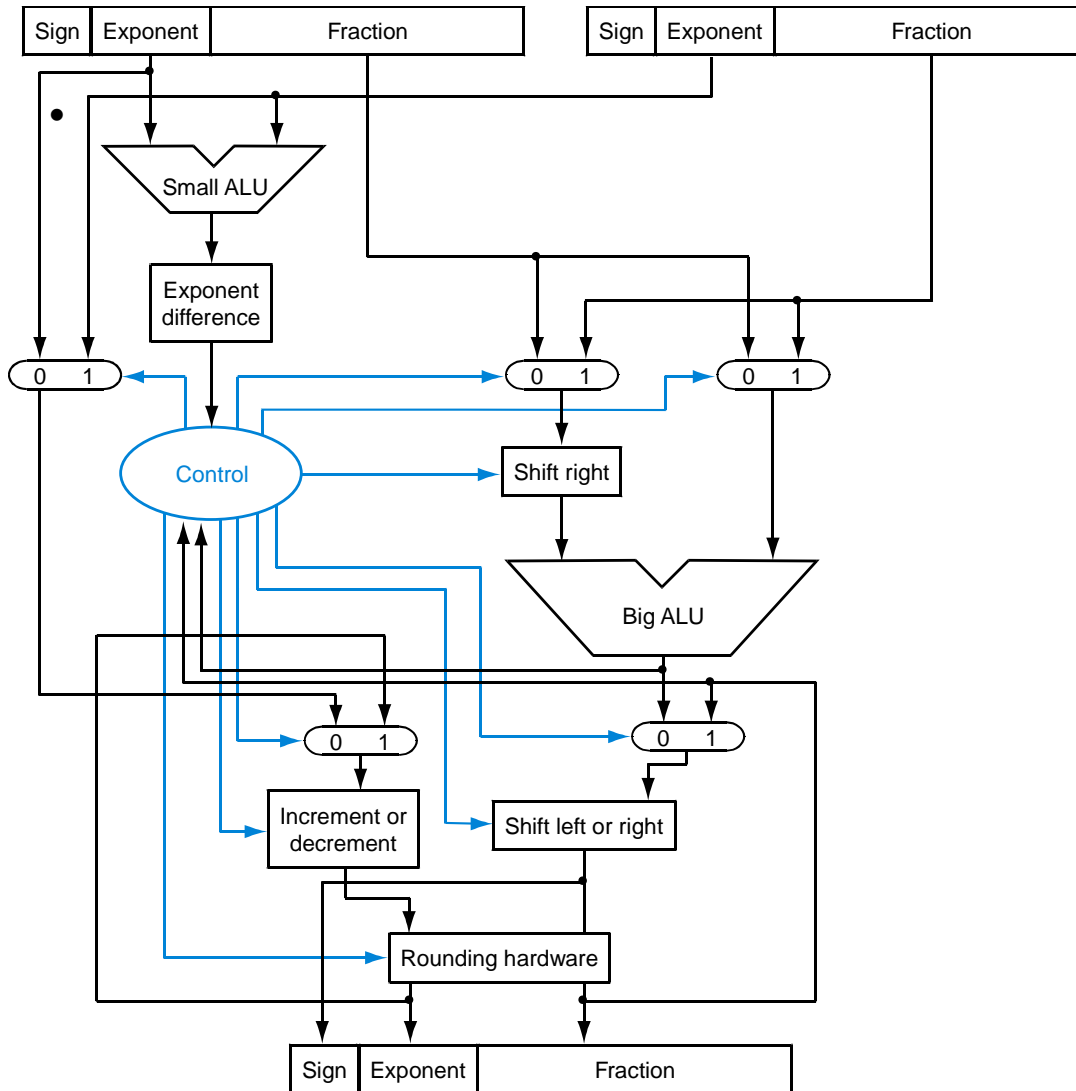
Double-Precision Range

- ❑ Exponents 0000...00 and 1111...11 reserved
- ❑ Smallest value
 - Exponent: 000000000001
 \Rightarrow actual exponent = $1 - 1023 = -1022$
 - Fraction: 000...00 \Rightarrow significand = 1.0
 - $\pm 1.0 \times 2^{-1022} \approx \pm 2.2 \times 10^{-308}$
- ❑ Largest value
 - Exponent: 111111111110
 \Rightarrow actual exponent = $2046 - 1023 = +1023$
 - Fraction: 111...11 \Rightarrow significand ≈ 2.0
 - $\pm 2.0 \times 2^{+1023} \approx \pm 1.8 \times 10^{+308}$

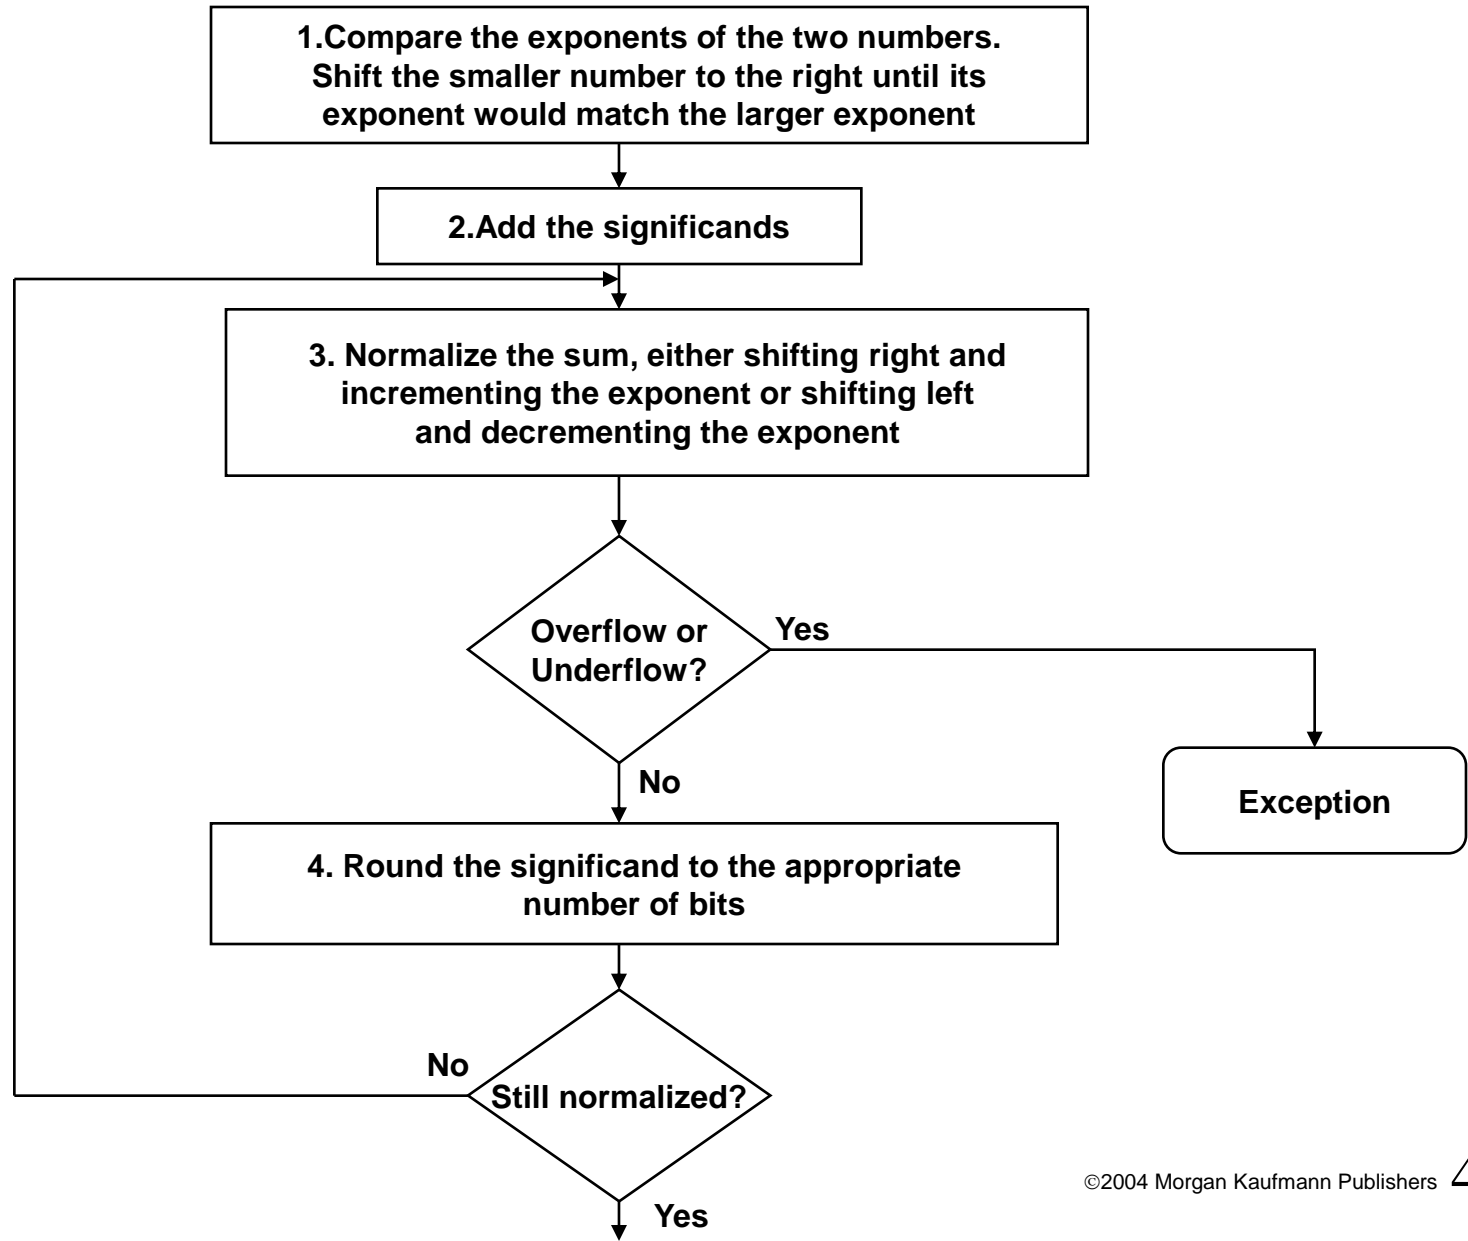
IEEE 754 floating point standard

- ❑ Why do we use floating point representation?
 - Very large numbers, very small numbers
- ❑ What do we lose?
 - Accuracy, intervals not uniform
- ❑ Think about overflow, underflow
- ❑ In 32-bit standard, how many bits for significand?
 - Range vs. accuracy trade-off
- ❑ How do we represent 0, infinity?
 - $\text{Exp} = 255, \text{sig} \neq 0 \quad \rightarrow \text{NaN}$
 - $\text{Exp} = 255, \text{sig} = 0 \quad \rightarrow +, - \text{infinity}$
 - $\text{Exp} = 0, \text{sig} = 0 \quad \rightarrow 0$
 - $\text{Exp} = 0, \text{sig} \neq 0 \quad \rightarrow \text{even smaller numbers}$

Floating point addition



Floating point addition




FP Adder Hardware

- ❑ Much more complex than integer adder
- ❑ Doing it in one clock cycle would take too long
 - Much longer than integer operations
 - Slower clock would penalize all instructions
- ❑ FP adder usually takes several cycles
 - Can be pipelined

Integer and FP ALUs (반복)

- ❑ Processor for general-purpose computers:



- ❑ Processors for embedded systems:



MIPS Assembly Language: floating-point

integer

floating-point register

- single

- double

MIPS floating-point operands

| Name | Example | Comments |
|------------------------------|--|---|
| 32 floating-point registers | \$f0, \$f1, \$f2, . . . , \$f31 | MIPS floating-point registers are used in pairs for double precision numbers. |
| 2 ³⁰ memory words | Memory[0], Memory[4], . . . , Memory[4294967292] | Accessed only by data transfer instructions. MIPS uses byte addresses, so sequential word addresses differ by 4. Memory holds data structures, such as arrays, and spilled registers, such as those saved on procedure calls. |

MIPS floating-point assembly language

| Category | Instruction | Example | Meaning | Comments |
|--------------------|---------------------------------------|----------------------|---|---------------------------------------|
| Arithmetic | FP add single | add.s \$f2,\$f4,\$f6 | \$f2 = \$f4 + \$f6 | FP add (single precision) |
| | FP subtract single | sub.s \$f2,\$f4,\$f6 | \$f2 = \$f4 - \$f6 | FP sub (single precision) |
| | FP multiply single | mul.s \$f2,\$f4,\$f6 | \$f2 = \$f4 × \$f6 | FP multiply (single precision) |
| | FP divide single | div.s \$f2,\$f4,\$f6 | \$f2 = \$f4 / \$f6 | FP divide (single precision) |
| | FP add double | add.d \$f2,\$f4,\$f6 | \$f2 = \$f4 + \$f6 | FP add (double precision) |
| | FP subtract double | sub.d \$f2,\$f4,\$f6 | \$f2 = \$f4 - \$f6 | FP sub (double precision) |
| | FP multiply double | mul.d \$f2,\$f4,\$f6 | \$f2 = \$f4 × \$f6 | FP multiply (double precision) |
| | FP divide double | div.d \$f2,\$f4,\$f6 | \$f2 = \$f4 / \$f6 | FP divide (double precision) |
| Data transfer | load word copr. 1 | lwc1 \$f1,100(\$s2) | \$f1 = Memory[\$s2 + 100] | 32-bit data to FP register |
| | store word copr. 1 | swc1 \$f1,100(\$s2) | Memory[\$s2 + 100] = \$f1 | 32-bit data to memory |
| Conditional branch | branch on FP true | bclt 25 | if (cond == 1) go to PC + 4 + 100 | PC-relative branch if FP cond. |
| | branch on FP false | bclf 25 | if (cond == 0) go to PC + 4 + 100 | PC-relative branch if not cond. |
| | FP compare single (eq,ne,lt,le,gt,ge) | c.lt.s \$f2,\$f4 | if (\$f2 < \$f4) cond = 1; else cond = 0 | FP compare less than single precision |
| | FP compare double (eq,ne,lt,le,gt,ge) | c.lt.d \$f2,\$f4 | if (\$f2 < \$f4) cond = 1; else cond = 0 | FP compare less than double precision |

MIPS Assembly Language: floating-point

MIPS floating-point machine language

| Name | Format | Example | | | | | | Comments |
|------------|--------|---------|--------|--------|--------|--------|--------|-------------------------------|
| add.s | R | 17 | 16 | 6 | 4 | 2 | 0 | add.s \$f2,\$f4,\$f6 |
| sub.s | R | 17 | 16 | 6 | 4 | 2 | 1 | sub.s \$f2,\$f4,\$f6 |
| mul.s | R | 17 | 16 | 6 | 4 | 2 | 2 | mul.s \$f2,\$f4,\$f6 |
| div.s | R | 17 | 16 | 6 | 4 | 2 | 3 | div.s \$f2,\$f4,\$f6 |
| add.d | R | 17 | 17 | 6 | 4 | 2 | 0 | add.d \$f2,\$f4,\$f6 |
| sub.d | R | 17 | 17 | 6 | 4 | 2 | 1 | sub.d \$f2,\$f4,\$f6 |
| mul.d | R | 17 | 17 | 6 | 4 | 2 | 2 | mul.d \$f2,\$f4,\$f6 |
| div.d | R | 17 | 17 | 6 | 4 | 2 | 3 | div.d \$f2,\$f4,\$f6 |
| lwc1 | I | 49 | 20 | 2 | 100 | | | lwc1 \$f2,100(\$s4) |
| swc1 | I | 57 | 20 | 2 | 100 | | | swc1 \$f2,100(\$s4) |
| bc1t | I | 17 | 8 | 1 | 25 | | | bc1t 25 |
| bc1f | I | 17 | 8 | 0 | 25 | | | bc1f 25 |
| c.lt.s | R | 17 | 16 | 4 | 2 | 0 | 60 | c.lt.s \$f2,\$f4 |
| c.lt.d | R | 17 | 17 | 4 | 2 | 0 | 60 | c.lt.d \$f2,\$f4 |
| Field size | | 6 bits | 5 bits | 5 bits | 5 bits | 5 bits | 6 bits | All MIPS instructions 32 bits |

To Think about

double

.

☐ When you do C, Java programming

- When do you use **float**? imbedded system
- When do you use **double**? double 가 powerful 가
- Do you ever think about long integer or arithmetic overflow?

☐ FP numbers and FP ALU

- 계산을 위한 인공적인 고안의 결과
- No beauty of integers
 - Not accurate from mathematics perspective

Associativity

- ❑ Parallel programs may interleave operations in unexpected orders
 - Assumptions of associativity may fail

| | | $(x+y)+z$ | $x+(y+z)$ |
|---|-----------|-----------|-----------|
| x | -1.50E+38 | | -1.50E+38 |
| y | 1.50E+38 | 0.00E+00 | |
| z | 1.0 | 1.0 | 1.50E+38 |
| | | 1.00E+00 | 0.00E+00 |

- ❑ Need to validate parallel programs under varying degrees of parallelism – numerical analysis
- ❖ FP numbers only approximations of real numbers

Who Cares About FP Accuracy?

- ❑ Important for scientific code
 - But for everyday consumer use?
 - “My bank balance is out by 0.0002¢!” ☹
- ❑ The Intel Pentium FDIV bug
 - The market expects accuracy
 - See Colwell, *The Pentium Chronicles*

Sections 3.6

Parallelism and Computer Arithmetic: Subword Parallelism

Sections 3.7

**Real stuff: streaming
SIMD extensions and
advanced vector
extensions in x86**

Sections 3.8

**Going faster: subword
parallelism and matrix
multiply**

Sections 3.10

Concluding Remarks

(가볍게 읽어 보실 것)

Concluding Remarks

- Bits have no inherent meaning
 - Interpretation depends on the instructions applied
- Computer representations of numbers
 - Finite range and precision
 - Need to account for this in programs

Concluding Remarks

- ISAs support arithmetic
 - Signed and unsigned integers
 - Floating-point approximation to reals
- Bounded range and precision
 - Operations can overflow and underflow
- MIPS ISA
 - Core instructions: 54 most frequently used
 - 100% of SPECINT, 97% of SPECFP
 - Other instructions: less frequent

MIPS Core, MIPS-32, Pseudo MIPS

| MIPS core instructions | Name | Format | MIPS arithmetic core | Name | Format |
|--------------------------------------|-------|--------|-------------------------------------|-------|--------|
| add | add | R | multiply | mult | R |
| add immediate | addi | I | multiply unsigned | multu | R |
| add unsigned | addu | R | divide | div | R |
| add immediate unsigned | addiu | I | divide unsigned | divu | R |
| subtract | sub | R | move from Hi | mfhi | R |
| subtract unsigned | subu | R | move from Lo | mflo | R |
| AND | AND | R | move from system control (EPC) | mfc0 | R |
| AND immediate | ANDi | I | floating-point add single | add.s | R |
| OR | OR | R | floating-point add double | add.d | R |
| OR immediate | ORi | I | floating-point subtract single | sub.s | R |
| NOR | NOR | R | floating-point subtract double | sub.d | R |
| shift left logical | sll | R | floating-point multiply single | mul.s | R |
| shift right logical | srl | R | floating-point multiply double | mul.d | R |
| load upper immediate | lui | I | floating-point divide single | div.s | R |
| load word | lw | I | floating-point divide double | div.d | R |
| store word | sw | I | load word to floating-point single | lwc1 | I |
| load halfword unsigned | lhu | I | store word to floating-point single | swc1 | I |
| store halfword | sh | I | load word to floating-point double | ldc1 | I |
| load byte unsigned | lbu | I | store word to floating-point double | sdcl | I |
| store byte | sb | I | branch on floating-point true | bclt | I |
| load linked (<i>atomic update</i>) | ll | I | branch on floating-point false | bclf | I |
| store cond. (<i>atomic update</i>) | sc | I | floating-point compare single | c.x.s | R |
| branch on equal | beq | I | (x = eq, neq, lt, le, gt, ge) | | |
| branch on not equal | bne | I | floating-point compare double | c.x.d | R |
| jump | j | J | (x = eq, neq, lt, le, gt, ge) | | |
| jump and link | jal | J | | | |
| jump register | jr | R | | | |
| set less than | slt | R | | | |
| set less than immediate | slti | I | | | |
| set less than unsigned | sltu | R | | | |
| set less than immediate unsigned | sltiu | I | | | |

Figure 3.26

MIPS

IPS

| Remaining MIPS-32 | Name | Format | Pseudo MIPS | Name | Format |
|--|---------|--------|---|--------|----------|
| exclusive or ($rs \oplus rt$) | xor | R | absolute value | abs | rd,rs |
| exclusive or immediate | xori | I | negate (<i>signed or unsigned</i>) | negs | rd,rs |
| shift right arithmetic | sra | R | rotate left | rol | rd,rs,rt |
| shift left logical variable | sllv | R | rotate right | ror | rd,rs,rt |
| shift right logical variable | srlv | R | multiply and don't check oflw (<i>signed or uns.</i>) | mults | rd,rs,rt |
| shift right arithmetic variable | srav | R | multiply and check oflw (<i>signed or uns.</i>) | multos | rd,rs,rt |
| move to Hi | mthi | R | divide and check overflow | div | rd,rs,rt |
| move to Lo | mtlo | R | divide and don't check overflow | divu | rd,rs,rt |
| load halfword | lh | I | remainder (<i>signed or unsigned</i>) | rems | rd,rs,rt |
| load byte | lb | I | load immediate | li | rd,imm |
| load word left (<i>unaligned</i>) | lwl | I | load address | la | rd,addr |
| load word right (<i>unaligned</i>) | lwr | I | load double | ld | rd,addr |
| store word left (<i>unaligned</i>) | swl | I | store double | sd | rd,addr |
| store word right (<i>unaligned</i>) | swr | I | unaligned load word | ulw | rd,addr |
| load linked (<i>atomic update</i>) | ll | I | unaligned store word | usw | rd,addr |
| store cond. (<i>atomic update</i>) | sc | I | unaligned load halfword (<i>signed or uns.</i>) | ulhs | rd,addr |
| move if zero | movz | R | unaligned store halfword | ush | rd,addr |
| move if not zero | movn | R | branch | b | Label |
| multiply and add (<i>S or uns.</i>) | madds | R | branch on equal zero | beqz | rs,L |
| multiply and subtract (<i>S or uns.</i>) | msubs | I | branch on compare (<i>signed or unsigned</i>) | bxs | rs,rt,L |
| branch on \geq zero and link | bgezal | I | ($x = lt, le, gt, ge$) | | |
| branch on $<$ zero and link | bltzal | I | set equal | seq | rd,rs,rt |
| jump and link register | jlr | R | set not equal | sne | rd,rs,rt |
| branch compare to zero | bxz | I | set on compare (<i>signed or unsigned</i>) | sxs | rd,rs,rt |
| branch compare to zero likely | bxzl | I | ($x = lt, le, gt, ge$) | | |
| ($x = lt, le, gt, ge$) | | | load to floating point (<i>s or d</i>) | l.f | rd,addr |
| branch compare reg likely | bxl | I | store from floating point (<i>s or d</i>) | s.f | rd,addr |
| trap if compare reg | tx | R | | | |
| trap if compare immediate | txi | I | | | |
| ($x = eq, neq, lt, le, gt, ge$) | | | | | |
| return from exception | rfe | R | | | |
| system call | syscall | I | | | |
| break (<i>cause exception</i>) | break | I | | | |
| move from FP to integer | mfc1 | R | | | |
| move to FP from integer | mtc1 | R | | | |
| FP move (<i>s or d</i>) | mov.f | R | | | |
| FP move if zero (<i>s or d</i>) | movz.f | R | | | |
| FP move if not zero (<i>s or d</i>) | movn.f | R | | | |
| FP square root (<i>s or d</i>) | sqr.f | R | | | |
| FP absolute value (<i>s or d</i>) | abs.f | R | | | |
| FP negate (<i>s or d</i>) | neg.f | R | | | |
| FP convert (<i>w, s, or d</i>) | cvt.f | R | | | |
| FP compare un (<i>s or d</i>) | c.xn.f | R | | | |

Figure 3.27

MIPS Core, MIPS-32, Pseudo MIPS

- Frequency of use in SPEC CPU2006 benchmark

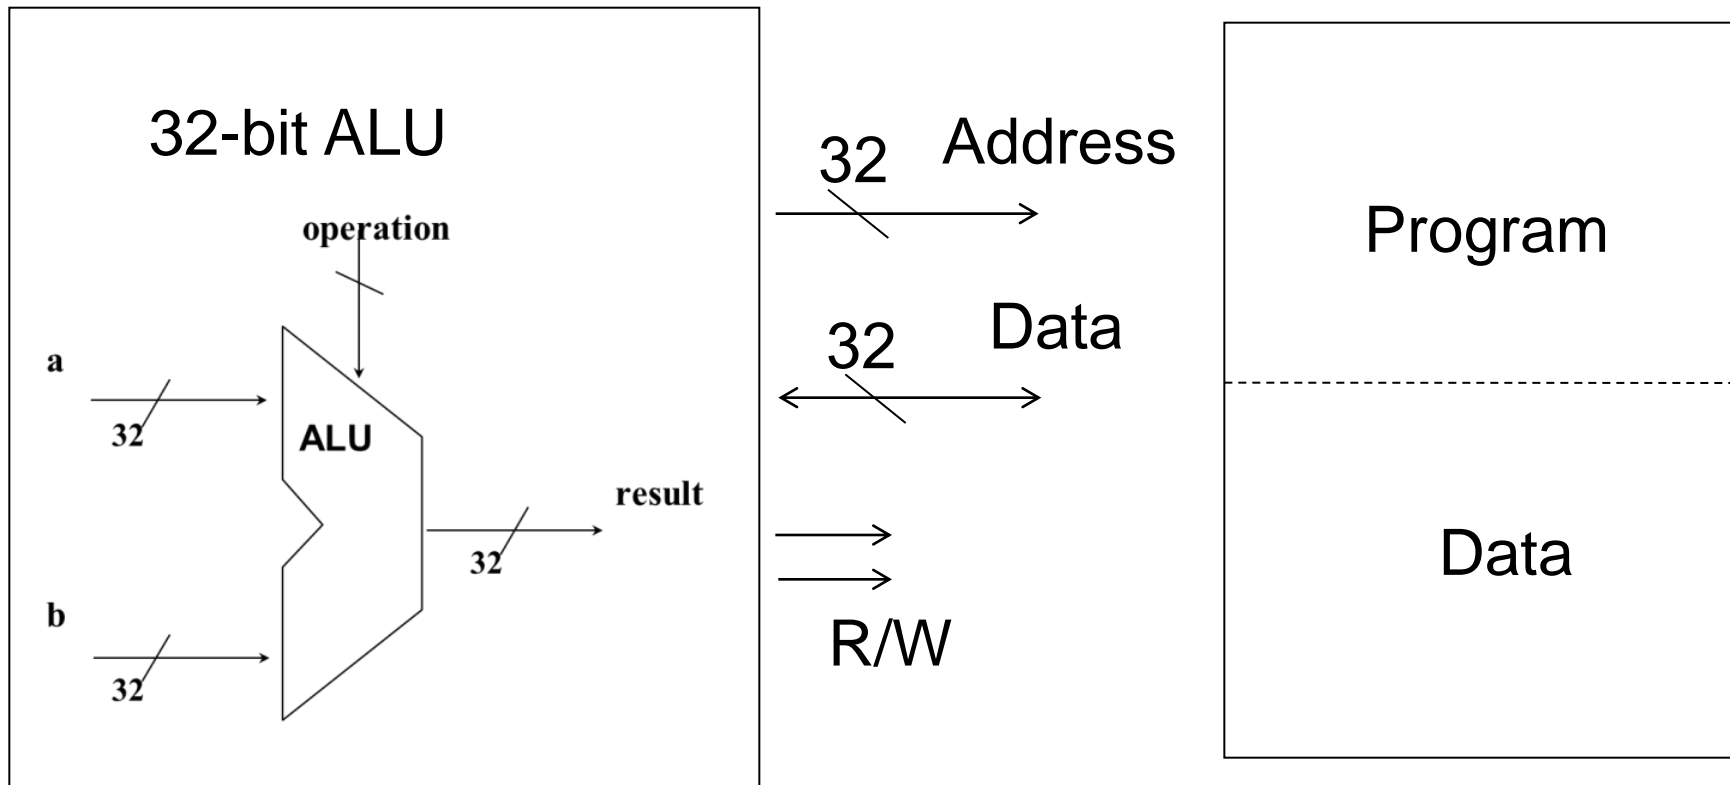
| Instruction subset | Integer | Fl. pt. |
|----------------------|---------|---------|
| MIPS core | 98% | 31% |
| MIPS arithmetic core | 2% | 66% |
| Remaining MIPS-32 | 0% | 3% |

- For the rest of book, focus on MIPS core (integer instruction set excluding multiply and divide)
 - To make the explanation of computer design easier

32-bit Computer

CPU

Memory: 32-bit wide



I/O: Monitor/keyboard, LAN-Internet, ...

Will implement fetch-decode-execute in Chapter 4