

Sections 3.6 – 3.8

**Subword Parallelism,
MMX, SSE, AVX in x86,
Going Faster**

Subword Parallelism

- History of support for small integers
 - Bytes, half words
 - Only load and store instructions
- Popularity of multimedia applications
 - 8-bit color, 16-bit audio
 - Arithmetic instructions to support narrower operations
 - Operate in parallel
 - Wide registers and multiple functional units

Subword Parallelism

- Graphics and audio applications can take advantage of performing simultaneous operations on short vectors
 - Example: 128-bit adder:
 - Sixteen 8-bit adds
 - Eight 16-bit adds
 - Four 32-bit adds
- Also called data-level parallelism, vector parallelism, or Single Instruction Multiple Data (SIMD)

The Intel x86 ISA (참고자료)

- Evolution with backward compatibility
 - 8080 (1974): 8-bit microprocessor
 - Accumulator, plus 3 index-register pairs
 - 8086 (1978): 16-bit extension to 8080
 - Complex instruction set (CISC)
 - 8087 (1980): floating-point coprocessor
 - Adds FP instructions and register stack
 - 80286 (1982): 24-bit addresses, MMU
 - Segmented memory mapping and protection
 - 80386 (1985): 32-bit extension (now IA-32)
 - Additional addressing modes and operations
 - Paged memory mapping as well as segments

The Intel x86 ISA (참고자료)

- Further evolution...
 - i486 (1989): pipelined, on-chip caches and FPU
 - Compatible competitors: AMD, Cyrix, ...
 - Pentium (1993): superscalar, 64-bit datapath
 - Later versions added 57 MMX (Multi-Media eXtension) instructions (1997)
 - The infamous FDIV bug
 - Pentium Pro (1995), Pentium II (1997)
 - New microarchitecture (see Colwell, *The Pentium Chronicles*)
 - Pentium III (1999)
 - Added 70 SSE (Streaming SIMD Extensions) instructions and associated registers
 - Cache prefetch instructions, streaming store instructions
 - Pentium 4 (2001)
 - New microarchitecture
 - Added 144 SSE2 instructions

The Intel x86 ISA (참고자료)

- And further...
 - AMD64 (2003): extended architecture to 64 bits
 - EM64T – Extended Memory 64 Technology (2004)
 - AMD64 adopted by Intel (with refinements)
 - Added 13 SSE3 instructions
 - Intel Core (2006)
 - Added SSE4 instructions, virtual machine support
 - AMD64 (announced 2007): SSE5 instructions
 - Intel declined to follow, instead...
 - AVX: Advanced Vector Extension (announced 2008)
 - Longer SSE registers, more instructions
- If Intel didn't extend with compatibility, its competitors would!
 - Technical elegance ≠ market success

MMX and SSE

■ MMX

- Use (64 bits of) existing FP registers
- Small integer data types (8bit color, 16bit audio)
 - Packed byte (8×8), packed word (4×16), (2×32), (1×64)

■ SSE

- Add new 8×128 -bit FP registers (XMM0-XMM7)
- Packed single precision FP (4×32)

SSE2, ..., AVX

- SSE2
 - Adds 8×128 -bit registers
 - Generalize SSE
 - Can be used for multiple FP/INT operands
 - 2×64 -bit FP (`addpd %xmm0, %xmm4`)
 - 4×32 -bit FP (`addps %xmm0, %xmm4`)
 - 2×64 -bit, 4×32 -bit, 8×16 -bit, 16×8 -bit INT
- AVX (advanced vector extension)
 - Double the width of registers
 - e.g., 4×64 -bit double precision FP

SSE/SSE2 (skip)

Data transfer	Arithmetic	Compare
MOV{A/U}{SS/PS/SD/PD} xmm, mem/xmm	ADD{SS/PS/SD/PD} xmm, mem/xmm	CMP{SS/PS/SD/PD}
	SUB{SS/PS/SD/PD} xmm, mem/xmm	
MOV {H/L} {PS/PD} xmm, mem/xmm	MUL{SS/PS/SD/PD} xmm, mem/xmm	
	DIV{SS/PS/SD/PD} xmm, mem/xmm	
	SQRT{SS/PS/SD/PD} mem/xmm	
	MAX {SS/PS/SD/PD} mem/xmm	
	MIN{SS/PS/SD/PD} mem/xmm	

Going Faster

- Demonstrate the performance benefit of adapting software to the underlying hardware

Going Faster

- With and without AVX, Intel Core i7
- DGEMM (Double Precision General Matrix Multiply)
 - AVX version 3.85 times fast
 - 4 double precision FP operations in parallel
- With Turbo mode turned off
 - Turbo mode
 - Temporarily run at higher clock until chip gets too hot
 - Particularly useful when using only single core of multicore chip

Matrix Multiply, DGEMM (참고자료)

■ Unoptimized code:

```
1. void dgemm (int n, double* A, double* B, double* C)
2. {
3.   for (int i = 0; i < n; ++i)
4.     for (int j = 0; j < n; ++j)
5.       {
6.         double cij = C[i+j*n]; /* cij = C[i][j] */
7.         for(int k = 0; k < n; k++ )
8.           cij += A[i+k*n] * B[k+j*n]; /* cij += A[i][k]*B[k][j] */
9.         C[i+j*n] = cij; /* C[i][j] = cij */
10.      }
11. }
```

Matrix Multiply, DGEMM (참고자료)

■ x86 assembly code:

```

1. vmovsd (%r10),%xmm0    # Load 1 element of C into %xmm0
2. mov %rsi,%rcx           # register %rcx = %rsi
3. xor %eax,%eax           # register %eax = 0
4. vmovsd (%rcx),%xmm1    # Load 1 element of B into %xmm1
5. add %r9,%rcx            # register %rcx = %rcx + %r9
6. vmulsd (%r8,%rax,8),%xmm1,%xmm1 # Multiply %xmm1,
   element of A
7. add $0x1,%rax           # register %rax = %rax + 1
8. cmp %eax,%edi           # compare %eax to %edi
9. vaddsd %xmm1,%xmm0,%xmm0 # Add %xmm1, %xmm0
10. jg 30 <dgemm+0x30>     # jump if %eax > %edi
11. add $0x1,%r11d         # register %r11 = %r11 + 1
12. vmovsd %xmm0, (%r10)  # Store %xmm0 into C element

```

Matrix Multiply, DGEMM (참고자료)

■ Optimized C code:

```

1. #include <x86intrin.h>
2. void dgemm (int n, double* A, double* B, double* C)
3. {
4.     for ( int i = 0; i < n; i+=4 )
5.         for ( int j = 0; j < n; j++ ) {
6.             __m256d c0 = _mm256_load_pd(C+i+j*n); /* c0 = C[i][j]
               */
7.             for( int k = 0; k < n; k++ )
8.                 c0 = _mm256_add_pd(c0, /* c0 += A[i][k]*B[k][j] */
9.                                     _mm256_mul_pd(_mm256_load_pd(A+i+k*n),
10.                                                    _mm256_broadcast_sd(B+k+j*n)));
11.             _mm256_store_pd(C+i+j*n, c0); /* C[i][j] = c0 */
12.         }
13. }

```

Matrix Multiply, DGEMM (참고자료)

■ Optimized x86 assembly code:

```

1. vmovapd (%r11),%ymm0      # Load 4 elements of C into %ymm0
2. mov %rbx,%rcx             # register %rcx = %rbx
3. xor %eax,%eax             # register %eax = 0
4. vbroadcastsd (%rax,%r8,1),%ymm1 # Make 4 copies of B element
5. add $0x8,%rax             # register %rax = %rax + 8
6. vmulpd (%rcx),%ymm1,%ymm1 # Parallel mul %ymm1, 4 A elements
7. add %r9,%rcx              # register %rcx = %rcx + %r9
8. cmp %r10,%rax             # compare %r10 to %rax
9. vaddpd %ymm1,%ymm0,%ymm0  # Parallel add %ymm1, %ymm0
10. jne 50 <dgemm+0x50>      # jump if not %r10 != %rax
11. add $0x1,%esi            # register %esi = %esi + 1
12. vmovapd %ymm0, (%r11)    # Store %ymm0 into 4 C elements

```

C Intrinsics

- Intrinsic function
 - Function available for use in a given programming language whose implementation is specially handled by compiler
 - Often used to explicitly implement vectorization and parallelization in languages which do not address such constructs
 - E.g., MMX, SSE, OpenMP

x86 FP Architecture (skip)

- Originally based on 8087 FP coprocessor
 - $8 \times$ 80-bit extended-precision registers
 - Used as a push-down stack
 - Registers indexed from TOS: ST(0), ST(1), ...
- FP values are 32-bit or 64 in memory
 - Converted on load/store of memory operand
 - Integer operands can also be converted on load/store
- Very difficult to generate and optimize code
 - Result: poor FP performance

x86 FP Instructions (skip)

Data transfer	Arithmetic	Compare	Transcendental
FILD mem/ST(i) FISTP mem/ST(i) FLDPI FLD1 FLDZ	F _I ADDP mem/ST(i) F _I SUBRP mem/ST(i) F _I MULP mem/ST(i) F _I DIVRP mem/ST(i) FSQRT FABS FRNDINT	F _I COMP F _I UCOMP FSTSW AX/mem	FPATAN F2XMI FCOS FPTAN FPREM FPSIN FYL2X

- Optional variations
 - **I**: integer operand
 - **P**: pop operand from stack
 - **R**: reverse operand order
 - But not all combinations allowed

Related to Chapter 6 (especially Section 6.3)