

COMPUTER ORGANIZATION AND DESIGN



The Hardware/Software Interface

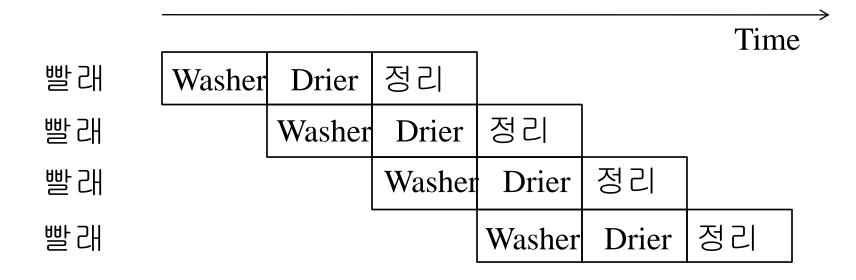
Chapter 4

Implementing ISA (Fetch, Decode, Execute)

Part 2: Pipelining

Pipelined Laundry (반복)

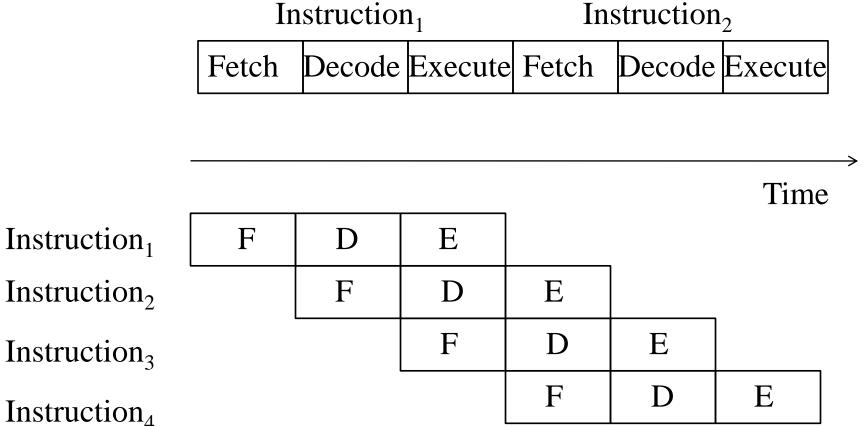
- □ 3-stage pipelining problem
 - Four loads: speedup = 12/6 = 2
 - Infinite loads: speedup = 3 = number of stages
- □ Need more hardware? What if we have more hardware?



† Pipelining: general speedup technique

Pipelining (반복)

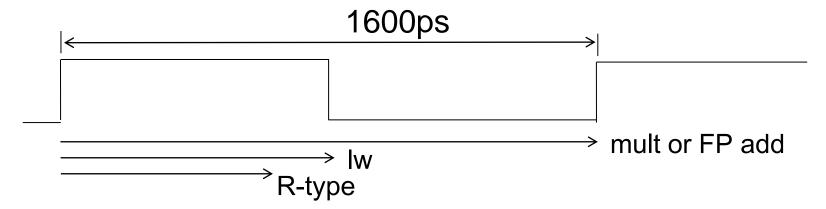
- □ 3-stage pipeline
 - Overlapped execution, instruction-level parallelism



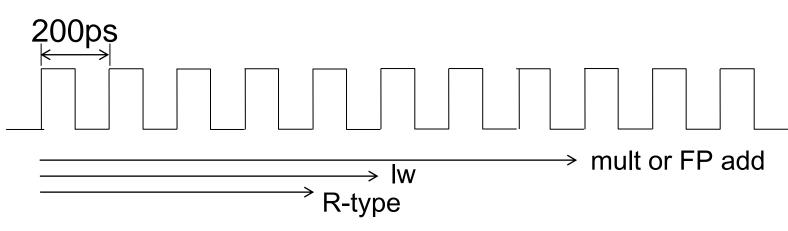
Performance: SingleC/MultiC/Pipelined

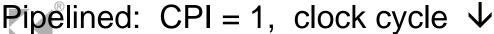
(반복)

 \Box Single-cycle: CPI = 1, clock cycle



■ Multi-cycle: CPI ↑, clock cycle ↓, overall performance ↑





MIPS Pipeline

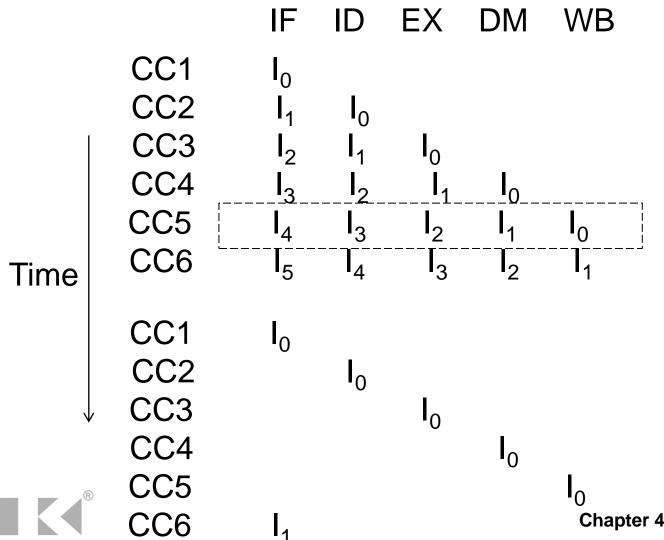
- Five stages, one step per stage
 - 1. IF: Instruction fetch from memory
 - 2. ID: Instruction decode & register read
 - 3. EX: Execute operation or calculate address
 - 4. MEM: Access memory operand
 - 5. WB: Write result back to register

Instr	Instr fetch	Register read	ALU op	Memory access	Register write	Total time
lw	200ps	100 ps	200ps	200ps	100 ps	800ps
SW	200ps	100 ps	200ps	200ps		700ps
R-format	200ps	100 ps	200ps		100 ps	600ps
beq	200ps	100 ps	200ps			500ps



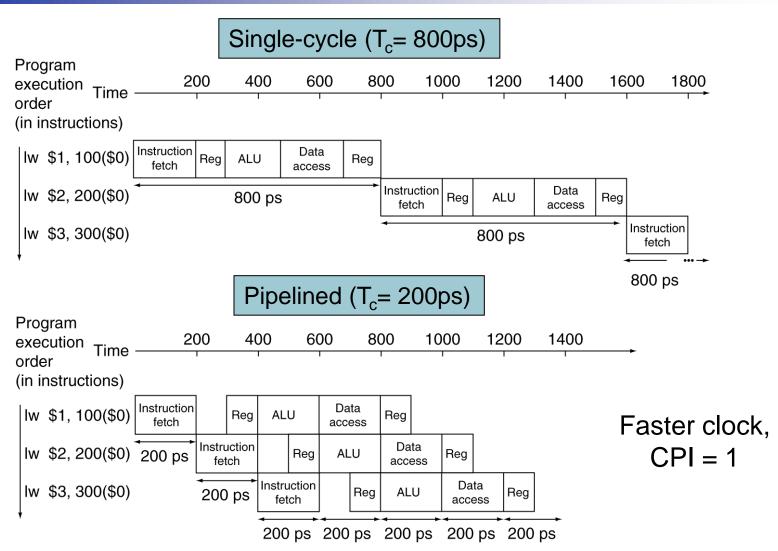
Pipelining vs. Multicycle (반복)

Each stage run different instructions





Pipeline Performance





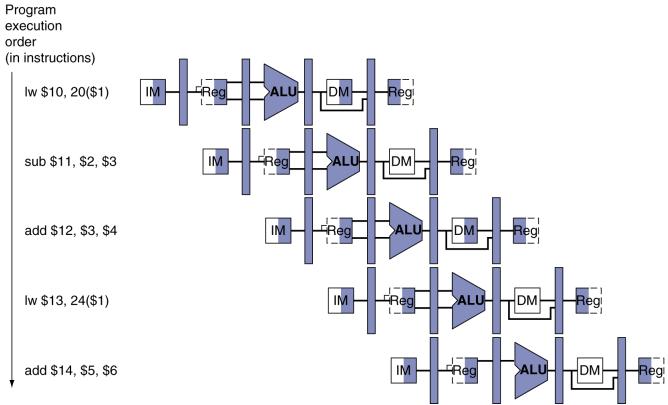
Traditional form

Time (in clock cycles) CC₁ CC 2 CC3 CC 4 CC₆ CC 7 CC8 CC9 CC₅ Program execution order (in instructions) Instruction Instruction Data lw \$10, 20(\$1) Execution Write back decode fetch access Instruction Instruction Data sub \$11, \$2, \$3 Write back Execution fetch decode access Data Instruction Instruction add \$12, \$3, \$4 Write back Execution fetch decode access Instruction Instruction Data lw \$13, 24(\$1) Execution Write back decode access fetch Instruction Instruction Data add \$14, \$5, \$6 Execution Write back fetch decode access



Form showing resource usage







- Each stage run different instructions
 - Only one set of hardware

		IF	ID	EX	DM	WB	
	CC1	I _o					
	CC2	I ₁	I ₀				
	CC3	I_2	I ₁	I_0			
	CC4	I_3	l ₂	I ₁	I ₀		
	CC5	I ₄	I ₃	I ₂	I ₁	 Ι ₀	
\	CC6	I ₅	 ₄	I ₃	 	 I ₁	



Pipeline Speedup

- If all stages are balanced
 - i.e., all take the same time
 - Time between instructions_{pipelined}
 - = Time between instructions_{nonpipelined}

Number of stages

- If not balanced, speedup is less
- Speedup due to increased throughput
 - Latency (time for each instruction) does not decrease



RISC Strategy

- Clock cycle
- CPI
- Instruction count



Pipelining

- What makes it easy
 - All instructions are the same length
 - Just a few instruction formats
 - Memory operands appear only in loads and stores
- What makes it hard?
 - Structural hazards: suppose we had only one memory
 - Control hazards: need to worry about branch
 - Data hazards: instruction depends on previous one
- We'll build a simple pipeline and look at these issues



MIPS Pipeline (skip)

- Assume time for stages is
 - 100ps for register read or write (why?)
 - 200ps for other stages
- Compare pipelined datapath with single-cycle datapath

Instr	Instr fetch	Register read	ALU op	Memory access	Register write	Total time
lw	200ps	100 ps	200ps	200ps	100 ps	800ps
sw	200ps	100 ps	200ps	200ps		700ps
R-format	200ps	100 ps	200ps		100 ps	600ps
beq	200ps	100 ps	200ps			500ps





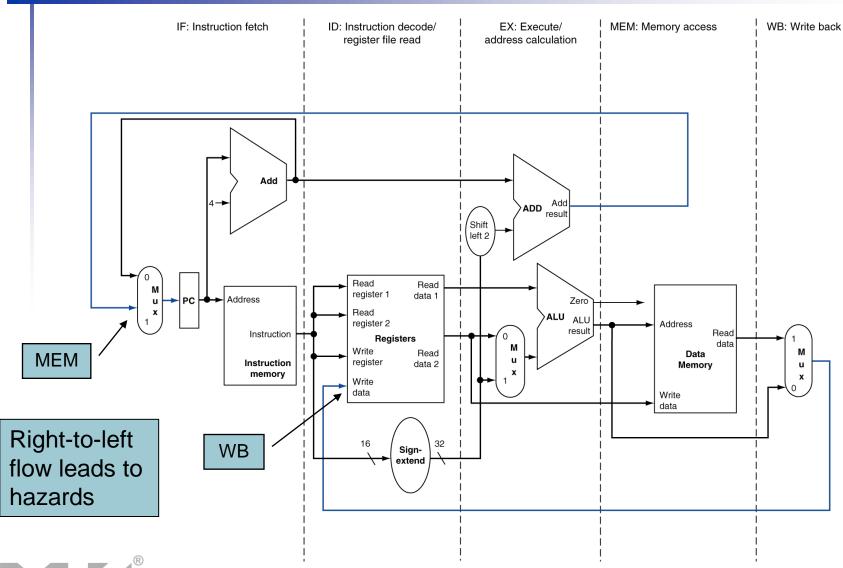
COMPUTER ORGANIZATION AND DESIGN



The Hardware/Software Interface

Pipelined Datapath and Control

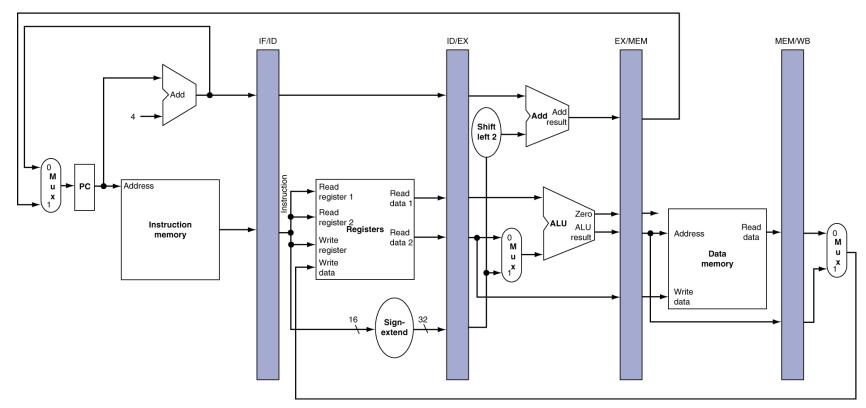
MIPS Pipelined Datapath





Pipeline registers

- Need registers between stages
 - To hold information produced in previous cycle



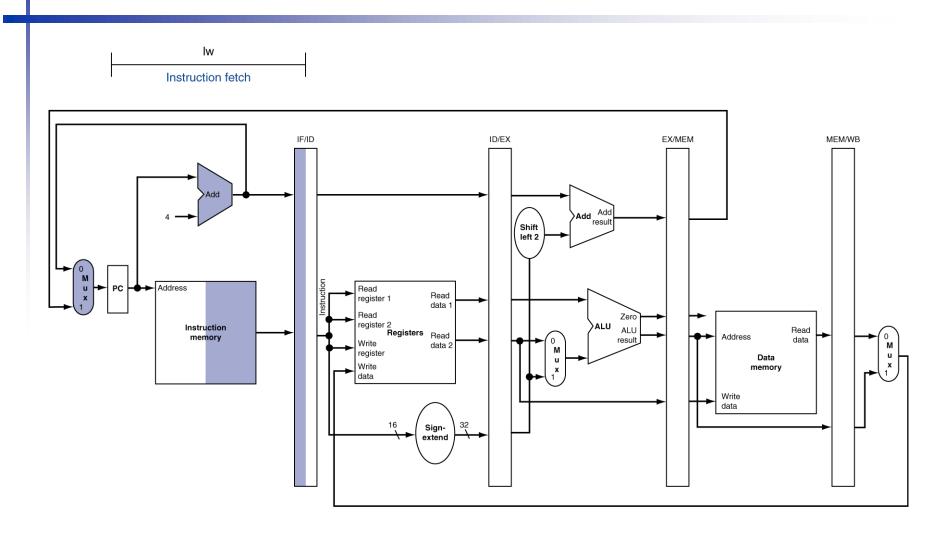


Pipeline Operation

- Cycle-by-cycle flow of instructions through the pipelined datapath
 - We'll look at diagrams for load & store



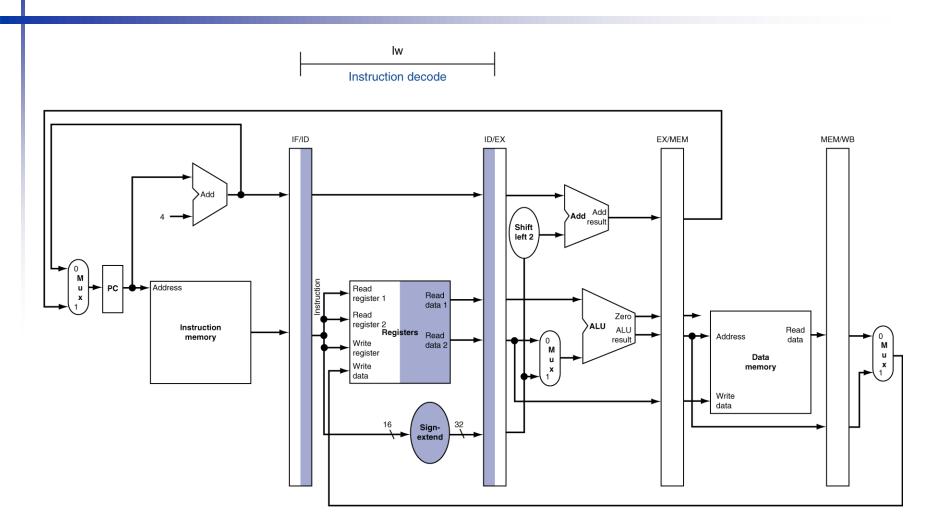
IF for Load, Store, ...



☐ What kinds of information do we carry with IF/ID latch?



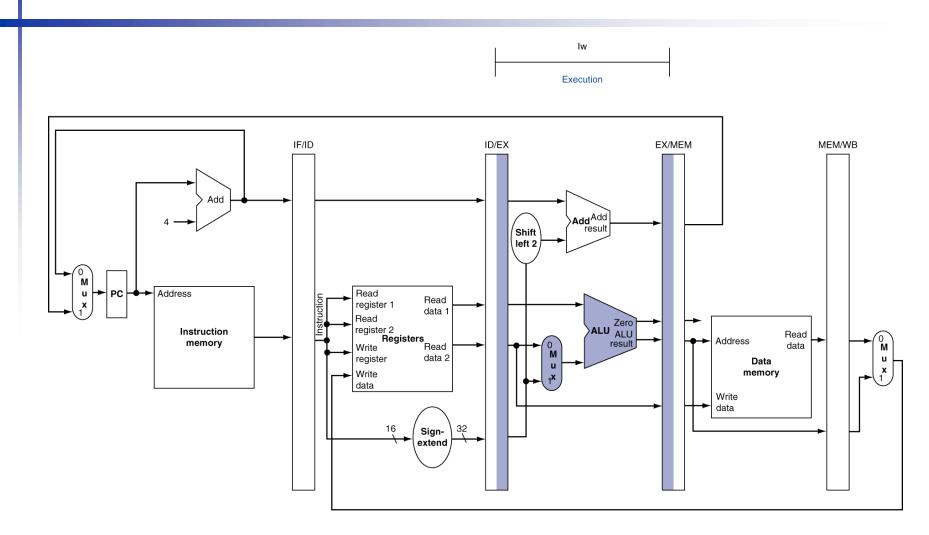
ID for Load, Store, ...



☐ What kinds of information do we carry with ID/EX latch?



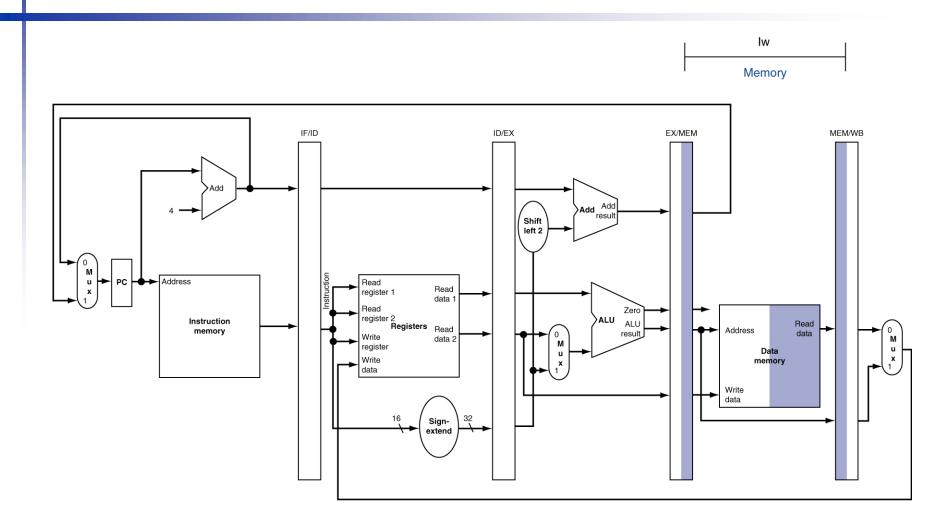
EX for Load



☐ What kinds of information do we carry with EX/MEM latch?



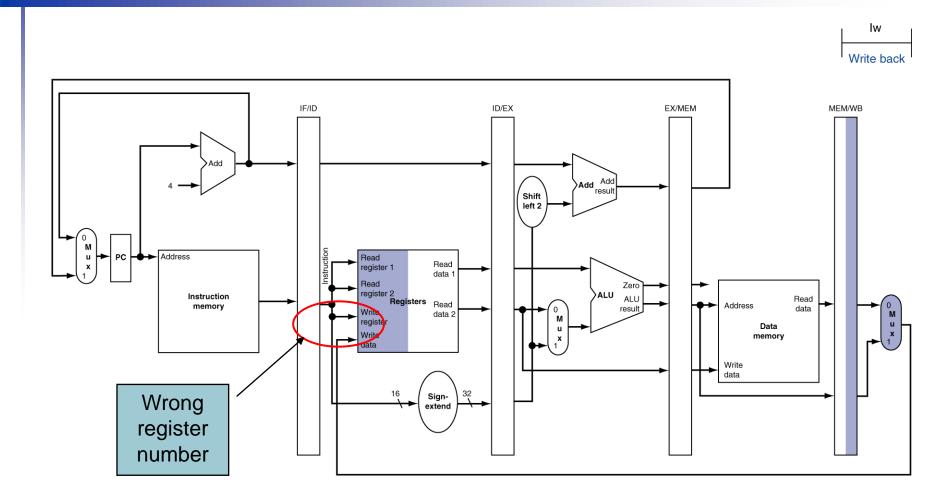
MEM for Load



☐ What kinds of information do we carry with MEM/WB latch?

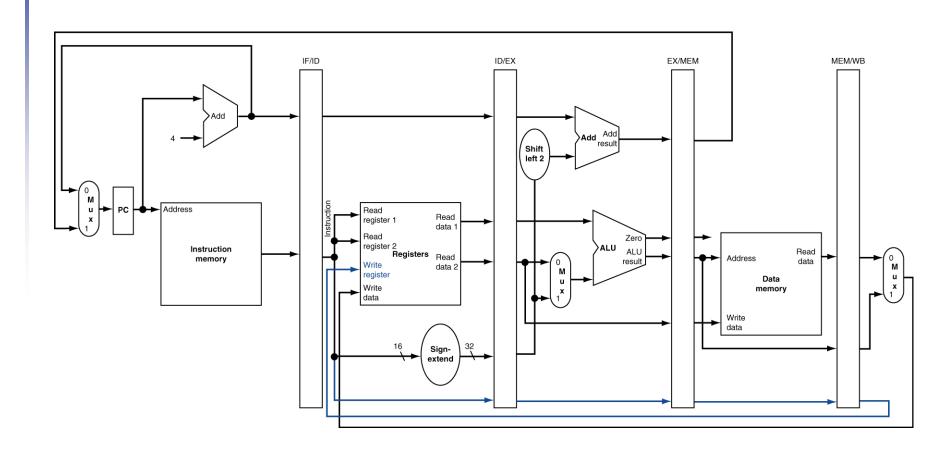


WB for Load



■ What kinds of information do we carry with MEM/WB latch?

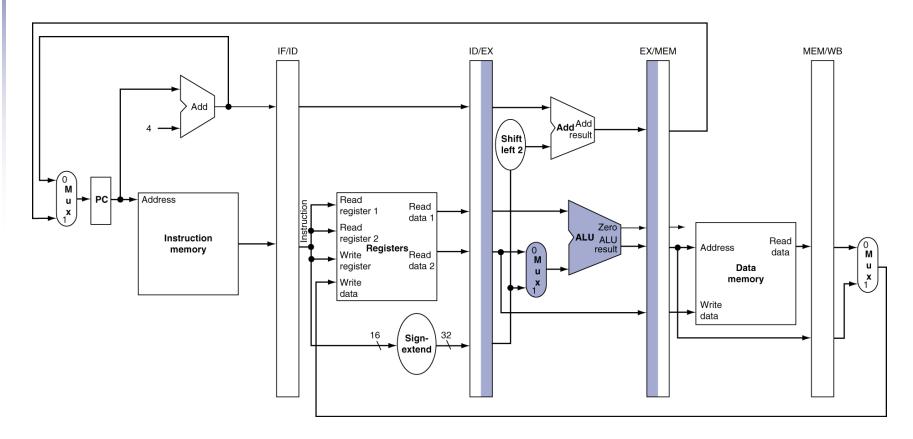
Corrected Datapath for Load





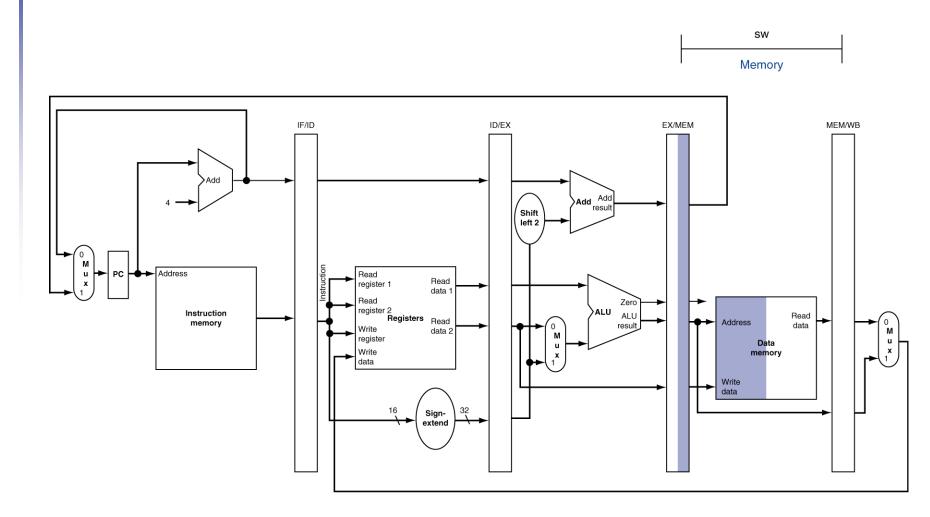
EX for Store





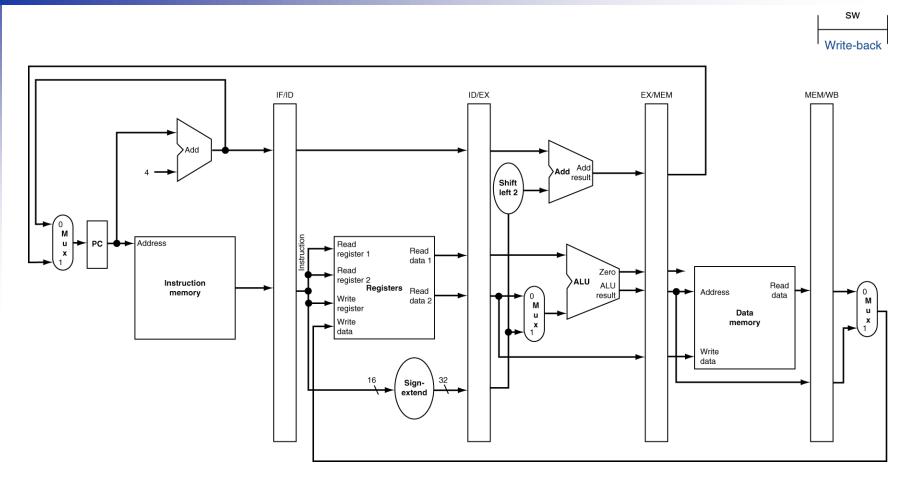


MEM for Store





WB for Store



- Doesn't matter whether there is meaningful operation
 - Important: at every cycle, start (or end) one instruction

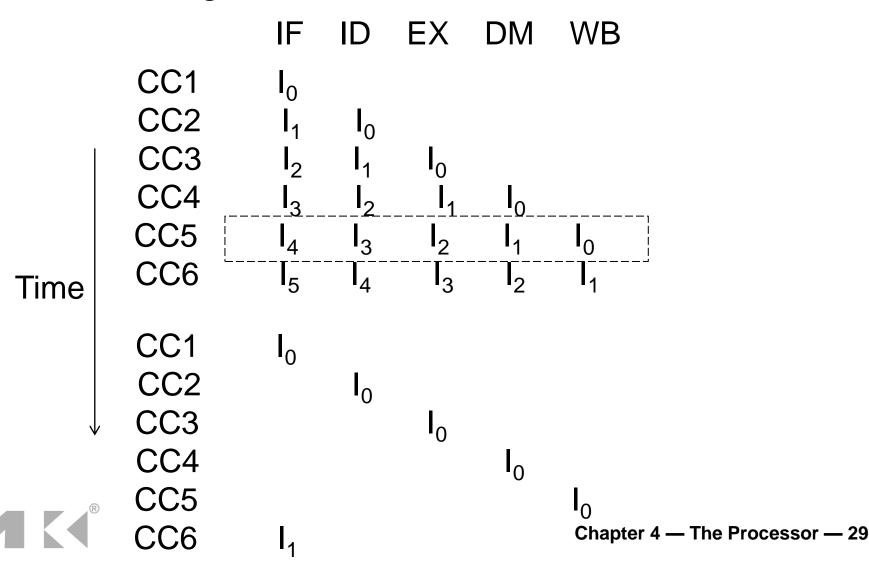
ALU instruction and Branch

- Operation of R-type instructions ("add")
 - No operation in MEM
 - Not affect overall performance
 - Important: at every cycle, start one instruction
- Operation of "addi"
- Operation of branch instructions ("beq")
 - Finished in MEM, no operation in WB

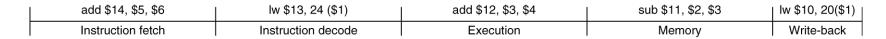


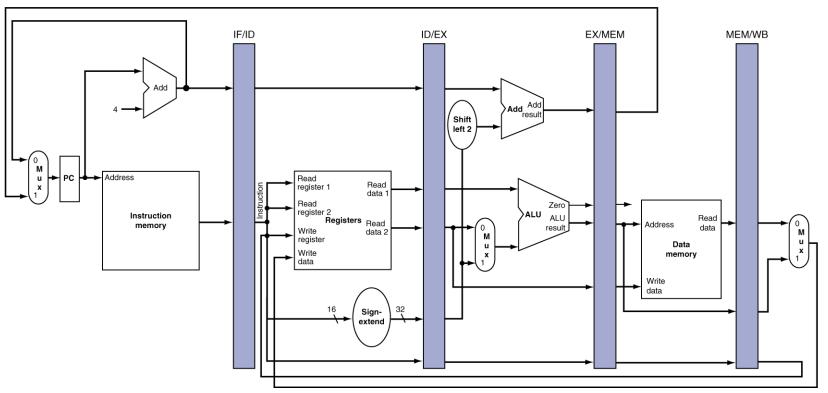
Pipelining vs. Multicycle

Each stage run different instructions



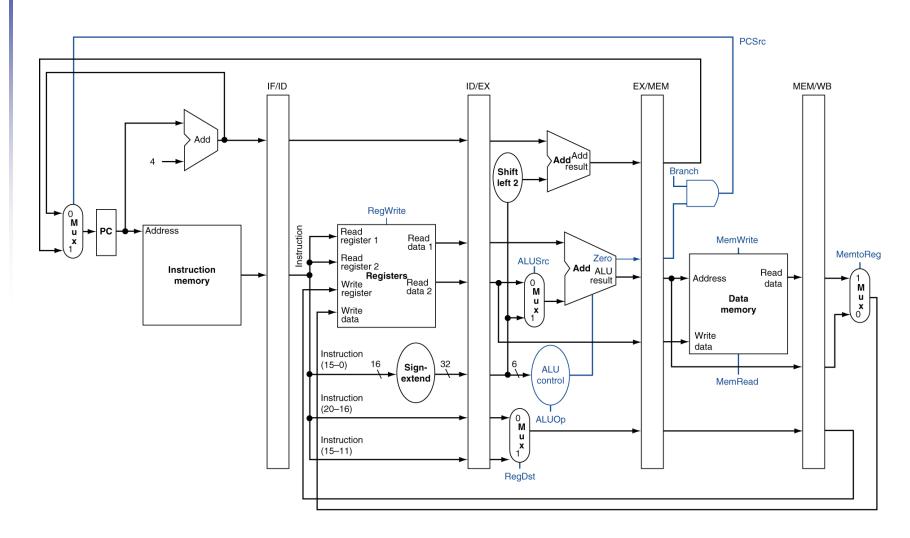
State of pipeline in a given cycle







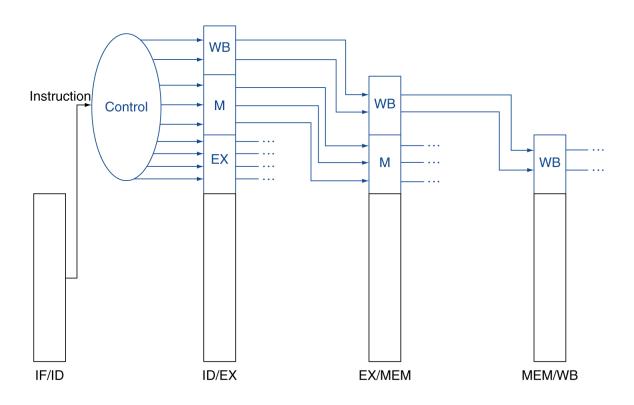
Pipelined Control (Simplified)





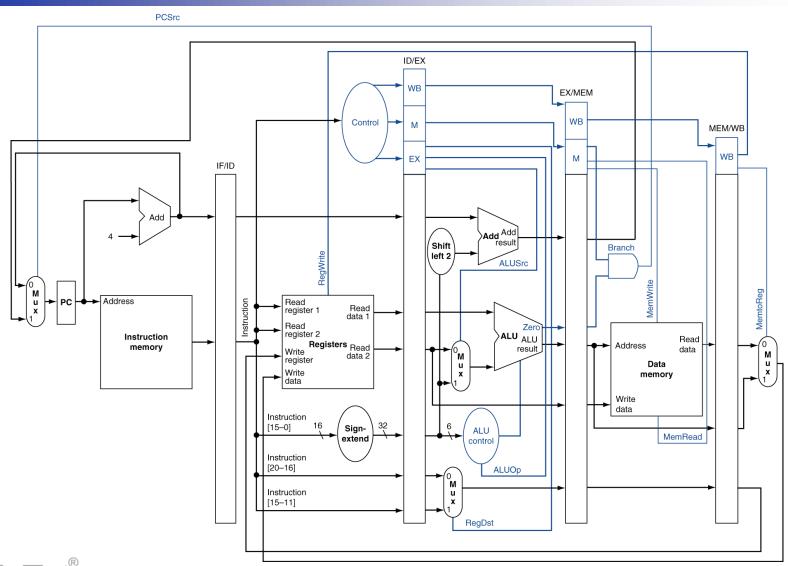
Pipelined Control

- Control signals derived from instruction
 - As in single-cycle implementation





Pipelined Control



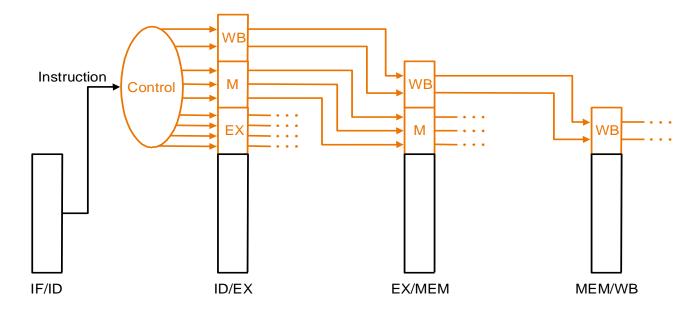


Pipeline Control

Pass control signals along just like the data

	Execution/Address Calculation stage control lines				Memory access stage control lines			stage control lines	
Instruction	Reg Dst	ALU Op1	ALU Op0	ALU Src	Branch	Mem Read	Mem Write	Reg write	Mem to Reg
R-format	1	1	0	0	0	0	0	1	0
lw	0	0	0	1	0	1	0	1	1
SW	X	0	0	1	0	0	1	0	X
beq	X	0	1	0	1	0	0	0	X

† Note the register destination signal





Pipelining and ISA Design

- MIPS ISA designed for pipelining
 - All instructions are 32-bits
 - Easier to fetch and decode in one cycle
 - c.f. x86: 1- to 17-byte instructions
 - Few and regular instruction formats
 - Can decode and read registers in one step
 - Load/store addressing
 - Can calculate address in 3rd stage, access memory in 4th stage
 - Alignment of memory operands
 - Memory access takes only one cycle





COMPUTER ORGANIZATION AND DESIGN



The Hardware/Software Interface

Pipeline Hazards

Hazards

- Situations that prevent starting the next instruction in the next cycle
- 1) Structure hazards
 - A required resource is busy
- 2) Data hazard
 - Need to wait for previous instruction to complete its data read/write
- 3) Control hazard
 - Deciding on control action depends on previous instruction



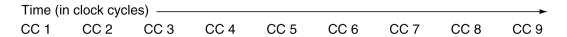
Structural Hazards

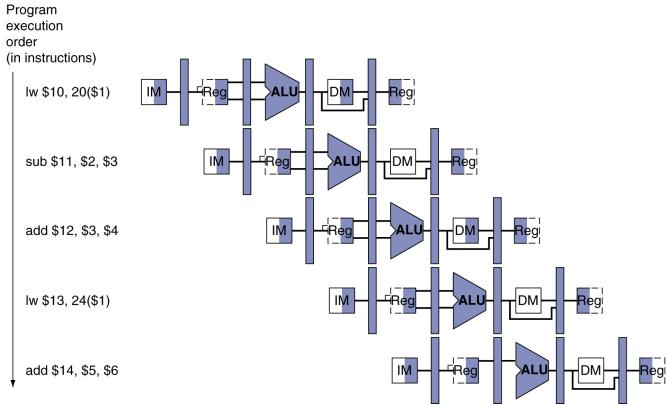
- Resource conflicts
 - Different stages want to use same hardware resource
 - Can avoid with hardware duplication
- Our datapath has no structural hazards, but what if
 - Instruction and data memory not separated
 - fetch vs. DM of lw/sw
 - Not separate adder for PC + 4
 - Register access
 - ID vs. write back
- Sometimes, can be better not to pipeline
 - Return on investment (e.g., divider)



Structural Hazards

Integrated IM and DM – impact on CPI?







Data Hazards in ALU Instructions

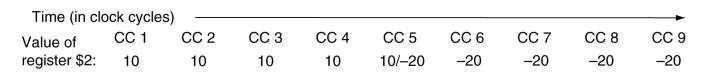
Consider this sequence:

```
sub $2, $1,$3
and $12,$2,$5
or $13,$6,$2
add $14,$2,$2
sw $15,100($2)
```

- We can resolve hazards with forwarding
 - How do we detect when to forward?



Dependencies & Forwarding

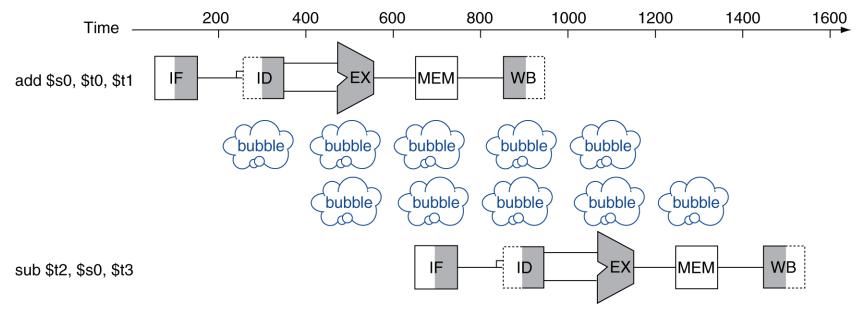


Program execution order (in instructions) sub \$2, \$1, \$3 IM and \$12, \$2, \$5 or \$13, \$6, \$2 add \$14, \$2,\$2 sw \$15, 100(\$2)



Data Hazards

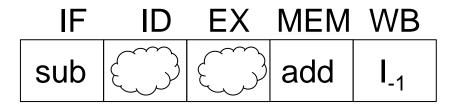
- An instruction depends on completion of data access by a previous instruction
 - add \$s0, \$t0, \$t1
 sub \$t2, \$s0, \$t3

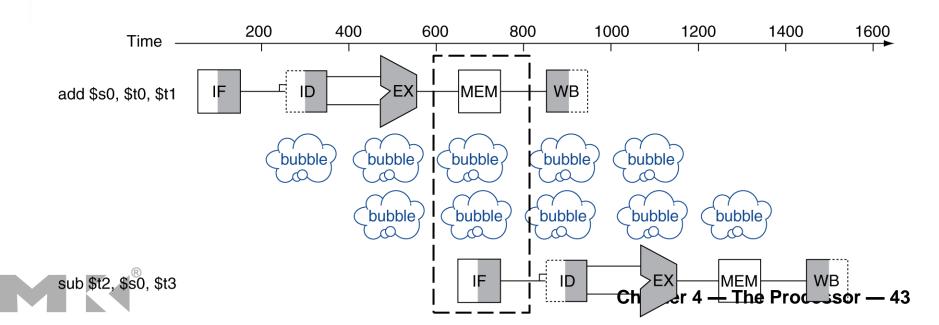




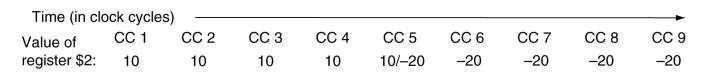
Pipeline Stall, Bubble

- Lose two cycles: increase in CPI
 - What is bubble: NOP instruction





Dependencies & Forwarding

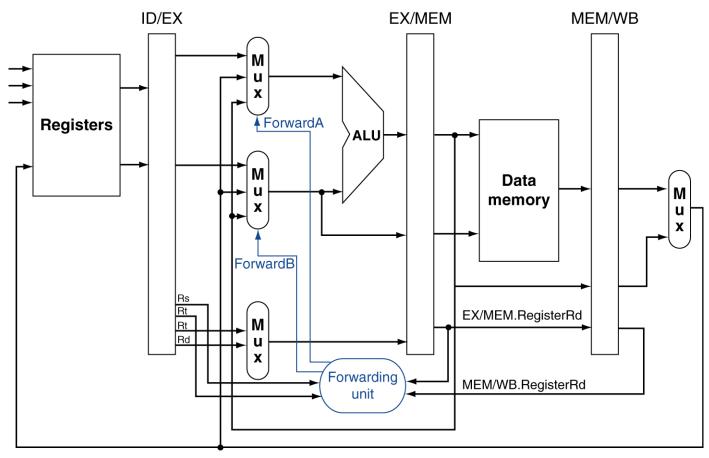


Program execution order (in instructions) sub \$2, \$1, \$3 IM and \$12, \$2, \$5 or \$13, \$6, \$2 add \$14, \$2,\$2 sw \$15, 100(\$2)



Forwarding Paths

More datapath and control



b. With forwarding



Think about Forwarding

- Source: instructions that generate new value (producer)
 - New values generated by ALU or data memory read
 - Non-producer: sw, beq
- Data receiver: data consumer
 - All instructions
 - For ALU operations or write to data memory
- To sum up
 - Forwarding source:
 - ALU result, delayed ALU results, result of Iw
 - Destination of forwarding: ALU inputs, write data to DM



Think about Forwarding

- Forwarding to ALU inputs
 - From EX/MEM latch (I₋₁) or MEM/WB latch (I₋₂)
 - Results of I₋₃ available from registers
- Will not cover forwarding to DM write data input
 - Your personal exercise



Detecting the Need to Forward

- ALU operand register numbers in EX stage are given by
 - ID/EX.RegisterRs, ID/EX.RegisterRt
- Data hazards when
 - 1a. EX/MEM.RegisterRd = ID/EX.RegisterRs
 - 1b. EX/MEM.RegisterRd = ID/EX.RegisterRt
 - 2a. MEM/WB.RegisterRd = ID/EX.RegisterRs
 - 2b. MEM/WB.RegisterRd = ID/EX.RegisterRt

Fwd from EX/MEM pipeline reg

Fwd from MEM/WB pipeline reg



Detecting the Need to Forward

- But only if forwarding instruction will write to a register!
 - EX/MEM.RegWrite, MEM/WB.RegWrite
- And only if Rd for that instruction is not \$zero
 - EX/MEM.RegisterRd ≠ 0,
 MEM/WB.RegisterRd ≠ 0



Forwarding Conditions

- EX hazard
 - if (EX/MEM.RegWrite and (EX/MEM.RegisterRd ≠ 0) and (EX/MEM.RegisterRd = ID/EX.RegisterRs)) ForwardA = 10
 - if (EX/MEM.RegWrite and (EX/MEM.RegisterRd ≠ 0) and (EX/MEM.RegisterRd = ID/EX.RegisterRt)) ForwardB = 10
- MEM hazard
 - if (MEM/WB.RegWrite and (MEM/WB.RegisterRd ≠ 0) and (MEM/WB.RegisterRd = ID/EX.RegisterRs)) ForwardA = 01
 - if (MEM/WB.RegWrite and (MEM/WB.RegisterRd ≠ 0) and (MEM/WB.RegisterRd = ID/EX.RegisterRt))





Think about Forwarding

- Difference between add and addi instructions
 - Two ALU inputs vs. one ALU input
 - Why do we ignore it?
 - ALU do not use that input anywat
 - Note that datapath for 16-bit immediate part not shown in figure



Double Data Hazard

Consider the sequence:

```
add $1,$1,$2
add $1,$1,$3
add $1,$1,$4
```

- Both hazards occur
 - Want to use the most recent
- Revise MEM hazard condition
 - Only fwd if EX hazard condition isn't true



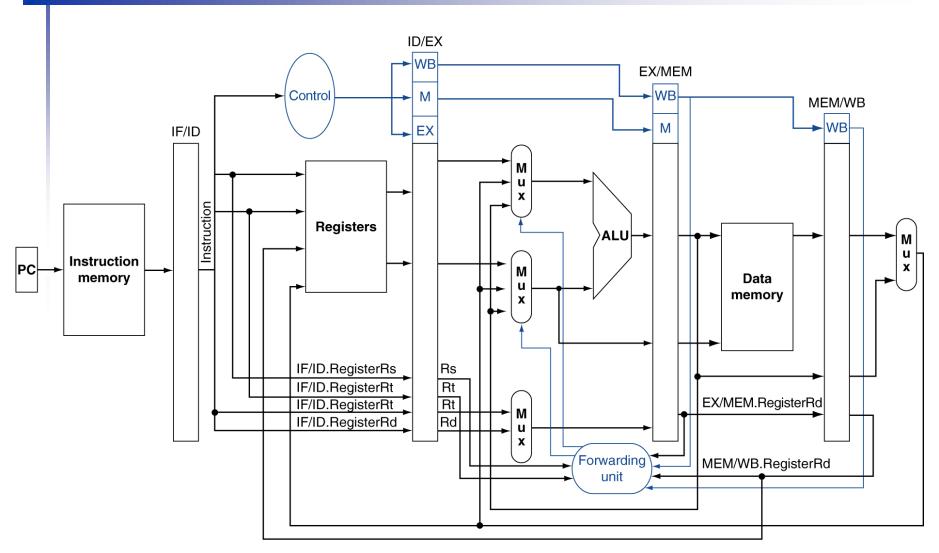
Revised Forwarding Condition

- MEM hazard
 - if (MEM/WB.RegWrite and (MEM/WB.RegisterRd ≠ 0) and not (EX/MEM.RegWrite and (EX/MEM.RegisterRd ≠ 0) and (EX/MEM.RegisterRd = ID/EX.RegisterRs)) and (MEM/WB.RegisterRd = ID/EX.RegisterRs)) ForwardA = 01
 - if (MEM/WB.RegWrite and (MEM/WB.RegisterRd ≠ 0)

ForwardB = 01



Datapath with Forwarding



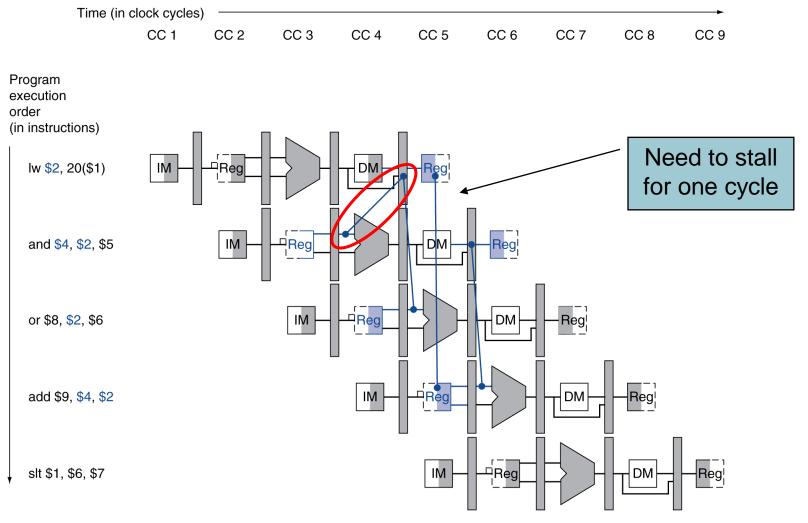


Load-Use Data Hazard

Only data hazard that require stall in our 5-stage design



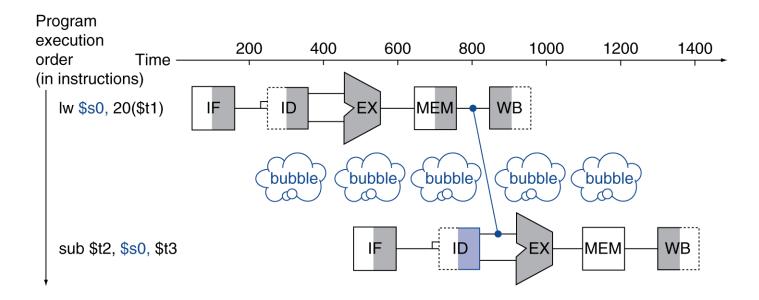
Load-Use Data Hazard





Load-Use Data Hazard

- Can't always avoid stalls by forwarding
 - If value not computed when needed
 - Can't forward backward in time!





Load-Use Hazard Detection

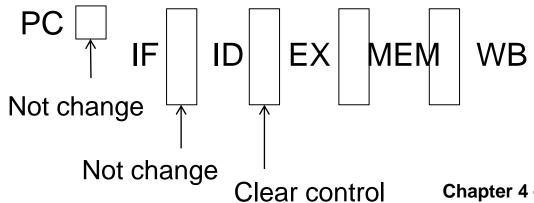
- Check when using instruction is decoded in ID stage
- ALU operand register numbers in ID stage are given by
 - IF/ID.RegisterRs, IF/ID.RegisterRt
- Load-use hazard when
 - ID/EX.MemRead and ((ID/EX.RegisterRt = IF/ID.RegisterRs) or (ID/EX.RegisterRt = IF/ID.RegisterRt))
- If detected, stall and insert bubble



How to Stall the Pipeline

	IF	ID	EX	MEM	WB
CC1	or	and	lw	 ₋₁	l ₋₂
CC2		and	- - ##-		
CC3	add	or	and	l nop	lw
CC4		add		and	nop)
CC5	I ₊₁	slt	add	or	and

☐ To insert NOP (or bubble), at the end of clock, do:



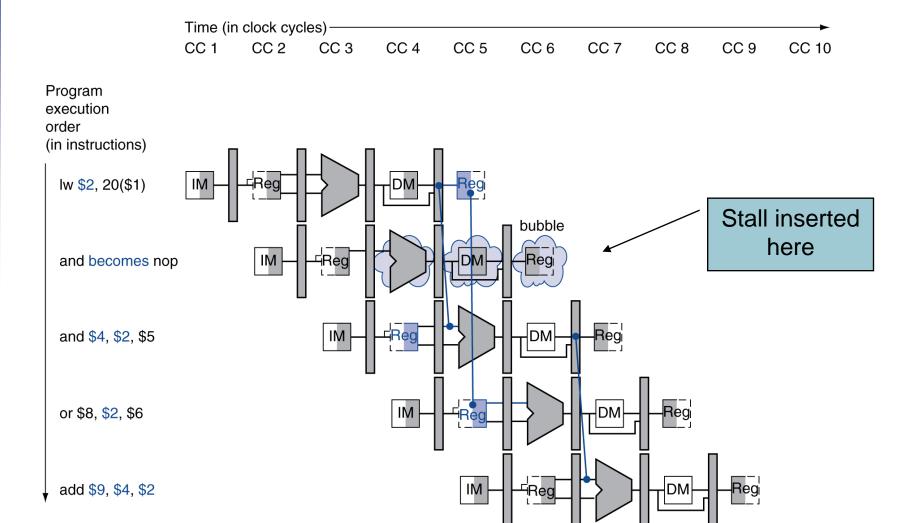


How to Stall the Pipeline

- Force control values in ID/EX register to 0
 - EX do "nop" (no-operation)
- Prevent update of PC and IF/ID register
 - ID stage: same instruction is decoded again
 - IF stage: same instruction is fetched again
 - 1-cycle stall allows MEM to read data for \(\)\
 - Can subsequently forward to EX stage

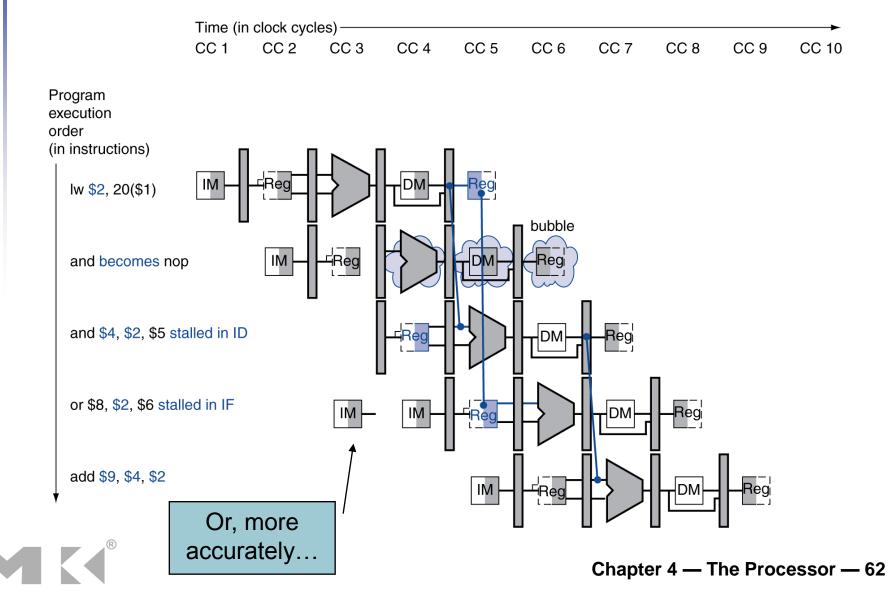


Stall/Bubble in the Pipeline (skip)

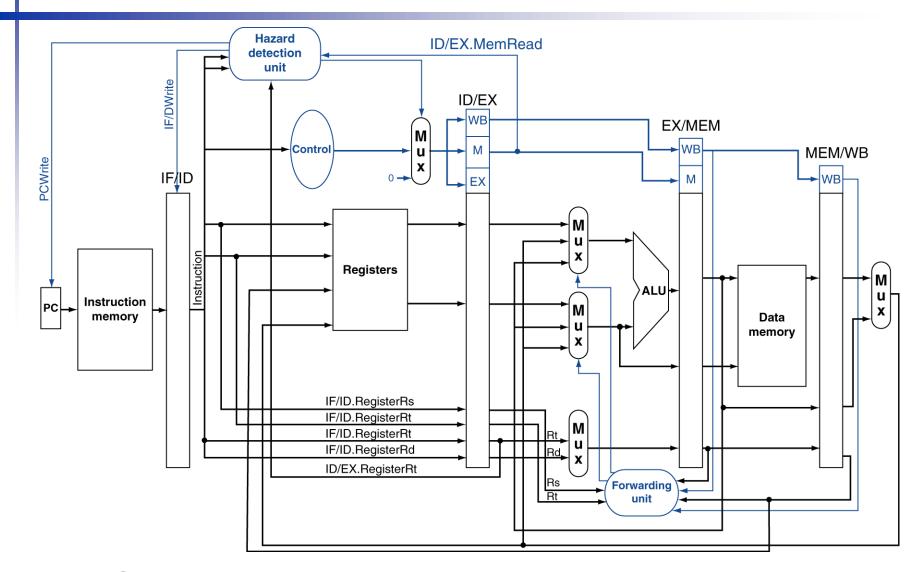




Stall/Bubble in the Pipeline (skip)



Datapath with Hazard Detection







Software Solution

Compiler can help



Software Solution

- Let compiler guarantee no hazards
 - Where do we insert the "nops" ?

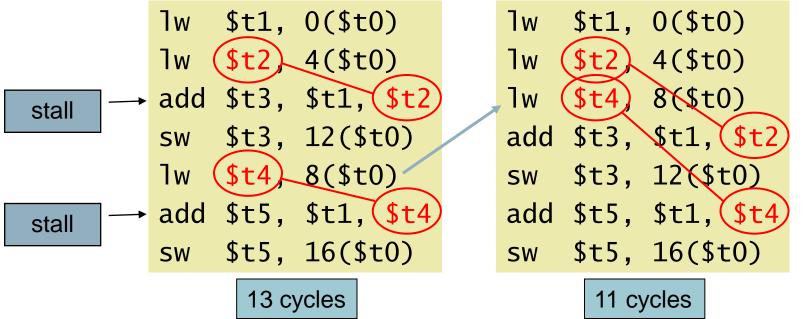
```
sub $2, $1, $3
and $12, $2, $5
or $13, $6, $2
add $14, $2, $2
sw $15, 100($2)
```

- Problem: this really slows us down! Layer violation!
 - Use forwarding!



Code Scheduling to Avoid Stalls

- Reorder code to avoid use of load result in the next instruction
- C code for A = B + E; C = B + F;





Software Solution

- Compiler can arrange code to avoid load-use stalls
 - Requires knowledge of the pipeline structure
 - Layer violation and dependency
- Hardware can also do it without difficulty
 - Runtime (dynamic) out-of-order execution
 - No compiler dependency





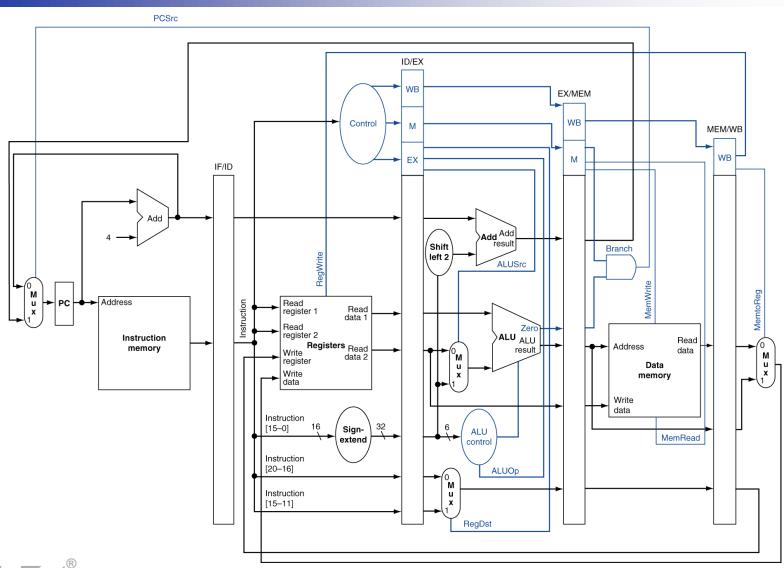
COMPUTER ORGANIZATION AND DESIGN



The Hardware/Software Interface

Branch Hazards

Pipelined Datapath and Control





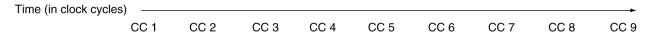
Control Hazards

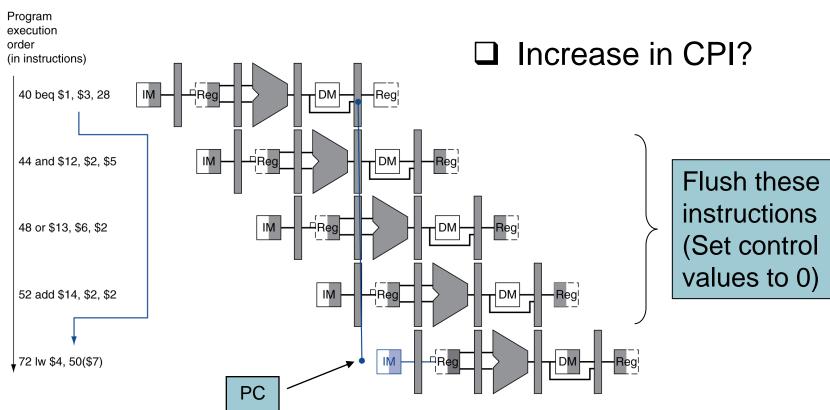
- Branch determines flow of control
 - Fetching next instruction depends on branch outcome
 - Pipeline can't always fetch correct instruction
 - Still working on ID stage of branch
- In MIPS pipeline
 - Need to compare registers and compute target early in the pipeline
 - Add hardware to do it in ID stage



Branch Hazards

If branch outcome determined in MEM







Reducing Branch Delay

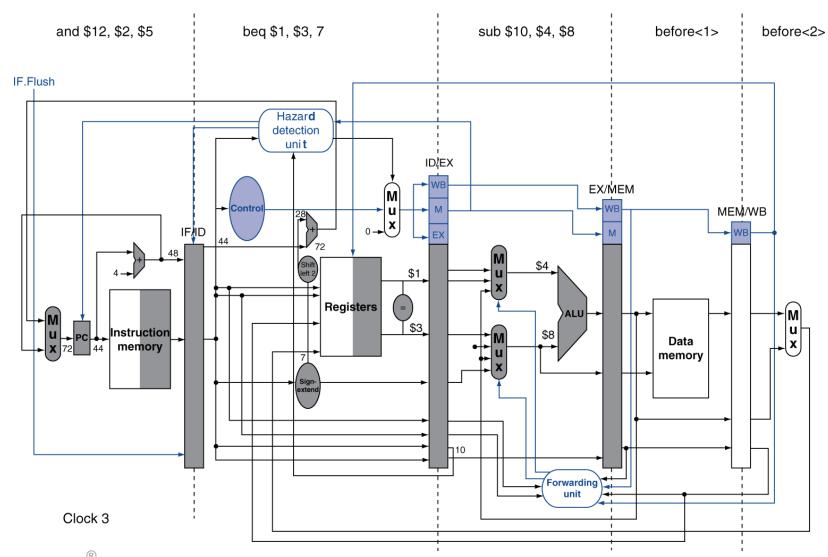
- Move hardware to determine outcome to ID stage
 - Target address adder
 - Register comparator
- Example: branch taken

```
36: sub $10, $4, $8
40: beq $1, $3, 7
44: and $12, $2, $5
48: or $13, $2, $6
52: add $14, $4, $2
56: slt $15, $6, $7
72: lw $4, 50($7)
```

// see next slide



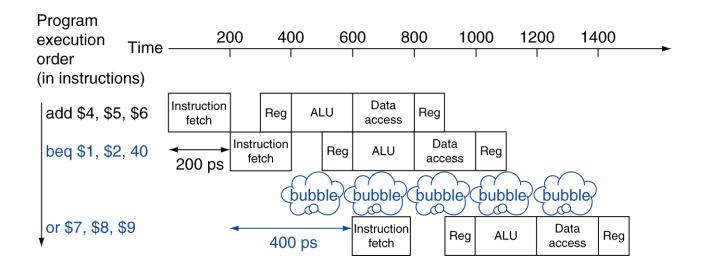
Moving Branch to ID Stage





Stall on Branch

 Wait until branch outcome determined before fetching next instruction



Increase in CPI?



Moving Branch to ID Stage

- Compute branch target address in ID
 - MIPS branch rely on simple test (equal, sign)
- Evaluate branch decision in ID
 - Additional forwarding
 - From ALU/MEM, MEM/WB latches
 - Additional stall
 - IF.Flush control signal
- Think about ARM conditional instruction (skip)



Branch Prediction

- Longer pipelines can't readily determine branch outcome early
 - Stall penalty becomes unacceptable
- Predict outcome of branch
 - Only stall if prediction is wrong
- In MIPS pipeline
 - Can predict branches not taken
 - Fetch instruction after branch, with no delay

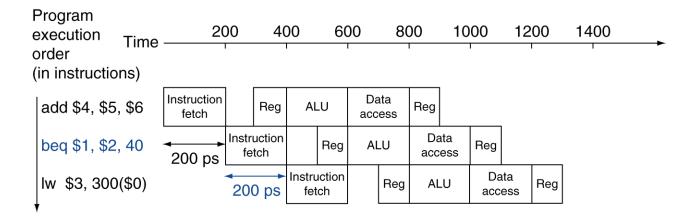


MIPS with (Static) Predict Not Taken

200

400

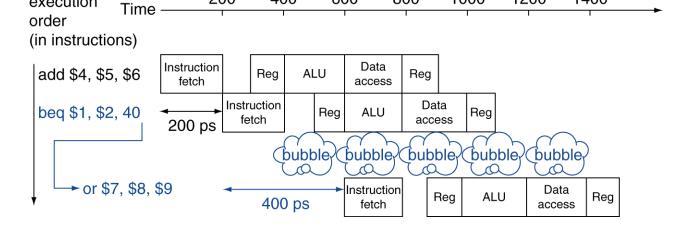
Prediction correct



Prediction incorrect

Program

execution



600

800

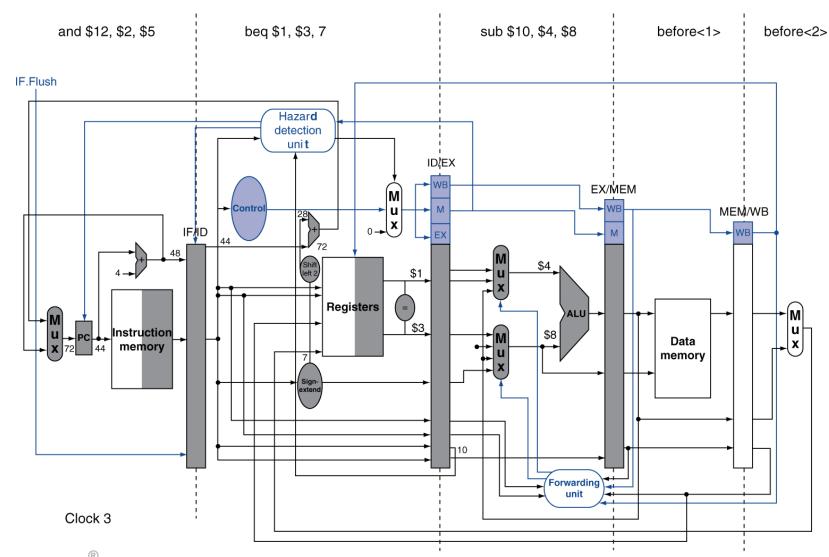
1000

1200



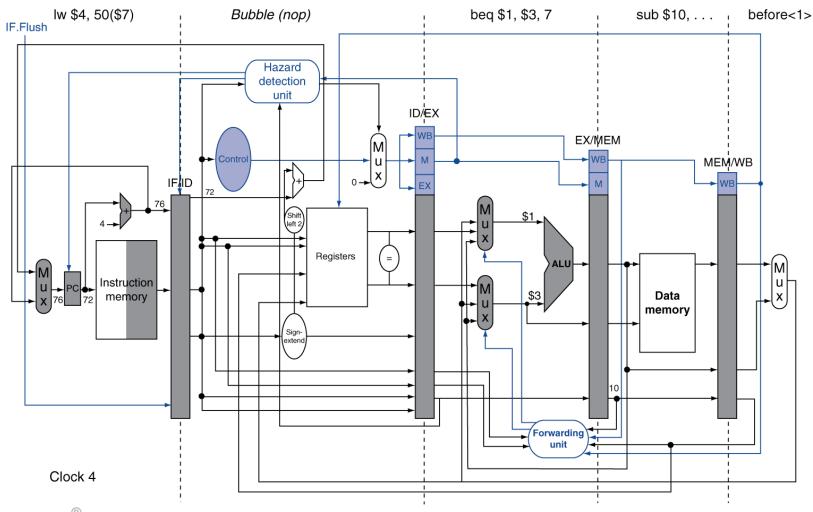
1400

Example: Branch Taken





Example: Branch Taken





Data Hazards for Branches

 If a comparison register is a destination of 2nd or 3rd preceding ALU instruction

Can resolve using forwarding



Data Hazards for Branches

- If a comparison register is a destination of preceding ALU instruction or 2nd preceding load instruction
 - Need 1 stall cycle



Data Hazards for Branches

- If a comparison register is a destination of immediately preceding load instruction
 - Need 2 stall cycles



More–Realistic Branch prediction

- The remaining part is optional
- Branch prediction is important in deep pipeline



More-Realistic Branch Prediction

- Static branch prediction
 - Based on typical branch behavior
 - Example: loop and if-statement branches
 - Predict backward branches taken
 - Predict forward branches not taken
- Dynamic branch prediction
 - Hardware measures actual branch behavior
 - e.g., record recent history of each branch
 - Assume future behavior will continue the trend
 - When wrong, stall while re-fetching, and update history



Dynamic Branch Prediction

- In deeper and superscalar pipelines, branch penalty is more significant
- Use dynamic prediction
 - Branch prediction buffer (aka branch history table)
 - Indexed by recent branch instruction addresses
 - Stores outcome (taken/not taken)
 - To execute a branch
 - Check table, expect the same outcome
 - Start fetching from fall-through or target
 - If wrong, flush pipeline and flip prediction



1-Bit Predictor: Shortcoming

Inner loop branches mispredicted twice!

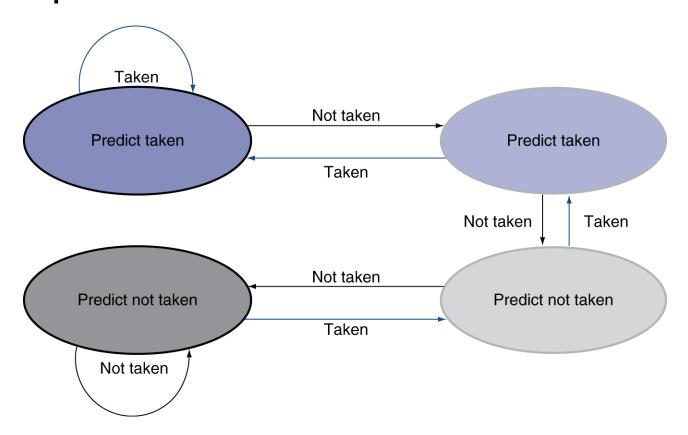
```
outer: ...
inner: ...
beq ..., ..., inner
...
beq ..., ..., outer
```

- Mispredict as taken on last iteration of inner loop
- Then mispredict as not taken on first iteration of inner loop next time around



2-Bit Predictor

 Only change prediction on two successive mispredictions





Calculating the Branch Target

- Even with predictor, still need to calculate the target address
 - 1-cycle penalty for a taken branch
- Branch target buffer
 - Cache of target addresses
 - Indexed by PC when instruction fetched
 - If hit and instruction is branch predicted taken, can fetch target immediately



Pipeline Summary

The BIG Picture

- Pipelining improves performance by increasing instruction throughput
 - Executes multiple instructions in parallel
 - Each instruction has the same latency
- Subject to hazards
 - Structure, data, control
- Instruction set design affects complexity of pipeline implementation

