

# **Chapter 6**

## **Parallel Processors from Client to Cloud**

**- A Very Quick Look**

# **Section 6.1**

## **Introduction**

# Hardware and Software

- Hardware
  - Serial: e.g., Pentium 4
  - Parallel: e.g., quad-core Xeon e5345
- Software
  - Sequential: e.g., matrix multiplication
  - Concurrent: e.g., operating system
- Sequential/concurrent software can run on serial/parallel hardware
  - Challenge: making effective use of parallel hardware

# What We've Already Covered

- §2.10: Parallelism and Instructions
  - Synchronization
- §3.6: Parallelism and Computer Arithmetic
  - Subword Parallelism
- §4.10: Parallelism and Advanced Instruction-Level Parallelism
- §5.10: Parallelism and Memory Hierarchies
  - Cache Coherence

## **Section 6.2**

**The difficulty of creating  
parallel processing  
programs**

# Parallel Programming

- Parallel software is the problem
- Difficult to write software that use multiple processors to complete one task faster
  - Gets worse as number of processors increase
- Why difficult?
  - Partitioning, load balancing
  - Coordination (synchronization, scheduling)
  - Communications overhead

# Parallel Programming

- Need to get significant performance improvement
  - Otherwise, just use a faster uniprocessor
- Sequential programming is simpler and
- Uniprocessor design techniques exploit instruction-level parallelism without involvement of programmers
  - Pipelining
  - Superscalar
  - Out-of-order execution

## **Section 6.3**

**SISD, MIMD, SIMD, SPMD,  
and Vector**



# Instruction and Data Streams

- An old classification of parallel hardware

		Data Streams	
		Single	Multiple
Instruction Streams	Single	<b>SISD:</b> Intel Pentium 4	<b>SIMD:</b> SSE instructions of x86
	Multiple	<b>MISD:</b> No examples today	<b>MIMD:</b> Intel Xeon e5345

- SPMD: Single Program Multiple Data
  - A parallel program on a MIMD computer
  - Conditional code for different processors

# MIMD and SIMD

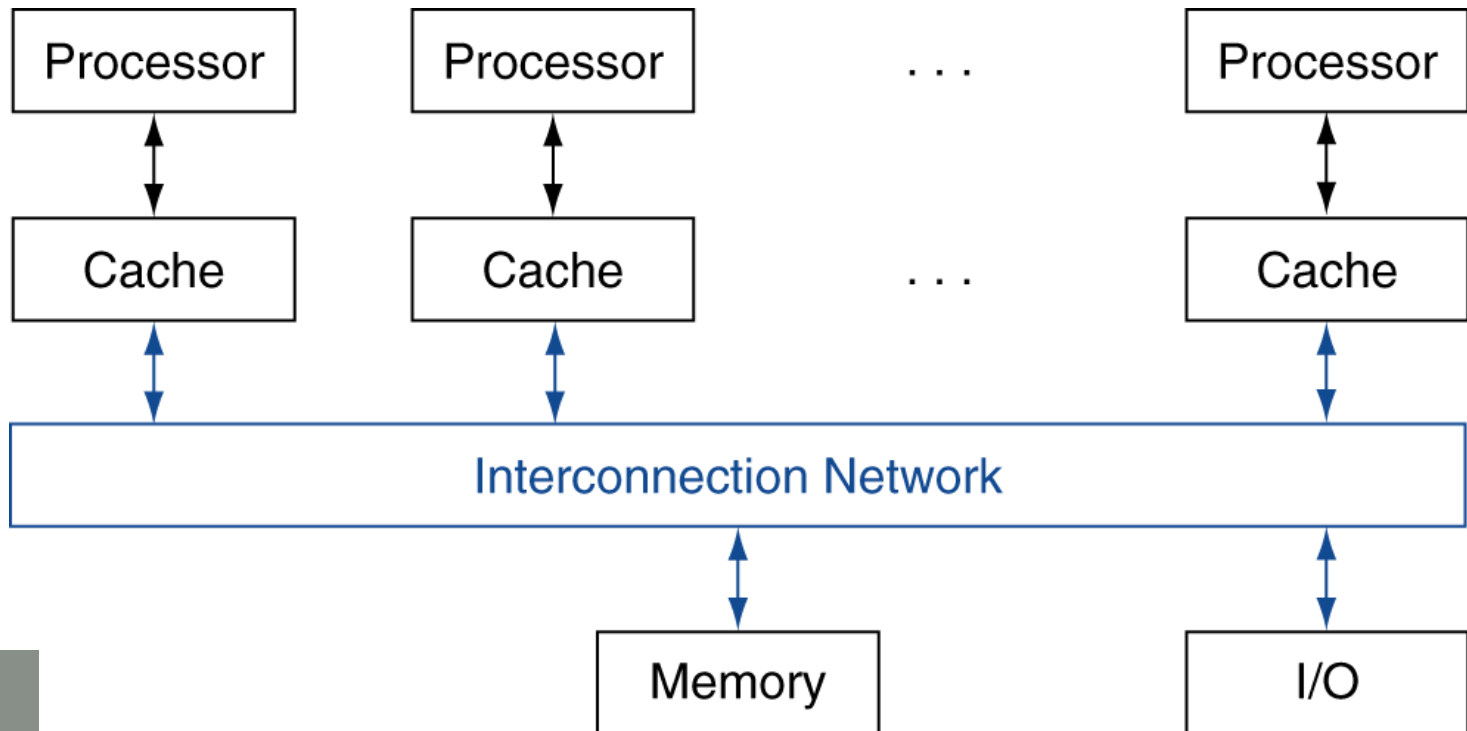
- MIMD: our focus
  - Two flavours
    - Shared memory multiprocessors (SMPs)
    - Clusters
- Let's examine SIMD first

## **Section 6.5**

**Multicore and other  
shared memory  
multiprocessors**


# Shared Memory

- SMP: shared memory multiprocessor
  - Hardware provides single physical address space for all processors
  - Synchronize shared variables using locks




# Uniprocessor System

- ❑ Multiple virtual addresses and single main memory
  - Execute processes sequentially
  - Virtual memory: protection and sharing




# Shared Memory Multiprocessor

- ❑ Multiple virtual addresses and single main memory
  - Run processes simultaneously (throughput, resp. time)
  - Virtual memory: protection and sharing



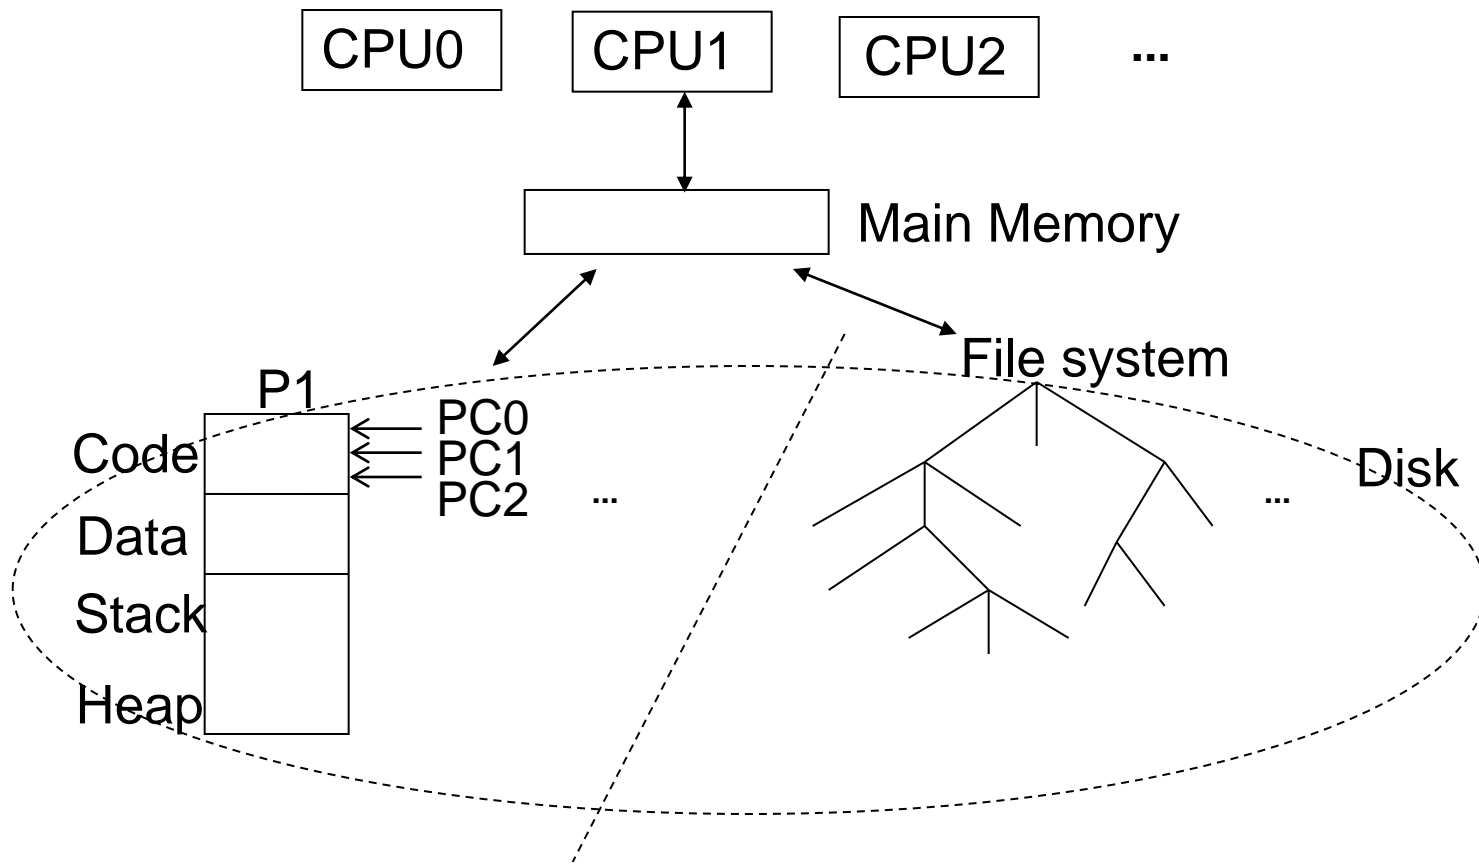
# Uniprocessor System

- ❑ Single virtual address and single main memory
  - Execute threads sequentially
    - Each has own PC, code area, data area



# Shared Memory Multiprocessor

- ❑ Single virtual address and single main memory
  - Run threads simultaneously (throughput, resp. time)
    - Each has own PC, code area, data area





# Example: Sum Reduction (skip)

- Sum 100,000 numbers on 100 processor UMA

- Each processor has ID:  $0 \leq P_n \leq 99$
- Partition 1000 numbers per processor
- Initial summation on each processor

```
sum[Pn] = 0;  
for (i = 1000*Pn;  
     i < 1000*(Pn+1); i = i + 1)  
    sum[Pn] = sum[Pn] + A[i];
```

- Now need to add these partial sums
  - Reduction: divide and conquer
  - Half the processors add pairs, then quarter, ...
  - Need to synchronize between reduction steps

# Example: Sum Reduction (skip)

```
half = 100;  
repeat
```

```
    synch();
```

```
    if (half%2 != 0 && Pn == 0)
```

```
        sum[0] = sum[0] + sum[half-1];
```


```
        /* Conditional sum needed when half is odd;
```

```
        Processor0 gets missing element */
```

```
    half = half/2; /* dividing line on who sums */
```

```
    if (Pn < half) sum[Pn] = sum[Pn] + sum[Pn+half];
```

```
until (half == 1);
```



# OpenMP (skip)

- A limited but popular example from hundreds of attempted parallel programming systems
- API, compiler directives, **environment variables and runtime libraries** that extend standard programming languages
  - Portable, scalable, simple programming model for SMP
  - Primary goal: parallelize loop & perform reduction
- Most C compilers support OpenMP
  - Use OpenMP API with UNIX C compiler  
cc -fopenmp foo.c
  - OpenMP extend C using *pragmas*  
#define P 100  
#pragma omp parallel num\_threads(P)

# OpenMP (skip)

- Parallelize loop (assuming *sum* is initialized to 0)

```
#pragma omp parallel for
    for (Pn = 0; Pn < P; Pn = Pn+1)
        for (i= 1000*Pn; i < 1000*(Pn+1); i = i+1)
            sum[Pn] = sum[Pn] + A[i];
```

- Perform reduction

```
#pragma omp parallel for reduction(+ Finalsum)
    for (i = 0; i < P; i += 1)
        Finalsum += sum[i];
```


- Programmers often use more sophisticated parallel programming systems than OpenMP

## **Section 5.10**

### **Parallelism and Memory Hierarchy: cache Coherence**

# Cache Coherence Problem

---



# Cache Coherence Problem

- Suppose two CPU cores share a physical address space
  - Write-through caches

Time step	Event	CPU A's cache	CPU B's cache	Memory
0				0
1	CPU A reads X	0		0
2	CPU B reads X	0	0	0
3	CPU A writes 1 to X	1	0	1

# Cache Coherence Protocols

- Operations performed by caches in multiprocessors to ensure coherence
  - Migration of data to local caches
    - Reduces bandwidth for shared memory
  - Replication of read-shared data
    - Reduces contention for access
- Snooping protocols
  - Each cache monitors bus reads/writes
- Directory-based protocols
  - Caches and memory record sharing status of blocks in a directory



## **Section 6.7**

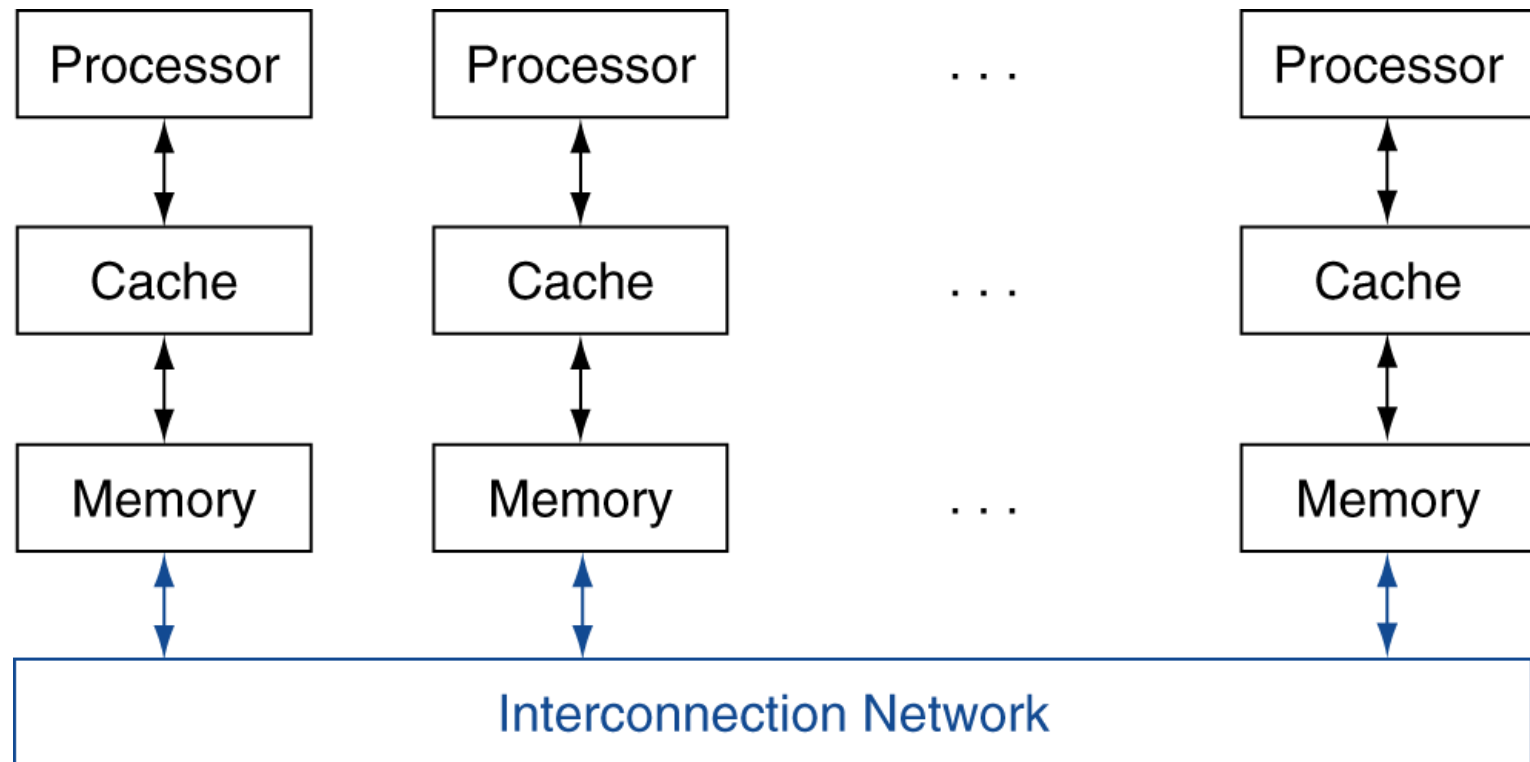
**Clusters, warehouse-scale computers, and other message-passing multiprocessors**

# Message Passing Machines

- Each processor has private physical address space
  - Alternative to sharing an address space
- Hardware sends/receives messages between processors (thus the name)
  - Explicit communication
  - Coordination is built in with message passing
- Interconnection network
  - Custom vs. local area networks like Ethernet
    - Many supercomputers use custom networks
    - Cost of custom network not justifiable for other applications

# Message Passing

- Easier to build much larger systems
- Internet services you depend on depend on these large scale systems



# Message Passing Machines (vs. SMPs)

## ■ Upside

- Much easier for hardware designers to build
- Explicit communication: an advantage to programmers
  - Fewer performance surprises than with implicit cache-coherent shared memory computers

## ■ Downside

- Harder to port sequential program to message passing computer
- Every communication must be identified in advance or the program doesn't work

# Sum Reduction (Again) (skip)

- Sum 100,000 on 100 processors
- First distribute 100 numbers to each
  - The do partial sums

```
sum = 0;
for (i = 0; i < 1000; i = i + 1)
    sum = sum + AN[i];
```
- Reduction
  - Half the processors send, other half receive and add
  - The quarter send, quarter receive and add, ...

# Sum Reduction (Again) (skip)

- Given send() and receive() operations

```
limit = 100; half = 100; /* 100 processors */
repeat
    half = (half+1)/2; /* send vs. receive
                        dividing line */
    if (Pn >= half && Pn < limit)
        send(Pn - half, sum);
    if (Pn < (limit/2))
        sum = sum + receive();
    limit = half; /* upper limit of senders */
until (half == 1); /* exit with final sum */
```

- Send/receive also provide synchronization
- Assumes send/receive take similar time to addition

# Message Passing vs. Shared Memory

- Which is right path to high-performance?
- No confusion in the marketplace
  - Multicore processors use shared physical memory
  - Nodes of a cluster communicate with each other with message passing
- What is a cluster?

# **Section 6.6**

## **Introduction to Graphics Processing Units**



# GPUs

- GPU: an exotic type of MIMD architecture
  - Different heritage and thus very different perspective on parallel programming challenge
  - Share address space
- Classic MIMD architectures
  - SMP
  - Clusters