

# Database Systems

## Lecture #10

Sang-Wook Kim  
Hanyang University

# Objectives



- ◆ To learn more other SQL commands
  - INSERT
  - DELETE
  - UPDATE
  - Views and indexes

# Outline



- ◆ INSERT Statement
- ◆ DELETE Statement
- ◆ UPDATE Statement
- ◆ Views
- ◆ Indexes

# INSERT Statement

- ◆ Used to add new tuples to a table
  - Modifies the table
- ◆ Types
  - Adding a single tuple
  - Adding a set of tuples

# INSERT Statement

- ◆ Adding a single tuple
  - **INSERT INTO** <table name>  
**VALUES** (<list of attributes values>);
  - Values should be listed *in the same order* in which the corresponding attributes were specified in the CREATE TABLE command

# INSERT Statement

## ◆ Adding a single tuple

### ● Example

```
U1:      INSERT INTO      EMPLOYEE  
        VALUES          ( 'Richard', 'K', 'Marini', '653298653', '1962-12-30', '98  
                           Oak Forest, Katy, TX', 'M', 37000, '653298653', 4 );
```

# INSERT Statement



- ◆ Adding a single tuple
  - Can specify some attributes partially
  - **INSERT INTO** <table name>(<list of attributes>)  
**VALUES** (<list of attribute values>);
  - Those attributes with NULL allowed or DEFAULT values can be *left out*
    - Will have NULL or DEFAULT value

# INSERT Statement

## ◆ Adding a single tuple

### ● Example

```
U1A:  INSERT INTO  EMPLOYEE (Fname, Lname, Dno, Ssn)
      VALUES      ('Richard', 'Marini', 4, '653298653');
```



# INSERT Statement

- ◆ Adding a set of tuples
  - Allows users to insert multiple tuples
    - From the *result of a query*
  - **INSERT INTO** <table name>(<list of attributes>)  
<select statement>;

# INSERT Statement

## ◆ Adding a set of tuples

- Example

```
U3A: CREATE TABLE WORKS_ON_INFO
      ( Emp_name      VARCHAR(15),
        Proj_name     VARCHAR(15),
        Hours_per_week DECIMAL(3,1) );
```

```
U3B:  INSERT INTO  WORKS_ON_INFO ( Emp_name, Proj_name,
        Hours_per_week )
        SELECT      E.Lname, P.Pname, W.Hours
        FROM         PROJECT P, WORKS_ON W, EMPLOYEE E
        WHERE        P.Pnumber=W.Pno AND W.Essn=E.Ssn;
```

- ◆ Adding a set of tuples
  - WORKS\_ON\_INFO table may not be up-to-date
    - Changes to PROJECT, WORKS\_ON, or EMPLOYEE table will not change WORKS\_ON\_INFO automatically
  - Use *views* to automatically reflect the changes of base tables

# DELETE Statement

- ◆ Removes tuples from a relation
  - Modifies the table
- ◆ Uses a WHERE clause to select the tuples to be deleted
- ◆ From only **one table** at a time

# DELETE Statement

- ◆ **DELETE FROM** <table name>  
**WHERE** <condition>;
- ◆ Missing WHERE clause
  - Unconditionally delete tuples
  - All the tuples in the relation are deleted

# DELETE Statement

## ◆ Examples

U4A:	DELETE FROM	EMPLOYEE
	WHERE	Lname='Brown';
U4B:	DELETE FROM	EMPLOYEE
	WHERE	Ssn='123456789';
U4C:	DELETE FROM	EMPLOYEE
	WHERE	Dno=5;
U4D:	DELETE FROM	EMPLOYEE;

# UPDATE Statement



- ◆ Modify attribute values of one or more tuples selected
  - Modifies the table
- ◆ Includes a WHERE clause to select the tuples to be modified
- ◆ From only one table at a time
- ◆ Additional SET clause in the UPDATE command
  - Specifies attributes to be modified and new values

# UPDATE Statement

- ◆ **UPDATE** <table name>  
    **SET** <list of 'attribute = value' pairs>  
    **WHERE** <condition>;
- ◆ Missing WHERE clause
  - Unconditionally modify tuples
  - All the tuples in the relation are modified



# UPDATE Statement



## ◆ Examples

**U5:**        **UPDATE**        PROJECT  
              **SET**            Plocation = 'Bellaire', Dnum = 5  
              **WHERE**        Pnumber=10;

# UPDATE Statement



## ◆ Examples

```
U6:      UPDATE      EMPLOYEE
          SET          Salary = Salary * 1.1
          WHERE        Dno = 5;
```

- New values for Salary is calculated by referring to the old values for Salary *before modification*

- ◆ Concept of a view in SQL
  - A single table *derived* from other tables
  - Considered to be a *virtual* table

## ◆ Characteristics of view

- View does not necessarily exist in a physical form
  - Limits the update operations that can be applied to views
- Querying views does not have any limitations
  - Retrieve related tuples from base tables

query tuple      -> 1. view  
                         2. query      select statement

## ◆ CREATE VIEW statement

- **CREATE VIEW** <view name> create      select .  
    **AS** <select statement>
- Note that the <view name> table does not exist physically
- Querying views retrieves tuples from <select statement>
- Basically, the attribute names for the view would be the same as those of the defining tables

# Creating Views



## ◆ Examples

V1:        **CREATE VIEW**        WORKS\_ON1  
          **AS SELECT**        Fname, Lname, Pname, Hours  
              **FROM**        EMPLOYEE, PROJECT, WORKS\_ON  
              **WHERE**        Ssn=Essn **AND** Pno=Pnumber;

**WORKS\_ON1**

Fname	Lname	Pname	Hours
-------	-------	-------	-------

## ◆ Examples

**V2:**      **CREATE VIEW**      DEPT\_INFO(Dept\_name, No\_of\_emps, Total\_sal)  
          **AS SELECT**        Dname, **COUNT** (\*), **SUM** (Salary)  
                              **FROM**        DEPARTMENT, EMPLOYEE  
                              **WHERE**       Dnumber=Dno  
                              **GROUP BY**   Dname;

### DEPT\_INFO

Dept_name	No_of_emps	Total_sal
-----------	------------	-----------

- Attribute names are specified explicitly

# Querying Views

- ◆ Queries on a view can be specified in the same way as a table
- ◆ Example
  - Retrieve the last name and first name of all employees who work on the 'ProductX' project

```
QV1:  SELECT  Fname, Lname
      FROM    WORKS_ON1
      WHERE   Pname='ProductX';
```



## ◆ Example

- Without WORKS\_ON1 view, QV1 is specified in a more complicate form

```
SELECT      Fname, Lname
FROM        EMPLOYEE, PROJECT, WORKS_ON
WHERE       Ssn=Essn AND Pno=Pnumber
              AND Pname='ProductX';
```

## ◆ Advantages of views

- Simplify the specification of certain queries
- Can be used as a security and authorization mechanism
  - Showing only a part of the physical tables
- Always up-to-date

- ◆ Disadvantage of views
  - Can cause a performance issue when a view is defined via a time-consuming query
  - Limited updates

# Dropping Views

- ◆ Dispose of a view when it is not needed any more
- ◆ **DROP VIEW** <view name>;
- ◆ Example

```
V1A: DROP VIEW WORKS_ON1;
```

## ◆ View update problem

- Updating views can be interpreted as updating underlying base tables
  - Only for some view updates
  - Can be ambiguous
- Research issue in the database field

# Updating Views

## ◆ Examples

- Update the PNAME attribute of 'John Smith' from 'ProductX' to 'ProductY'

```
UV1:  UPDATE WORKS_ON1
      SET      Pname = 'ProductY'
      WHERE    Lname='Smith' AND Fname='John'
            AND Pname='ProductX';
```

# Updating Views

## ◆ Examples: Approach 1

- Change the name of the 'ProductX' tuple in the PROJECT relation to 'ProductY'

```
UPDATE PROJECT SET    Pname = 'ProductY'  
WHERE      Pname = 'ProductX';
```

# Updating Views

## ◆ Examples: Approach 2

- Relate 'John Smith' to the 'ProductY' PROJECT tuple instead of the 'ProductX' PROJECT tuple

```

UPDATE WORKS_ON
SET      Pno = ( SELECT Pnumber
                  FROM   PROJECT
                  WHERE  Pname='ProductY' )

WHERE    Essn IN ( SELECT Ssn
                   FROM   EMPLOYEE
                   WHERE  Lname='Smith' AND Fname='John' )

AND
Pno = ( SELECT Pnumber
        FROM   PROJECT
        WHERE  Pname='ProductX' );
    
```



# Updating Views

## ◆ Examples: Impossible update

```
UV2:    UPDATE DEPT_INFO  
        SET      Total_sal=100000  
        WHERE    Dname='Research';
```

- Total\_sal is defined to be the sum of the individual employee salaries
  - Not stored in the physical base table

## ◆ Summary

- Cannot guarantee successful updates for some queries on a view
- Possibilities
  - A view with a single defining table is updatable if the view attributes contain the primary key of the base relation
  - Views defined on multiple tables using joins are generally not updatable.
  - Views defined using grouping and aggregate functions are not updatable

- ◆ Additional meta-data to support fast retrieval
- ◆ Indexing attributes:
  - Attributes used to index a table
- ◆ If indexing attributes are used in the condition of a retrieval query, the query can be processed very fast

# Creating Indexes

- ◆ **CREATE INDEX** <index name>  
**ON** <table name>(<list of attributes>);
- ◆ Create an index on <list of attributes> for  
<table name>

# Creating Indexes



## ◆ Example

- **CREATE INDEX** LnameIndex  
**ON** EMPLOYEE (Lname);

- ◆ Order of the attribute values in indexes
  - Ascending order (ASC) by default
  - Descending order can be specified by DESC keyword

# Creating Indexes



## ◆ Example

- **CREATE INDEX** LnameIndex  
**ON** EMPLOYEE (Lname **DESC**);

# Creating Indexes



- ◆ Index can be created on multiple attributes

- ◆ Example

- **CREATE INDEX** NameIndex  
**ON** EMPLOYEE (Lname, Fname, Minit);



# UNIQUE Keyword



- ◆ Used for specifying key constraint on indexing attributes
- ◆ Example
  - **CREATE UNIQUE INDEX** SsnIndex  
**ON** EMPLOYEE (Ssn);

# CLUSTER Keyword



- ◆ Used when the index to be created should also sort the data file records on the indexing attributes
- ◆ Can improve query performance if a query requires join or ranged condition on the indexing attributes

## ◆ Example

```
CREATE INDEX DnoIndex  
ON EMPLOYEE (Dno)  
CLUSTER ;
```

Dno                      group                      disk  
<disadvantage>  
1. cluster index  
2. insert                      가                      .  
(                      .)

# Dropping Indexes

- ◆ Drop an index when it is not needed any more
- ◆ **DROP INDEX** <index name>;
- ◆ Example
  - **DROP INDEX** DnoIndex;

## ◆ Database update commands

- INSERT
- DELETE
- UPDATE

## ◆ Views

## ◆ Indexes

# References



1. Dayal, Umeshwar, and Philip A. Bernstein. "On the Updatability of Relational Views." *VLDB*. Vol. 78. 1978.
2. Keller, Arthur M. *Updates to relational databases through views involving joins*. IBM Research Division, 1981.
3. Langerak, Rom. "View updates in relational databases with an independent scheme." *ACM Transactions on Database Systems (TODS)* 15.1 (1990): 40-66.
4. Blakeley, José A., Neil Coburn, and Per Larson. "Updating derived relations: Detecting irrelevant and autonomously computable updates." *ACM Transactions on Database Systems (TODS)* 14.3 (1989): 369-400.

# References



5. Negri, Mauro, Giuseppe Pelagatti, and Licia Sbattella. "Formal semantics of SQL queries." *ACM Transactions on Database Systems (TODS)* 16.3 (1991): 513-534.
6. Melton, Jim, and Alan R. Simon. *SQL: 1999: understanding relational language components*. Morgan Kaufmann, 2001.
7. Melton, Jim. *Advanced SQL: 1999: Understanding object-relational and other advanced features*. Morgan Kaufmann, 2002.

Have a nice day!