

Database Systems

Lecture #10

Sang-Wook Kim
Hanyang University

Objectives



- ◆ To learn more other SQL commands
 - INSERT
 - DELETE
 - UPDATE
 - Views and indexes

Outline



- ◆ INSERT Statement
- ◆ DELETE Statement
- ◆ UPDATE Statement
- ◆ Views
- ◆ Indexes

INSERT Statement

- ◆ Used to add new tuples to a table
 - Modifies the table
- ◆ Types
 - Adding a single tuple
 - Adding a set of tuples

INSERT Statement

- ◆ Adding a single tuple
 - **INSERT INTO** <table name>
VALUES (<list of attributes values>);
 - Values should be listed *in the same order* in which the corresponding attributes were specified in the CREATE TABLE command

INSERT Statement

◆ Adding a single tuple

● Example

```
U1:    INSERT INTO    EMPLOYEE  
        VALUES      ( 'Richard', 'K', 'Marini', '653298653', '1962-12-30', '98  
                        Oak Forest, Katy, TX', 'M', 37000, '653298653', 4 );
```

INSERT Statement



- ◆ Adding a single tuple
 - Can specify some attributes partially
 - **INSERT INTO** <table name>(<list of attributes>)
VALUES (<list of attribute values>);
 - Those attributes with NULL allowed or DEFAULT values can be *left out*
 - Will have NULL or DEFAULT value

INSERT Statement

◆ Adding a single tuple

● Example

```
U1A:  INSERT INTO  EMPLOYEE (Fname, Lname, Dno, Ssn)
      VALUES      ('Richard', 'Marini', 4, '653298653');
```


INSERT Statement

- ◆ Adding a set of tuples
 - Allows users to insert multiple tuples
 - From the *result of a query*
 - **INSERT INTO** <table name>(<list of attributes>)
<select statement>;

INSERT Statement

◆ Adding a set of tuples

● Example

U3A: **CREATE TABLE** WORKS_ON_INFO
 (Emp_name VARCHAR(15),
 Proj_name VARCHAR(15),
 Hours_per_week DECIMAL(3,1));

U3B: **INSERT INTO** WORKS_ON_INFO (Emp_name, Proj_name,
 Hours_per_week)
SELECT E.Lname, P.Pname, W.Hours
FROM PROJECT P, WORKS_ON W, EMPLOYEE E
WHERE P.Pnumber=W.Pno **AND** W.Essn=E.Ssn;

- ◆ Adding a set of tuples
 - WORKS_ON_INFO table may not be up-to-date
 - Changes to PROJECT, WORKS_ON, or EMPLOYEE table will not change WORKS_ON_INFO automatically
 - Use *views* to automatically reflect the changes of base tables

DELETE Statement

- ◆ Removes tuples from a relation
 - Modifies the table
- ◆ Uses a WHERE clause to select the tuples to be deleted
- ◆ From only **one table** at a time

DELETE Statement

- ◆ **DELETE FROM** <table name>
WHERE <condition>;
- ◆ Missing WHERE clause
 - Unconditionally delete tuples
 - All the tuples in the relation are deleted

DELETE Statement

◆ Examples

U4A:	DELETE FROM WHERE	EMPLOYEE Lname='Brown';
U4B:	DELETE FROM WHERE	EMPLOYEE Ssn='123456789';
U4C:	DELETE FROM WHERE	EMPLOYEE Dno=5;
U4D:	DELETE FROM	EMPLOYEE;

UPDATE Statement



- ◆ Modify attribute values of one or more tuples selected
 - Modifies the table
- ◆ Includes a WHERE clause to select the tuples to be modified
- ◆ From only one table at a time
- ◆ Additional SET clause in the UPDATE command
 - Specifies attributes to be modified and new values

UPDATE Statement

- ◆ **UPDATE** <table name>
 SET <list of 'attribute = value' pairs>
 WHERE <condition>;
- ◆ Missing WHERE clause
 - Unconditionally modify tuples
 - All the tuples in the relation are modified

UPDATE Statement



◆ Examples

U5: **UPDATE** PROJECT
 SET Plocation = 'Bellaire', Dnum = 5
 WHERE Pnumber=10;

UPDATE Statement

◆ Examples

```
U6:      UPDATE      EMPLOYEE
         SET          Salary = Salary * 1.1
         WHERE        Dno = 5;
```

- New values for Salary is calculated by referring to the old values for Salary *before modification*

- ◆ Concept of a view in SQL
 - A single table *derived* from other tables
 - Considered to be a *virtual* table

◆ Characteristics of view

- View does not necessarily exist in a physical form
 - Limits the update operations that can be applied to views
- Querying views does not have any limitations
 - Retrieve related tuples from base tables

◆ CREATE VIEW statement

- **CREATE VIEW** <view name>
AS <select statement>
- Note that the <view name> table does not exist physically
- Querying views retrieves tuples from <select statement>
- Basically, the attribute names for the view would be the same as those of the defining tables

Creating Views



◆ Examples

V1: **CREATE VIEW** WORKS_ON1
 AS SELECT Fname, Lname, Pname, Hours
 FROM EMPLOYEE, PROJECT, WORKS_ON
 WHERE Ssn=Essn **AND** Pno=Pnumber;

WORKS_ON1

Fname	Lname	Pname	Hours
-------	-------	-------	-------

Creating Views



◆ Examples

V2: **CREATE VIEW** DEPT_INFO(Dept_name, No_of_emps, Total_sal)
 AS SELECT Dname, **COUNT** (*), **SUM** (Salary)
 FROM DEPARTMENT, EMPLOYEE
 WHERE Dnumber=Dno
 GROUP BY Dname;

DEPT_INFO

Dept_name	No_of_emps	Total_sal
-----------	------------	-----------

- Attribute names are specified explicitly

Querying Views

- ◆ Queries on a view can be specified in the same way as a table
- ◆ Example
 - Retrieve the last name and first name of all employees who work on the 'ProductX' project

```
QV1:  SELECT  Fname, Lname
      FROM    WORKS_ON1
      WHERE   Pname='ProductX';
```


Querying Views

◆ Example

- Without WORKS_ON1 view, QV1 is specified in a more complicate form

```
SELECT      Fname, Lname
FROM        EMPLOYEE, PROJECT, WORKS_ON
WHERE       Ssn=Essn AND Pno=Pnumber
              AND Pname='ProductX';
```

◆ Advantages of views

- Simplify the specification of certain queries
- Can be used as a security and authorization mechanism
 - Showing only a part of the physical tables
- Always up-to-date

- ◆ Disadvantage of views
 - Can cause a performance issue when a view is defined via a time-consuming query
 - Limited updates

Dropping Views

- ◆ Dispose of a view when it is not needed any more
- ◆ **DROP VIEW** <view name>;
- ◆ Example

```
V1A: DROP VIEW WORKS_ON1;
```

◆ View update problem

- Updating views can be interpreted as updating underlying base tables
 - Only for some view updates
 - Can be ambiguous
- Research issue in the database field

Updating Views

◆ Examples

- Update the PNAME attribute of 'John Smith' from 'ProductX' to 'ProductY'

```
UV1:  UPDATE WORKS_ON1
      SET      Pname = 'ProductY'
      WHERE    Lname='Smith' AND Fname='John'
            AND Pname='ProductX';
```

Updating Views

◆ Examples: Approach 1

- Change the name of the 'ProductX' tuple in the PROJECT relation to 'ProductY'

```
UPDATE PROJECT SET    Pname = 'ProductY'  
WHERE      Pname = 'ProductX';
```

Updating Views

◆ Examples: Approach 2

- Relate 'John Smith' to the 'ProductY' PROJECT tuple instead of the 'ProductX' PROJECT tuple

```

UPDATE WORKS_ON
SET      Pno = ( SELECT Pnumber
                  FROM   PROJECT
                  WHERE  Pname='ProductY' )

WHERE    Essn IN ( SELECT Ssn
                   FROM   EMPLOYEE
                   WHERE  Lname='Smith' AND Fname='John' )

AND
Pno = ( SELECT Pnumber
        FROM   PROJECT
        WHERE  Pname='ProductX' );
  
```


Updating Views

◆ Examples: Impossible update

```
UV2:    UPDATE DEPT_INFO
        SET      Total_sal=100000
        WHERE    Dname='Research';
```

- Total_sal is defined to be the sum of the individual employee salaries
 - Not stored in the physical base table

◆ Summary

- Cannot guarantee successful updates for some queries on a view
- Possibilities
 - A view with a single defining table is updatable if the view attributes contain the primary key of the base relation
 - Views defined on multiple tables using joins are generally not updatable.
 - Views defined using grouping and aggregate functions are not updatable

- ◆ Additional meta-data to support fast retrieval
- ◆ Indexing attributes:
 - Attributes used to index a table
- ◆ If indexing attributes are used in the condition of a retrieval query, the query can be processed very fast

Creating Indexes

- ◆ **CREATE INDEX** <index name>
ON <table name>(<list of attributes>);
- ◆ Create an index on <list of attributes> for
<table name>

Creating Indexes



◆ Example

- **CREATE INDEX** LnameIndex
ON EMPLOYEE (Lname);

- ◆ Order of the attribute values in indexes
 - Ascending order (ASC) by default
 - Descending order can be specified by DESC keyword

Creating Indexes



◆ Example

- **CREATE INDEX** LnameIndex
ON EMPLOYEE (Lname **DESC**);

Creating Indexes



- ◆ Index can be created on multiple attributes

- ◆ Example

- **CREATE INDEX** NameIndex
ON EMPLOYEE (Lname, Fname, Minit);

UNIQUE Keyword



- ◆ Used for specifying key constraint on indexing attributes
- ◆ Example
 - **CREATE UNIQUE INDEX** SsnIndex
ON EMPLOYEE (Ssn);

CLUSTER Keyword



- ◆ Used when the index to be created should also sort the data file records on the indexing attributes
- ◆ Can improve query performance if a query requires join or ranged condition on the indexing attributes

CLUSTER Keyword



◆ Example

```
CREATE INDEX DnoIndex  
ON EMPLOYEE (Dno)  
CLUSTER ;
```

Dropping Indexes

- ◆ Drop an index when it is not needed any more
- ◆ **DROP INDEX** <index name>;
- ◆ Example
 - **DROP INDEX** DnoIndex;

◆ Database update commands

- INSERT
- DELETE
- UPDATE

◆ Views

◆ Indexes

References



1. Dayal, Umeshwar, and Philip A. Bernstein. "On the Updatability of Relational Views." *VLDB*. Vol. 78. 1978.
2. Keller, Arthur M. *Updates to relational databases through views involving joins*. IBM Research Division, 1981.
3. Langerak, Rom. "View updates in relational databases with an independent scheme." *ACM Transactions on Database Systems (TODS)* 15.1 (1990): 40-66.
4. Blakeley, José A., Neil Coburn, and Per Larson. "Updating derived relations: Detecting irrelevant and autonomously computable updates." *ACM Transactions on Database Systems (TODS)* 14.3 (1989): 369-400.

References



5. Negri, Mauro, Giuseppe Pelagatti, and Licia Sbattella. "Formal semantics of SQL queries." *ACM Transactions on Database Systems (TODS)* 16.3 (1991): 513-534.
6. Melton, Jim, and Alan R. Simon. *SQL: 1999: understanding relational language components*. Morgan Kaufmann, 2001.
7. Melton, Jim. *Advanced SQL: 1999: Understanding object-relational and other advanced features*. Morgan Kaufmann, 2002.

Have a nice day!