
PERFORMANCE PROFILING

Minsoo Ryu

**Real-Time Computing and Communications Lab.
Hanyang University**

msryu@hanyang.ac.kr

OUTLINE

- ❑ HISTORY
- ❑ UNDERSTANDING PROFILING
- ❑ UNDERSTANDING PERFORMANCE
- ❑ UNDERSTANDING PERFORMANCE PROFILING
- ❑ PERFORMANCE PROFILING TYPES
 - PERFORMANCE COUNTERS
 - EVENT BASED PROFILING
 - STATISTICAL PROFILING
 - INSTRUMENTATION PROFILING
 - HYPERVISOR/SIMULATOR
- ❑ PERFORMANCE PROFILING TOOLS

HISTORY

HISTORY

❑ 1970

- Usually based on timer interrupts which recorded the program status word (PSW) at set timer-intervals to detect "hot spots" in executing code
 - Instruction-set simulators permitted full trace and other performance-monitoring features in 1974

❑ 1980

- Profiler-driven program analysis on Unix dates back to at least 1979, when unix systems included a basic tool
 - Prof which listed each function and how much of program execution time it used
 - Gprof extended the concept to a complete call graph analysis in 1982

❑ 1990

- The ATOM(analysis tools with OM) platform converts a program into its own profiler
 - It inserts code into the program to be analyzed: instrumentation profiling
 - ✓ That inserted code outputs analysis data, at compile time

UNDERSTANDING PROFILING

UNDERSTANDING PROFILING

□ What is measured

- Time
- Control flow
 - Loop counts, Function calls
- Aliasing facts
- Cache stats
- Allocation information
 - Track allocation sites for objects
- Hardware stats
- Granularity of what is measured
 - Instructions, basic blocks, line of code, function, modules

UNDERSTANDING PROFILING

□ How are measurements taken

- Instrumentation
 - – The code is modified to take the measurements
 - – When?
 - ✓ At compile time: static time
 - ✓ At runtime: dynamic time
- Interruption
 - – An outside event triggers inspection and measurement
 - – Who?
 - ✓ Hardware
 - ✓ Timer
 - ✓ Another thread

UNDERSTANDING PROFILING

□ When are measurements taken

- All the time
 - – Expensive
- Sampling
 - – Cheaper
 - – When?
 - ✓ N_th function call
 - ✓ N_th basic block
 - ✓ Timer
 - ✓ Some property of the hardware

UNDERSTANDING PERFORMANCE

UNDERSTANDING PERFORMANCE

- ❑ We need visibility into what the machine is doing
 - To understand low-level performance
 - We mostly only control data and code
- ❑ We need to know how it interacts with HW
 - But to make an existing algorithm/data-structure work better
- ❑ Sometimes it is enough to know time
 - – Algorithm-level fix, data-structure-level fix
- ❑ Part of performance measurements
 - – CPU
 - – Memory
 - – Network
 - – Disk
 - – SQL DB
 - – User Input

UNDERSTANDING PERFORMANCE

□ Aspects of performance

▪ Availability

- Availability of a system is typically measured as a factor of its reliability
 - ✓ As reliability increases, so does availability

▪ Response time

- Response time is the total amount of time it takes to respond to a request for service
- The response time is the sum of three numbers
 - ✓ Service time + wait time + transmission time

▪ Throughput

- Throughput is the rate of production or the rate at which something can be processed

UNDERSTANDING PERFORMANCE

□ Aspects of performance

▪ Latency

- Latency is a time delay between the cause and the effect of some physical change in the system being observed.
- Latency is a result of the limited velocity with which any physical interaction can take place.
 - ✓ This velocity is always lower or equal to speed of light

▪ Scalability

- Scalability is the ability of a system or process to handle a growing amount of work in a capable manner or its ability to be enlarged to accommodate that growth

UNDERSTANDING PERFORMANCE PROFILING

UNDERSTANDING PERFORMANCE PROFILING

□ What is a hotspot

- Where in an application or system there is a significant amount of activity
 - Significant: activity that occurs infrequently probably does not have much impact on system performance
 - Activity: time spent or other internal processor event
- Examples of other events: cache misses, branch mispredictions, floating-point instructions retired, partial register stalls, and so on

UNDERSTANDING PERFORMANCE PROFILING

- ❑ In theory, the " algorithm " and " data structure " can be predicted through knowledge. but, performance profiling requires precise measurements for the possible prediction results
 - Even though " the same problem " , substantially significant differences occur
 - The hotspot can be found through precise measurements

- ❑ Performance profiling programming
 - Knowledge of all the layers involved
 - Experience in knowing when and how performance can be a problem
 - Skill in detecting and zooming in on the problems
 - A good dose of common sense

PERFORMANCE PROFILING TYPES

PERFORMANCE PROFILING TYPES

- ☐ PERFORMANCE COUNTERS
- ☐ EVENT BASED PROFILING
- ☐ STATISTICAL PROFILING
- ☐ INSTRUMENTATION PROFILING
- ☐ HYPERVISOR/SIMULATOR

PERFORMANCE COUNTERS

❑ Performance counters

- We need to count the events
- Program a counter to count a specific event
- HW provides special programmable counters
- Software can read and write the counters

❑ What's wrong with performance counters

- Should be solved the coarse-granularity problem
- Extend performance counters with a programmable limit value
 - Cause an interrupt when the counter overflows

EVENT BASED PROFILING

❑ Event based profiling is triggered by any one of the software-based events or any performance monitor event that occurs on the processor

❑ **Event-based profiling: advantages**

- The routine addresses are visible when interrupts are disabled
- The ability to vary the profiling event
- The ability to vary the sampling frequency

❑ **Event-based profiling: disadvantages**

- Limited data collection
 - Per run
 - Types of data
- In-exact
 - Sample point and event cause don't match

STATISTICAL PROFILING

- ❑ A method of profiling an application by taking samples of the program address at regular intervals while the application is executing

- ❑ Cycle accurate profiling methods
 - These samples are then associated with either a specific function or a specific memory range
 - A simple statistical analysis is performed to determine the areas where an applications spends the largest portions of its cycles
 - Statistical Profiling is a very quick and handy way to get a first look at which functions are consuming the largest proportions of cycles

STATISTICAL PROFILING

❑ The statistical profiling: advantages

- No installation required
- Wide coverage
- Low overhead

❑ The statistical profiling: disadvantages

- Approximate precision
- Limited report

INSTRUMENTATION PROFILING

- ❑ A method of profiling how to get the information by inserting an additional command to the code
 - The performance impact depends on the amount of detail in the data type
 - To measure the execution frequency of each line VS execution times of the function
 - Method
 - Programmer to insert code by hand
 - Automatic code insertion tool
 - How to insert the code with the help of the compiler
 - The final non-binary conversion code
 - Modify the binaries to run immediately before
 - Operate in protected mode while running

INSTRUMENTATION PROFILING

❑ The instrumentation profiling: advantages

- Perfect accuracy ❑
 - Where you visit immediately before and after your
 - Can calculate how much time you spent at each site
 - How many times you visited each site

❑ The instrumentation profiling: disadvantages

- Low granularity ❑
 - Too coarse
- High overhead ❑
- High touch ❑
 - I have to build all those area, which expands the space in each site you visit

HYPERVISOR/SIMULATOR

□ Hypervisor/Simulator

- Hypervisor: data are collected by running the unmodified program under a hypervisor
- Simulator: data collected interactively and selectively by running the unmodified program under an instruction set simulator

PERFORMANCE PROFILING TOOLS

PERFORMANCE PROFILING TOOLS

In order to do	You can use
Performance counters	Perf, Oprofile
Event based profiling	Ruby, JIT
Statistical profiling	Oprofile, Vtune
Instrumentation profiling	Gprof, Valgrind
Hypervisor/Simulator	SIMMON, OLIVER
Other tools exist for Network, Disk IO...	

수고하셨습니다.